

```
In [1]: import pandas as pd
```

```
In [2]: import os
import torch
from pycocotools.coco import COCO
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as T
from torchvision.models.detection import fasterrcnn_resnet50_fpn
import cv2
```

```
In [3]: torch.cuda.empty_cache()
```

```
In [4]: class CarDamageDataset(Dataset):
    def __init__(self, root, annFile, transforms=None):
        """
        root: Directory where images are stored
        annFile: Path to the COCO-format annotation file
        transforms: Any transformations to apply to the images
        """
        self.root = root
        self.coco = COCO(annFile)
        self.ids = list(self.coco.imgs.keys()) # List of image IDs
        self.transforms = transforms

        # Remove images with no annotations (i.e., no bounding boxes)
        self.ids = [img_id for img_id in self.ids if len(self.coco.getAnnIds(imgIds=img_id)) > 0]

    def __getitem__(self, idx):
        # Get image ID and associated file name
        img_id = self.ids[idx]
        img_info = self.coco.loadImgs(img_id)[0]
        path = img_info['file_name']

        # Load image using file name
        img = cv2.imread(os.path.join(self.root, path))
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

        # Get annotations for this image
        ann_ids = self.coco.getAnnIds(imgIds=img_id)
        annotations = self.coco.loadAnns(ann_ids)

        # Prepare bounding boxes and labels
        boxes = []
        labels = []
        areas = []
        iscrowd = []
        for ann in annotations:
            #print(ann['bbox'])
            xmin, ymin, width, height = ann['bbox']
            boxes.append([xmin, ymin, xmin + width, ymin + height])
            labels.append(ann['category_id']) # category ID for damaged part
            areas.append(ann['area']) # area of the bounding box
            iscrowd.append(ann['iscrowd']) # crowd annotation

        # Convert to tensors
        target = {
            "boxes": torch.tensor(boxes, dtype=torch.float32),
            "labels": torch.tensor(labels, dtype=torch.int64),
            "area": torch.tensor(areas, dtype=torch.float32),
            "iscrowd": torch.tensor(iscrowd, dtype=torch.int64)
        }

        # Apply transforms if any
        if self.transforms:
            img = self.transforms(img)

        return img, target

    def __len__(self):
        return len(self.ids)
```

```
In [5]: # Paths to image directories and COCO-format annotation files for each split
train_root = "Part_Detection_final_dataset/train"
val_root = "Part_Detection_final_dataset/valid"
test_root = "Part_Detection_final_dataset/test"
train_annFile = "Part_Detection_final_dataset/annotations/train_annotations.coco.json"
val_annFile = "Part_Detection_final_dataset/annotations/val_annotations.coco.json"
test_annFile = "Part_Detection_final_dataset/annotations/test_annotations.coco.json"

# Initialize datasets for each split
train_dataset = CarDamageDataset(root=train_root, annFile=train_annFile, transforms=T.ToTensor())
```

```

val_dataset = CarDamageDataset(root=val_root, annFile=val_annFile, transforms=T.ToTensor())
test_dataset = CarDamageDataset(root=test_root, annFile=test_annFile, transforms=T.ToTensor())

loading annotations into memory...
Done (t=0.18s)
creating index...
index created!
loading annotations into memory...
Done (t=0.03s)
creating index...
index created!
loading annotations into memory...
Done (t=0.01s)
creating index...
index created!

```

## Paths to image directories and COCO-format annotation files for each split

```

train_root = "Part_Detection_final_dataset/train" val_root = "Part_Detection_final_dataset/val" test_root =
"Part_Detection_final_dataset/test"

train_annFile = "Part_Detection_final_dataset/annotations/train_annotations.coco.json" val_annFile =
"Part_Detection_final_dataset/annotations/val_annotations.coco.json" test_annFile =
"Part_Detection_final_dataset/annotations/test_annotations.coco.json"

```

## Initialize datasets for each split

```

train_dataset = CarDamageDataset(root=train_root, annFile=train_annFile, transforms=T.ToTensor()) val_dataset =
CarDamageDataset(root=val_root, annFile=val_annFile, transforms=T.ToTensor()) test_dataset = CarDamageDataset(root=test_root,
annFile=test_annFile, transforms=T.ToTensor())

```

```
In [6]: # Create data loaders for each dataset split
train_loader = DataLoader(train_dataset, batch_size=4, shuffle=True, pin_memory=True, collate_fn=lambda x: tuple(
    val_loader = DataLoader(val_dataset, batch_size=4, shuffle=False, collate_fn=lambda x: tuple(zip(*x)))
    test_loader = DataLoader(test_dataset, batch_size=4, shuffle=True, collate_fn=lambda x: tuple(zip(*x)))
)
```

```
In [7]: from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
```

```
In [9]: # Model setup
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')
num_classes = 15 # Assuming 6 damage categories + 1 background class
model = fasterrcnn_resnet50_fpn(pretrained=True)
in_features = model.roi_heads.box_predictor.cls_score.in_features
model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)
model.to(device)

# Define the optimizer
optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9, weight_decay=0.0005)
```

```
C:\Users\harsh\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\torchvision\models\_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
  warnings.warn(
C:\Users\harsh\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\torchvision\models\_utils.py:223: UserWarning: Arguments other than a weight enum or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=FasterRCNN_ResNet50_FPN_Weights.COCO_V1`. You can also use `weights=FasterRCNN_ResNet50_FPN_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
```

```
In [10]: from tqdm import tqdm

# Training and validation loop
num_epochs = 15
for epoch in range(num_epochs):
    model.train()
    train_loss = 0

    # Add progress bar to the training loop
    progress_bar = tqdm(train_loader, desc=f"Epoch [{epoch+1}/{num_epochs}]")
    for images, targets in progress_bar:
        images = [img.to(device) for img in images]
        targets = [{k: v.to(device) for k, v in t.items()} for t in targets]

        # Forward pass
```

```

loss_dict = model(images, targets)
losses = sum(loss for loss in loss_dict.values())
train_loss += losses.item()

# Backward pass
optimizer.zero_grad()
losses.backward()
optimizer.step()

# Update progress bar with current loss
progress_bar.set_postfix(loss=losses.item())

print(f"Epoch [{epoch+1}/{num_epochs}], Train Loss: {train_loss/len(train_loader):.4f}")

```

```

Epoch [1/15]: 100%|██████████| 1501/1501 [1:28:14<00:00,  3.53s/it, loss=0.13]
Epoch [1/15], Train Loss: 0.2202
Epoch [2/15]: 100%|██████████| 1501/1501 [1:27:27<00:00,  3.50s/it, loss=0.158]
Epoch [2/15], Train Loss: 0.1518
Epoch [3/15]: 100%|██████████| 1501/1501 [1:27:40<00:00,  3.50s/it, loss=0.0695]
Epoch [3/15], Train Loss: 0.1292
Epoch [4/15]: 100%|██████████| 1501/1501 [1:27:16<00:00,  3.49s/it, loss=0.156]
Epoch [4/15], Train Loss: 0.1196
Epoch [5/15]: 100%|██████████| 1501/1501 [1:28:33<00:00,  3.54s/it, loss=0.0732]
Epoch [5/15], Train Loss: 0.1122
Epoch [6/15]: 100%|██████████| 1501/1501 [1:27:15<00:00,  3.49s/it, loss=0.0856]
Epoch [6/15], Train Loss: 0.1081
Epoch [7/15]: 100%|██████████| 1501/1501 [1:42:28<00:00,  4.10s/it, loss=0.072]
Epoch [7/15], Train Loss: 0.1057
Epoch [8/15]: 100%|██████████| 1501/1501 [1:43:44<00:00,  4.15s/it, loss=0.0905]
Epoch [8/15], Train Loss: 0.1007
Epoch [9/15]: 100%|██████████| 1501/1501 [1:45:44<00:00,  4.23s/it, loss=0.138]
Epoch [9/15], Train Loss: 0.0983
Epoch [10/15]: 100%|██████████| 1501/1501 [1:39:08<00:00,  3.96s/it, loss=0.102]
Epoch [10/15], Train Loss: 0.0971
Epoch [11/15]: 100%|██████████| 1501/1501 [1:48:17<00:00,  4.33s/it, loss=0.176]
Epoch [11/15], Train Loss: 0.1035
Epoch [12/15]: 100%|██████████| 1501/1501 [1:56:48<00:00,  4.67s/it, loss=0.0851]
Epoch [12/15], Train Loss: 0.1015
Epoch [13/15]: 100%|██████████| 1501/1501 [1:56:54<00:00,  4.67s/it, loss=0.0981]
Epoch [13/15], Train Loss: 0.0918
Epoch [14/15]: 100%|██████████| 1501/1501 [1:56:37<00:00,  4.66s/it, loss=0.103]
Epoch [14/15], Train Loss: 0.0898
Epoch [15/15]: 100%|██████████| 1501/1501 [1:56:50<00:00,  4.67s/it, loss=0.0917]
Epoch [15/15], Train Loss: 0.0898

```

```
In [13]: print(f"Images format: {type(images)}, Target format: {type(targets)}")
```

```
Images format: <class 'list'>, Target format: <class 'list'>
```

```
In [ ]:
```

```
In [15]: # Save the model's state_dict to a file
torch.save(model.state_dict(), 'Part_presetion_model52_state_dict.pkl')
```

```
In [16]: torch.save(model, 'Part_preditction_model_full_20241125.pkl')
```

```
In [8]: model = torch.load('Part_preditction_model_full_20241125.pkl', map_location=torch.device('cpu'))
device=torch.device('cpu')
model.to(device)
```

```
Out[8]: FasterRCNN(
  (transform): GeneralizedRCNNTransform(
    Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    Resize(min_size=(800,), max_size=1333, mode='bilinear')
  )
  (backbone): BackboneWithFPN(
    (body): IntermediateLayerGetter(
      (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (bn1): FrozenBatchNorm2d(64, eps=0.0)
      (relu): ReLU(inplace=True)
      (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    )
    (layer1): Sequential(
      (0): Bottleneck(
        (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): FrozenBatchNorm2d(64, eps=0.0)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): FrozenBatchNorm2d(64, eps=0.0)
        (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): FrozenBatchNorm2d(256, eps=0.0)
      )
    )
  )
)
```

```

        (relu): ReLU(inplace=True)
        (downsample): Sequential(
            (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): FrozenBatchNorm2d(256, eps=0.0)
        )
    )
(1): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(64, eps=0.0)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(64, eps=0.0)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(256, eps=0.0)
    (relu): ReLU(inplace=True)
)
(2): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(64, eps=0.0)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(64, eps=0.0)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(256, eps=0.0)
    (relu): ReLU(inplace=True)
)
)
(layer2): Sequential(
(0): Bottleneck(
    (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(128, eps=0.0)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(128, eps=0.0)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(512, eps=0.0)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): FrozenBatchNorm2d(512, eps=0.0)
    )
)
(1): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(128, eps=0.0)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(128, eps=0.0)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(512, eps=0.0)
    (relu): ReLU(inplace=True)
)
(2): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(128, eps=0.0)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(128, eps=0.0)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(512, eps=0.0)
    (relu): ReLU(inplace=True)
)
(3): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(128, eps=0.0)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(128, eps=0.0)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(512, eps=0.0)
    (relu): ReLU(inplace=True)
)
)
(layer3): Sequential(
(0): Bottleneck(
    (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(256, eps=0.0)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(256, eps=0.0)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(1024, eps=0.0)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
        (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): FrozenBatchNorm2d(1024, eps=0.0)
    )
)
(1): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

```

```

(bn1): FrozenBatchNorm2d(256, eps=0.0)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): FrozenBatchNorm2d(256, eps=0.0)
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): FrozenBatchNorm2d(1024, eps=0.0)
(relu): ReLU(inplace=True)
)
(2): Bottleneck(
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): FrozenBatchNorm2d(256, eps=0.0)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): FrozenBatchNorm2d(256, eps=0.0)
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): FrozenBatchNorm2d(1024, eps=0.0)
(relu): ReLU(inplace=True)
)
(3): Bottleneck(
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): FrozenBatchNorm2d(256, eps=0.0)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): FrozenBatchNorm2d(256, eps=0.0)
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): FrozenBatchNorm2d(1024, eps=0.0)
(relu): ReLU(inplace=True)
)
(4): Bottleneck(
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): FrozenBatchNorm2d(256, eps=0.0)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): FrozenBatchNorm2d(256, eps=0.0)
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): FrozenBatchNorm2d(1024, eps=0.0)
(relu): ReLU(inplace=True)
)
(5): Bottleneck(
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): FrozenBatchNorm2d(256, eps=0.0)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): FrozenBatchNorm2d(256, eps=0.0)
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): FrozenBatchNorm2d(1024, eps=0.0)
(relu): ReLU(inplace=True)
)
)
(layer4): Sequential(
(0): Bottleneck(
(conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): FrozenBatchNorm2d(512, eps=0.0)
(conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
(bn2): FrozenBatchNorm2d(512, eps=0.0)
(conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): FrozenBatchNorm2d(2048, eps=0.0)
(relu): ReLU(inplace=True)
(downsample): Sequential(
(0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
(1): FrozenBatchNorm2d(2048, eps=0.0)
)
)
(1): Bottleneck(
(conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): FrozenBatchNorm2d(512, eps=0.0)
(conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): FrozenBatchNorm2d(512, eps=0.0)
(conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): FrozenBatchNorm2d(2048, eps=0.0)
(relu): ReLU(inplace=True)
)
(2): Bottleneck(
(conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): FrozenBatchNorm2d(512, eps=0.0)
(conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): FrozenBatchNorm2d(512, eps=0.0)
(conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): FrozenBatchNorm2d(2048, eps=0.0)
(relu): ReLU(inplace=True)
)
)
)
(fpn): FeaturePyramidNetwork(
(inner_blocks): ModuleList(
(0): Conv2dNormActivation(
(0): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
)
)
)

```

In [9]: `import json`

```
In [10]: import matplotlib.pyplot as plt  
import matplotlib.patches as patches  
import numpy as np  
from PIL import Image
```

```
In [20]: # Function to load the COCO label map from an annotation file
def load_label_map(ann_file):
    with open(ann_file, 'r') as f:
        coco_data = json.load(f)

    label_map = {category['id']: category['name'] for category in coco_data['categories']}
    return label_map
```

```
In [21]: train_label_map = load_label_map(train_annFile)
print("Train label map:", train_label_map)
```

Train label map: {0: 'background', 1: 'Damage-Windshield', 2: 'Damage-boots', 3: 'Damage-door', 4: 'Damage-front-bumper', 5: 'Damage-headlight', 6: 'Damage-hood', 7: 'Damage-left-fender', 8: 'Damage-rear-bumper', 9: 'Damage-rear-light', 10: 'Damage-right-fender', 11: 'Damage-roof', 12: 'Damage-side-mirror', 13: 'Damage-window-glass', 14: 'flat-tire'}

```
In [22]: # Initialize a persistent color map for labels (only needs to be done once)
def initialize_color_map(label_map):
    # Use a colormap with high contrast (e.g., 'tab20' or 'hsv')
    cmap = plt.cm.get_cmap('hsv', len(label_map)) # Get `len(label_map)` distinct colors
    color_map = {label: cmap(i) for i, label in enumerate(label_map.keys())}
    return color_map

color_map = initialize_color_map(train_label_map) # Initialize once with consistent colors

# Visualization function for predictions using the pre-defined color map
def visualize_predictions(image, outputs, label_map, color_map, threshold=0.5):
    fig, ax = plt.subplots(1, figsize=(12, 9))
    ax.imshow(image)
```

```

pred_boxes = outputs['boxes']
pred_labels = outputs['labels']
pred_scores = outputs['scores']

for i, score in enumerate(pred_scores):
    if score > threshold:
        box = pred_boxes[i].cpu().numpy()
        label = pred_labels[i].cpu().item()
        label_name = label_map.get(label, "Unknown") # Get the class name from the label map
        score = score.cpu().item()

        # Get the pre-defined color for this label
        color = color_map.get(label, 'red') # Default to red if label is unknown

        # Draw bounding box with the specific color
        rect = patches.Rectangle(
            (box[0], box[1]), box[2] - box[0], box[3] - box[1],
            linewidth=2, edgecolor=color, facecolor='none'
        )
        ax.add_patch(rect)

        # Display label name and score with the same color
        ax.text(
            box[0], box[1], f"{label_name}: {score:.2f}",
            color='black', fontsize=10, bbox=dict(facecolor=color, alpha=0.5)
        )

plt.show()

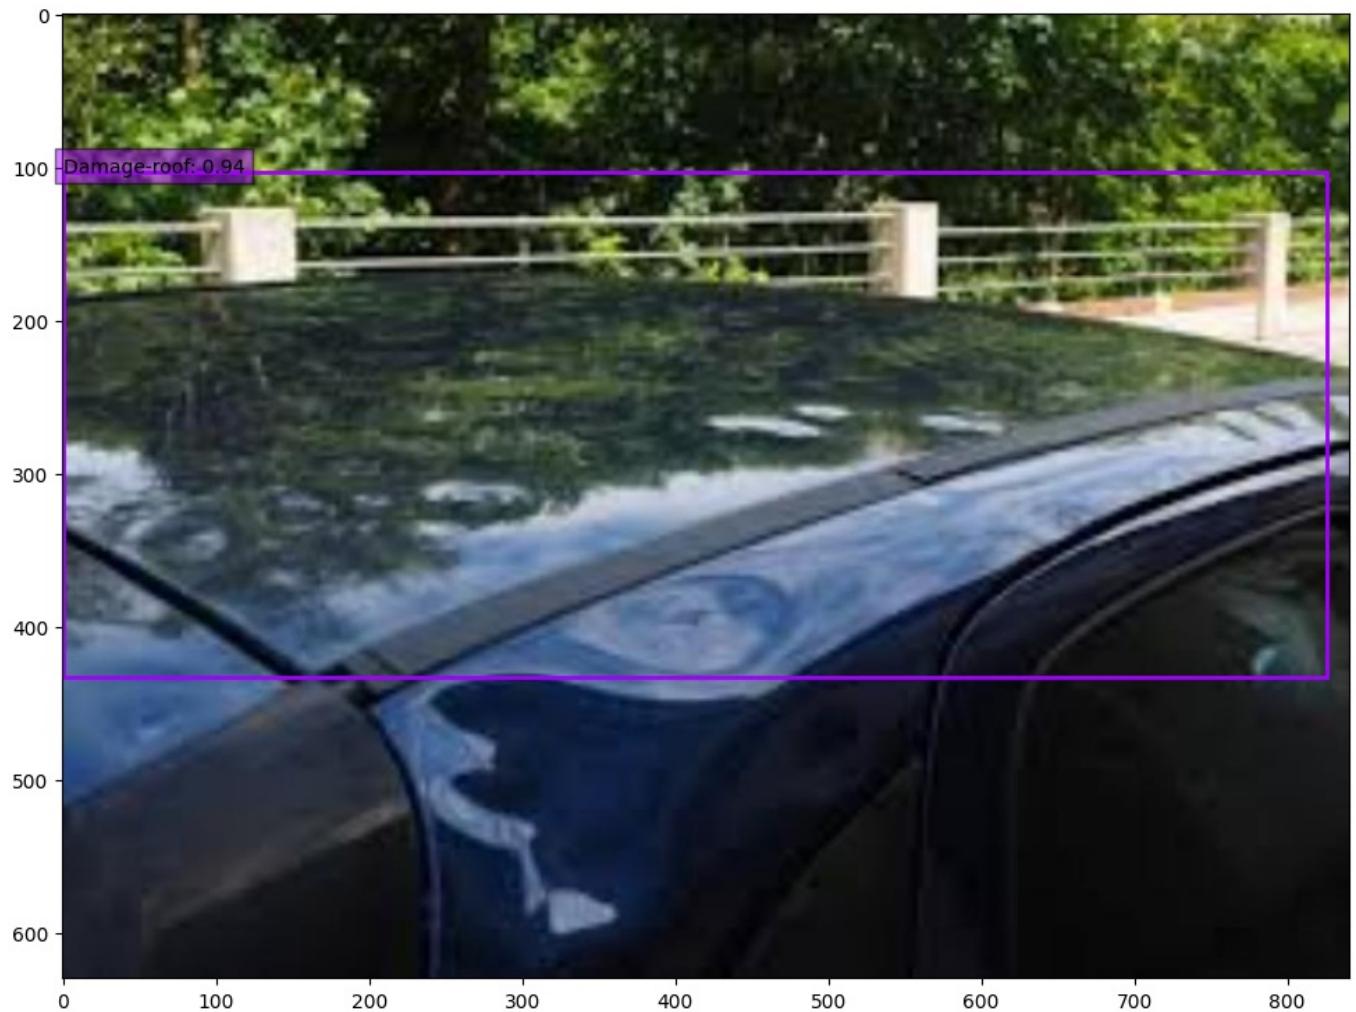
```

```
C:\Users\harsh\AppData\Local\Temp\ipykernel_1872\2259875598.py:4: MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib 3.7 and will be removed in 3.11. Use ``matplotlib.colormaps[name]`` or ``matplotlib.colormaps.get_cmap()`` or ``pyplot.get_cmap()`` instead.
cmap = plt.cm.get_cmap('hsv', len(label_map)) # Get `len(label_map)` distinct colors
```

```
In [23]: # Run inference and visualize predictions
def run_inference(model, data_loader, train_label_map, device):
    model.eval()
    with torch.no_grad():
        for images, _ in data_loader:
            images = [img.to(device) for img in images]
            outputs = model(images)

            for i, output in enumerate(outputs):
                image = images[i].cpu().permute(1, 2, 0).numpy()
                image = np.uint8(image * 255)
                image_pil = Image.fromarray(image)
                visualize_predictions(image_pil, output, train_label_map, color_map) # Pass the label map

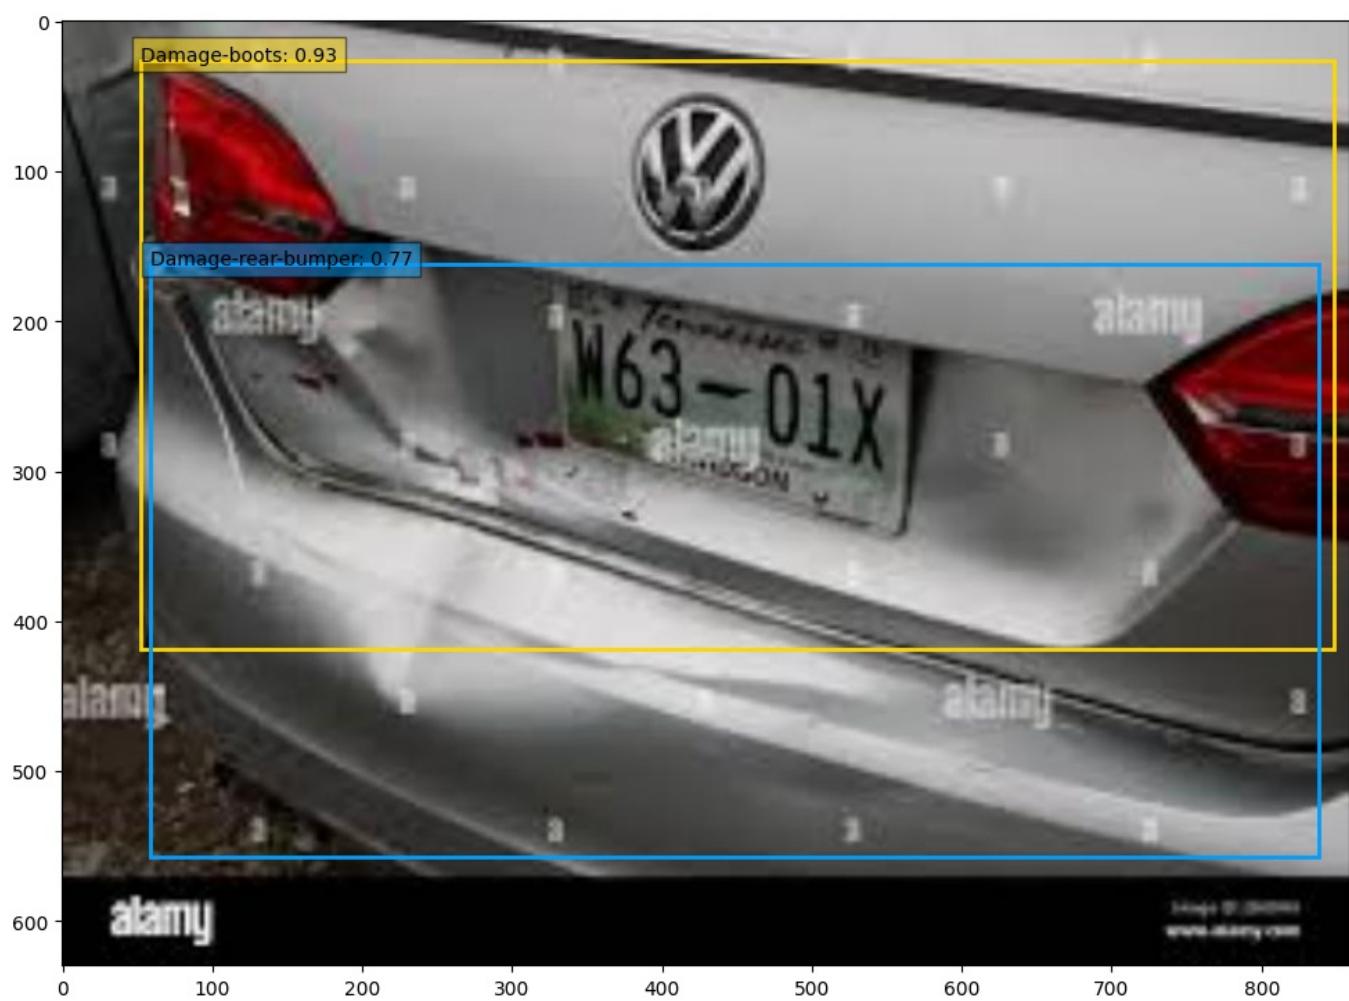
# Correct function call with train_label_map
run_inference(model, test_loader, train_label_map, device)
```

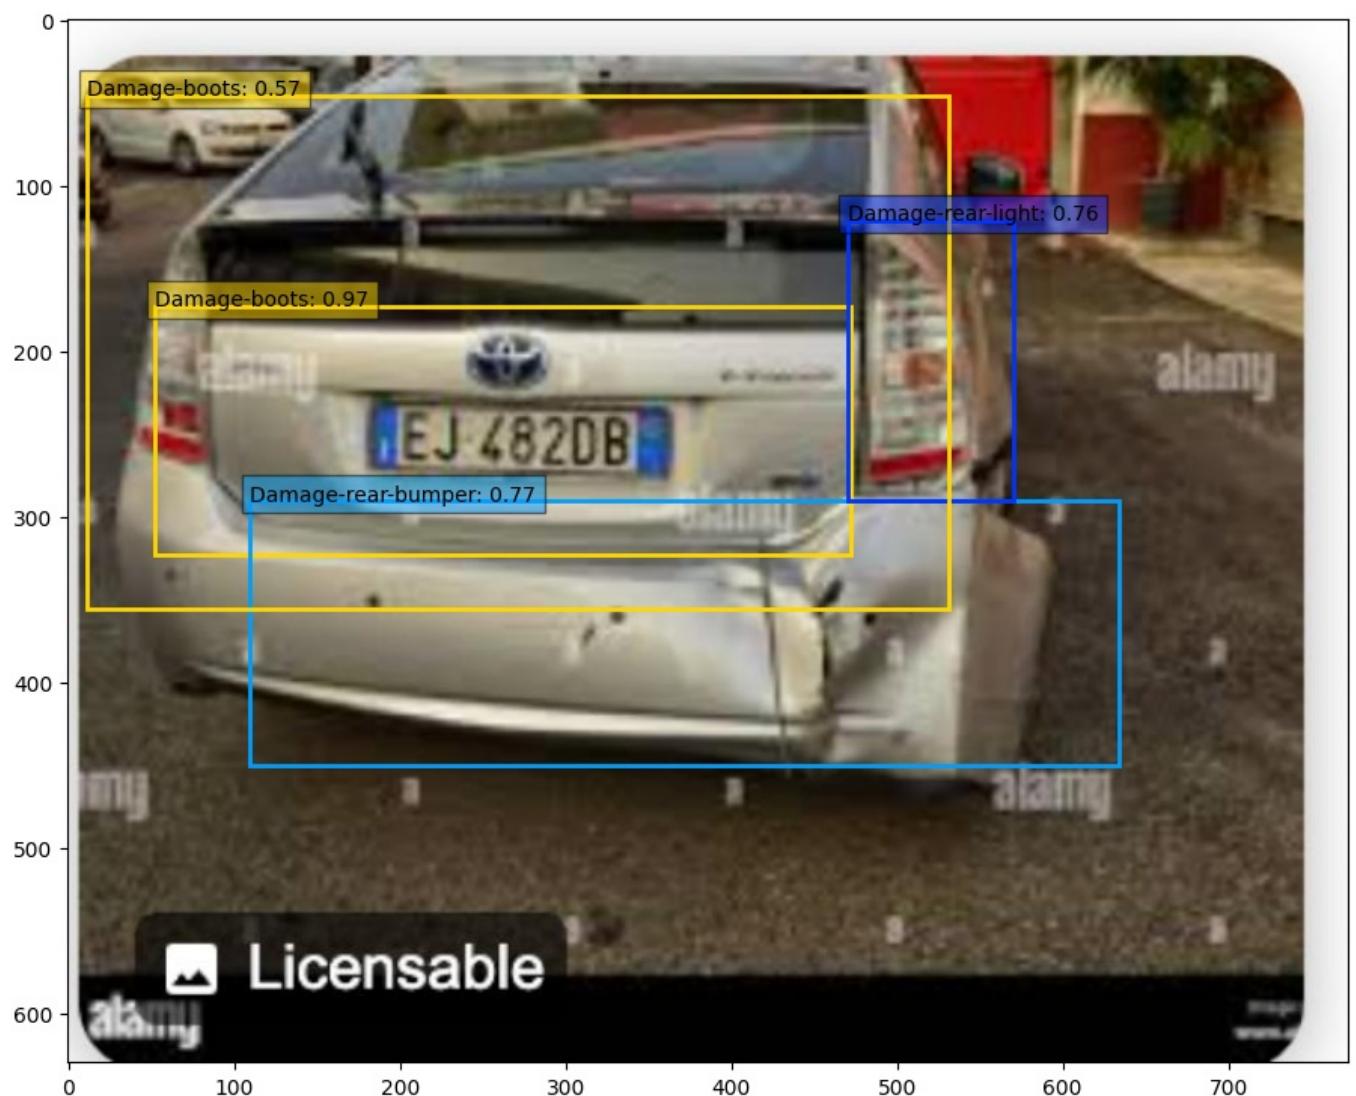


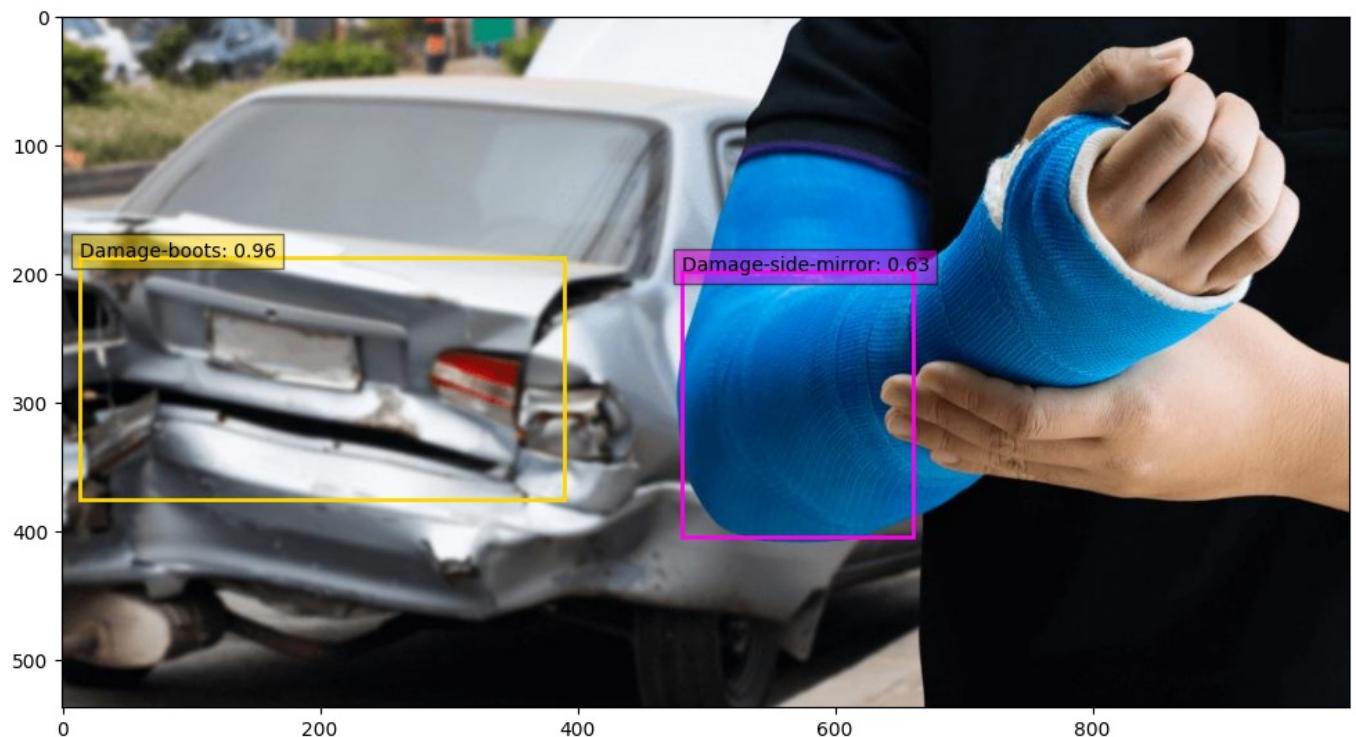
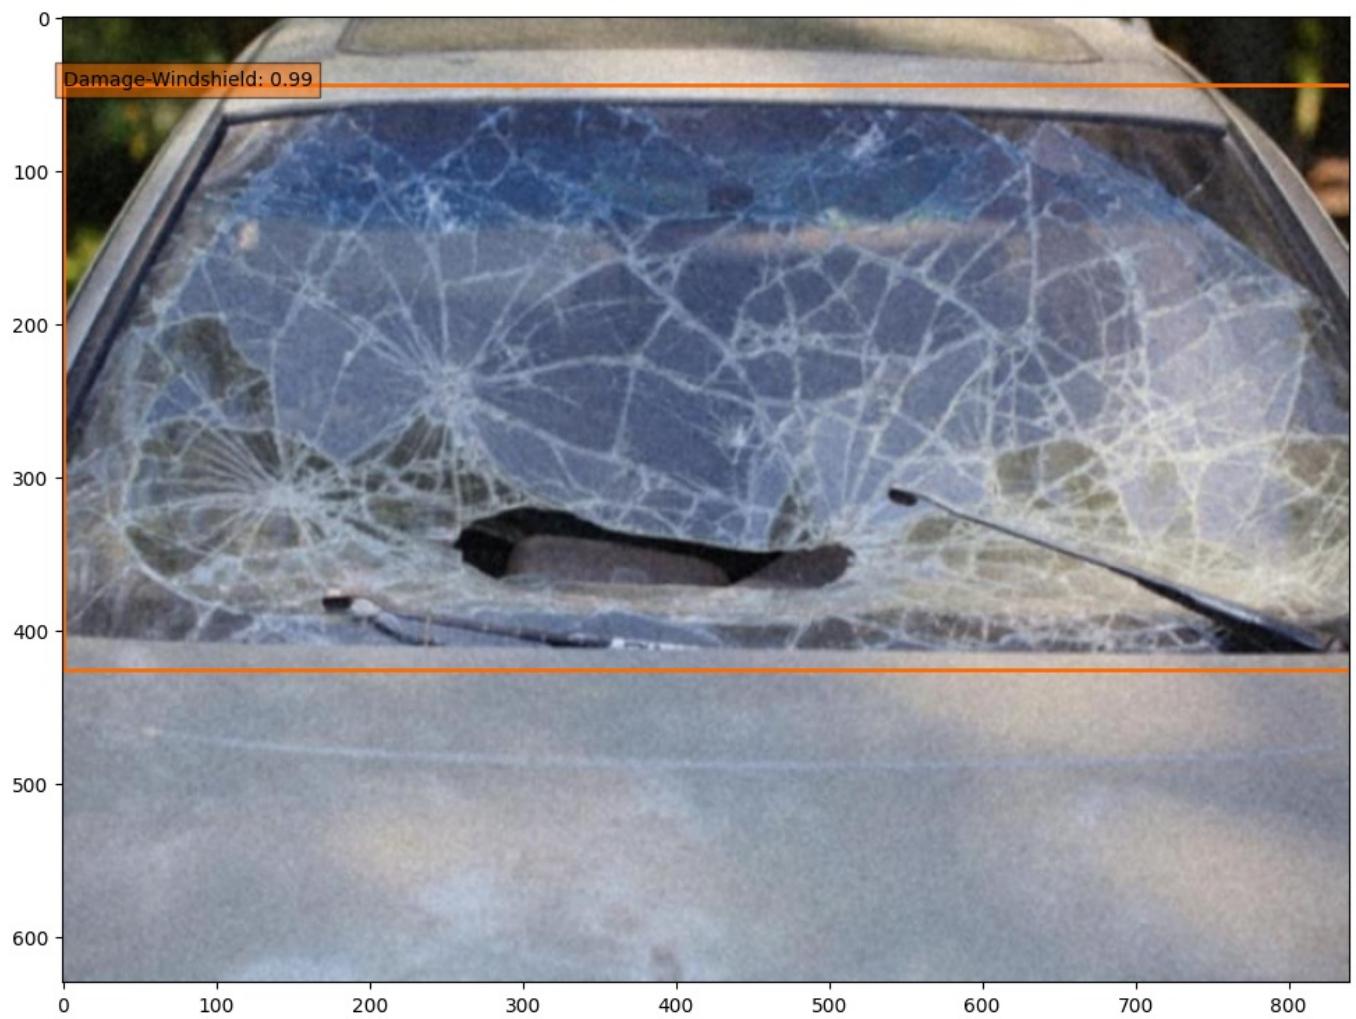


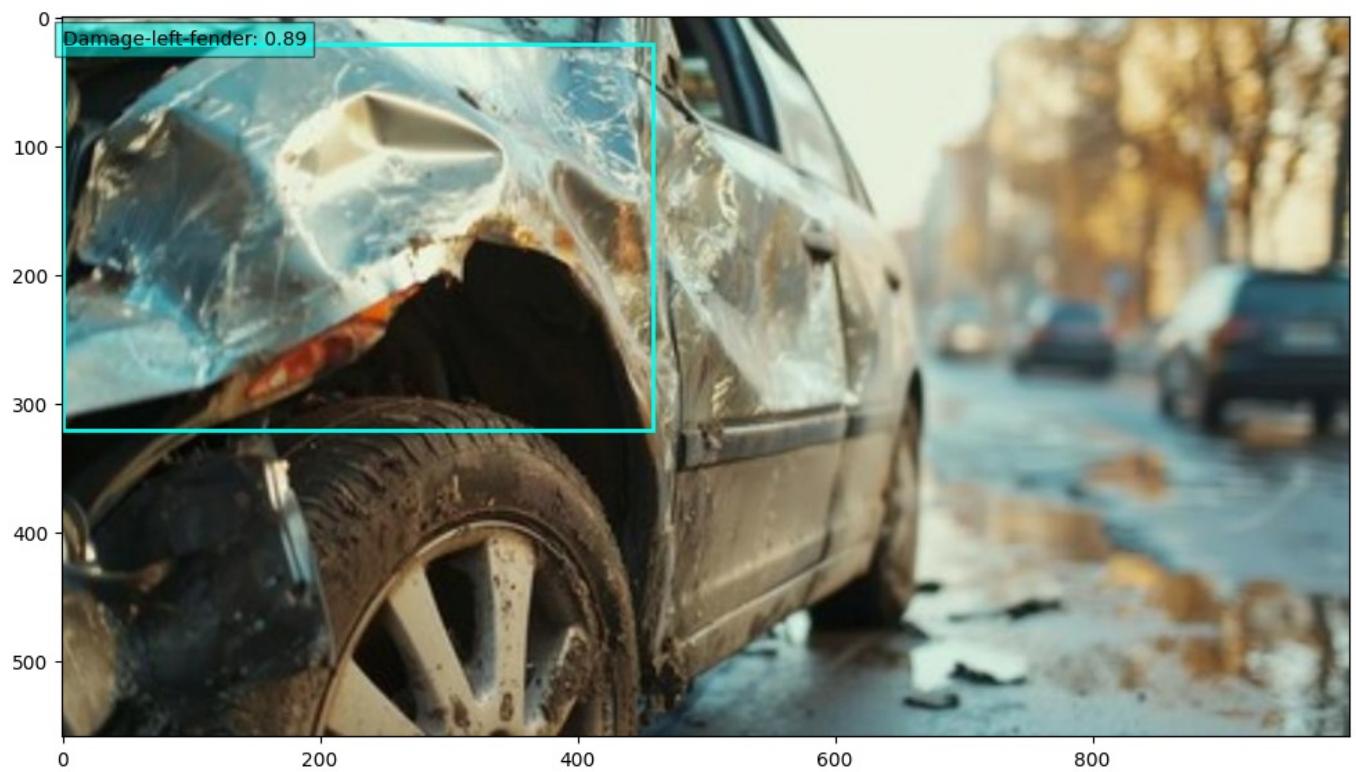
0 Damage-Windshield: 1.00

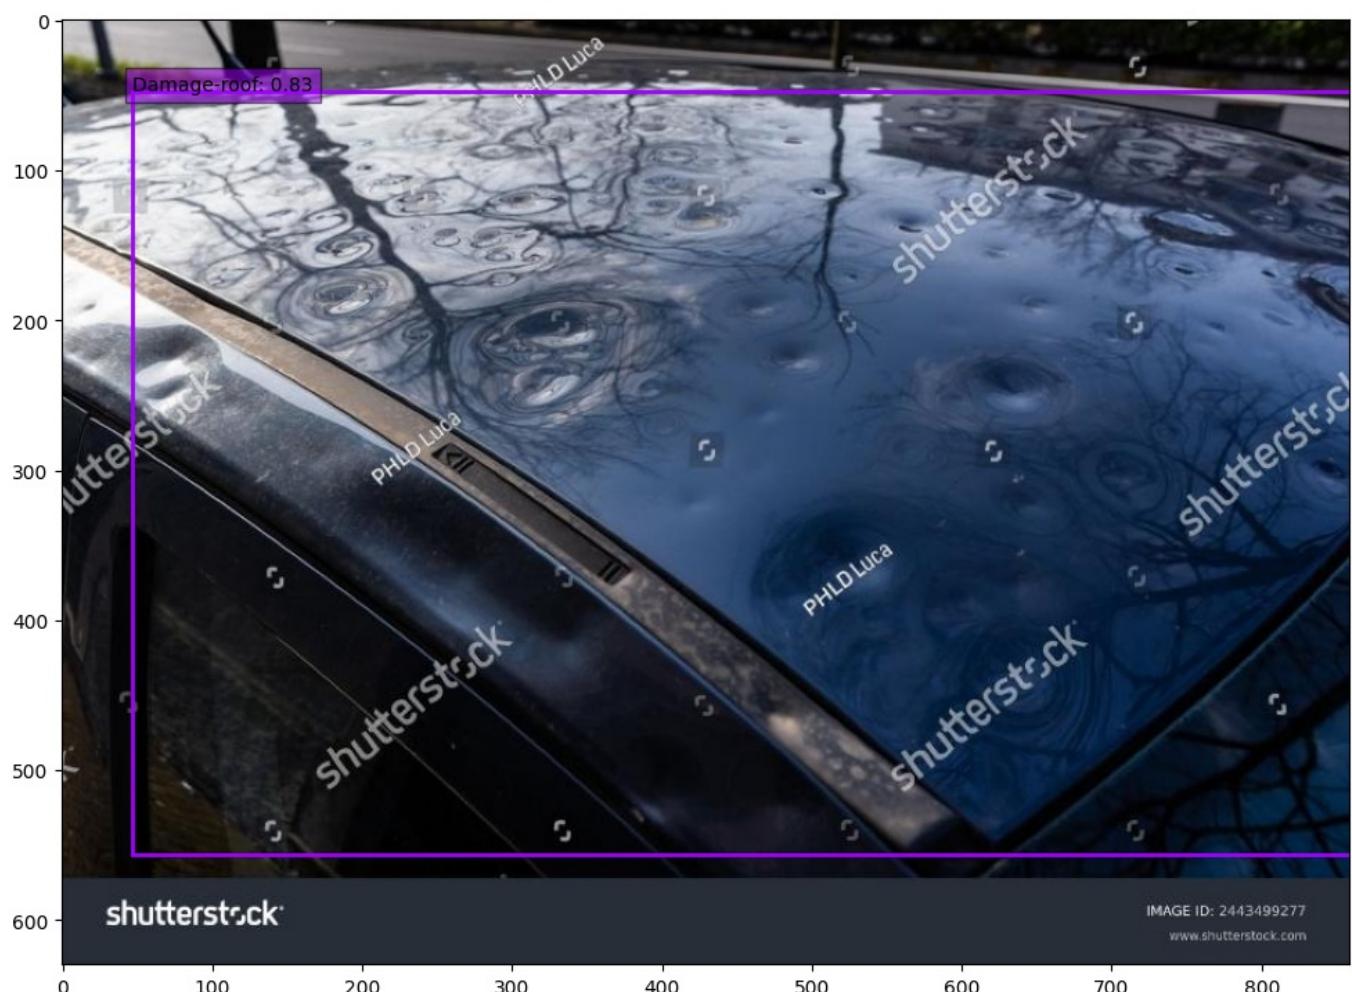
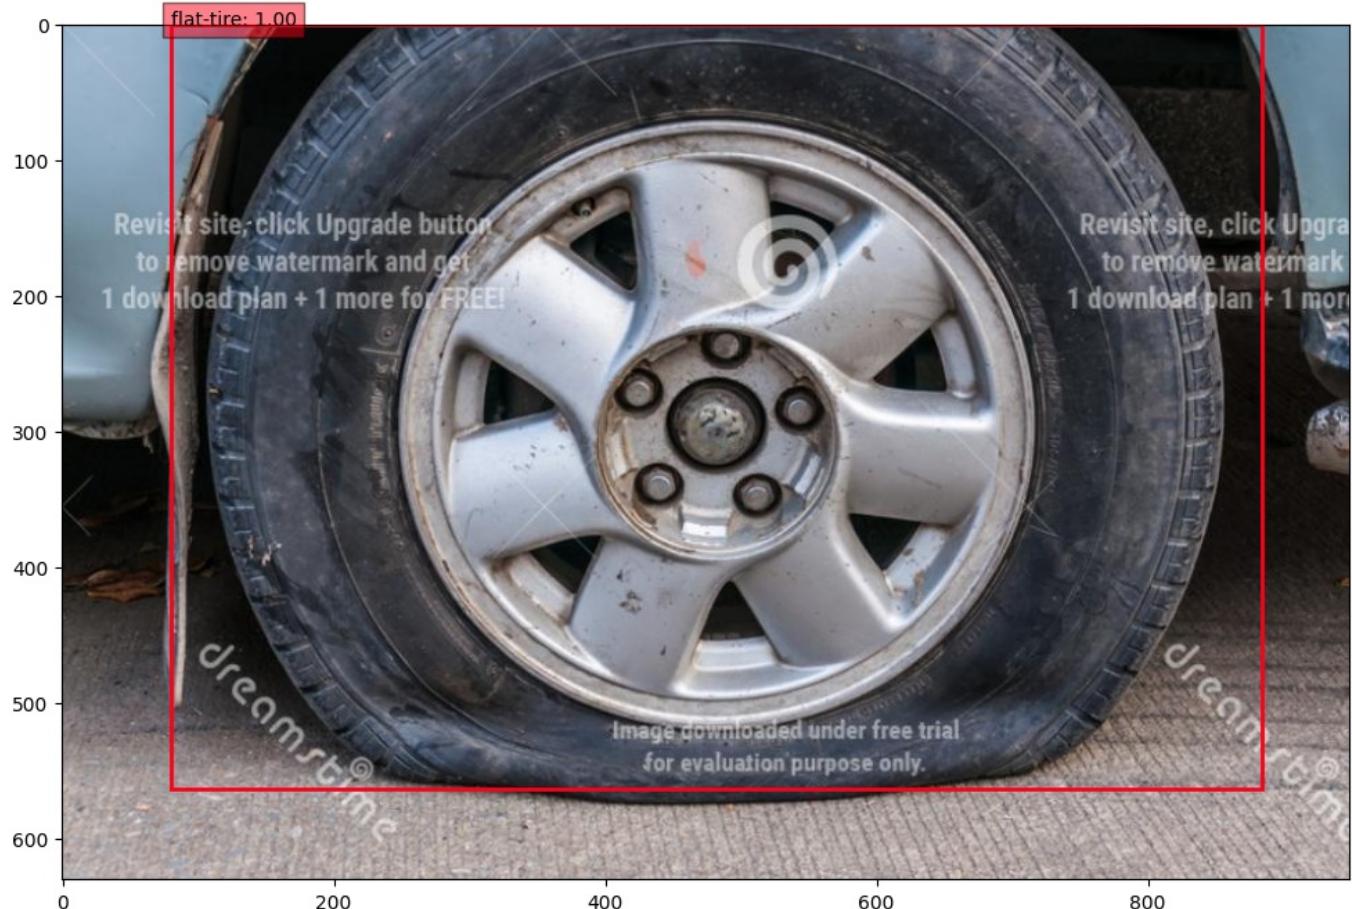


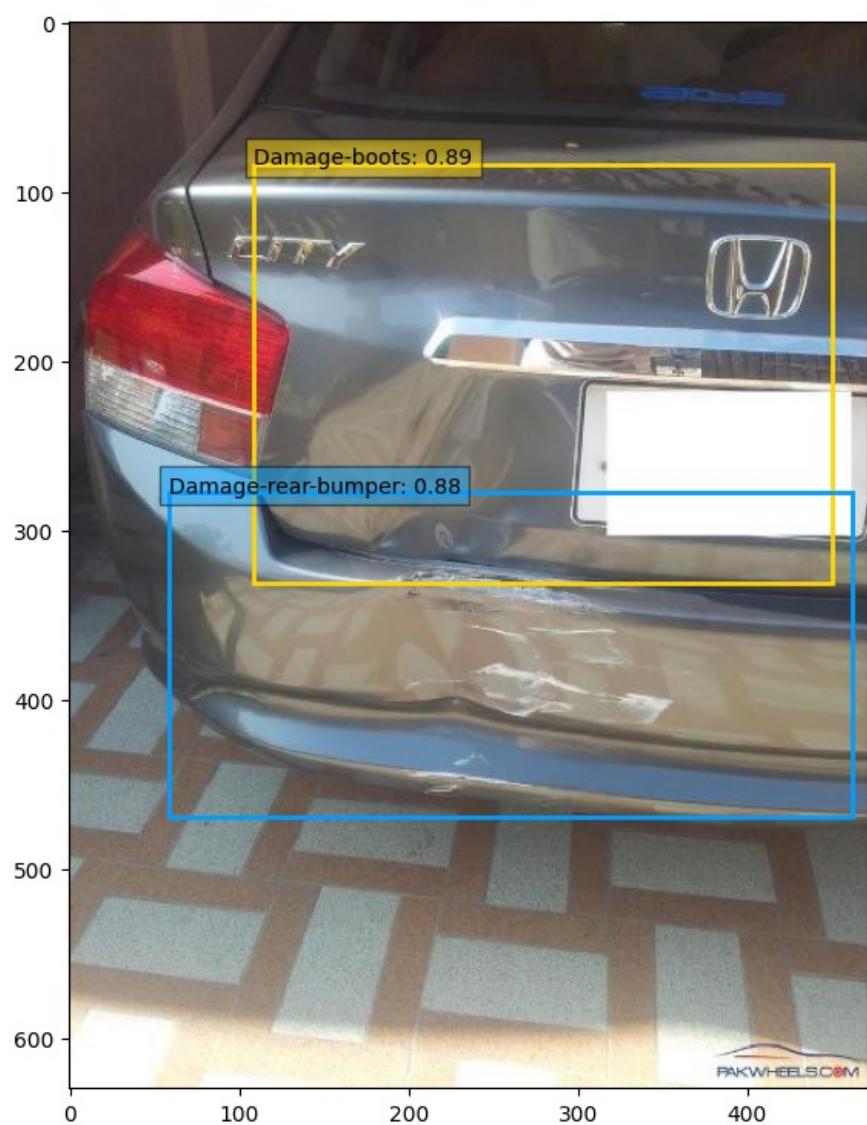




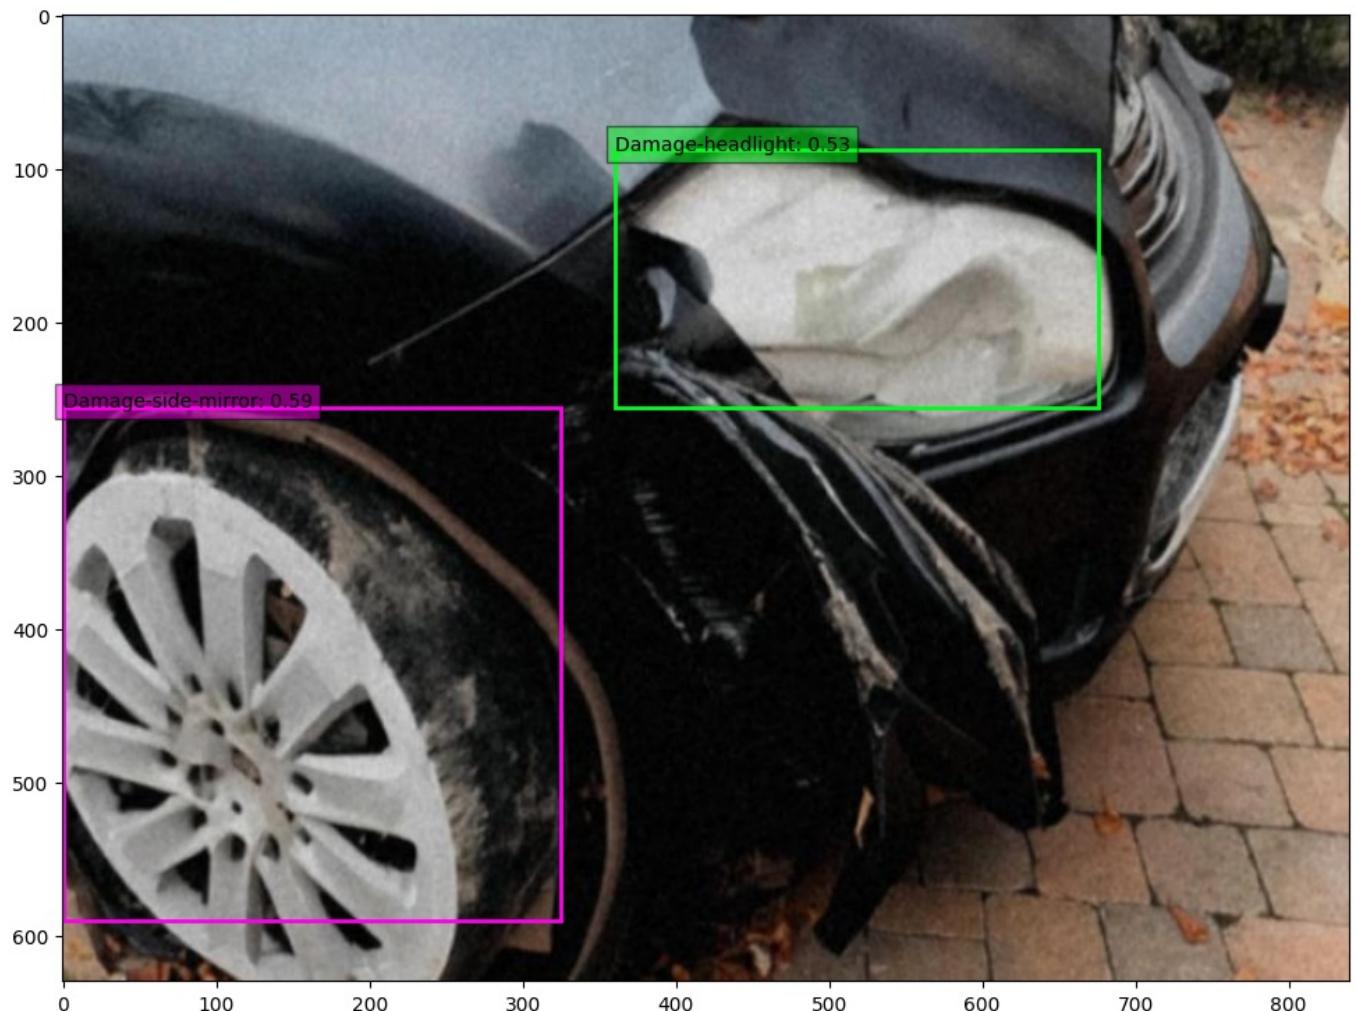


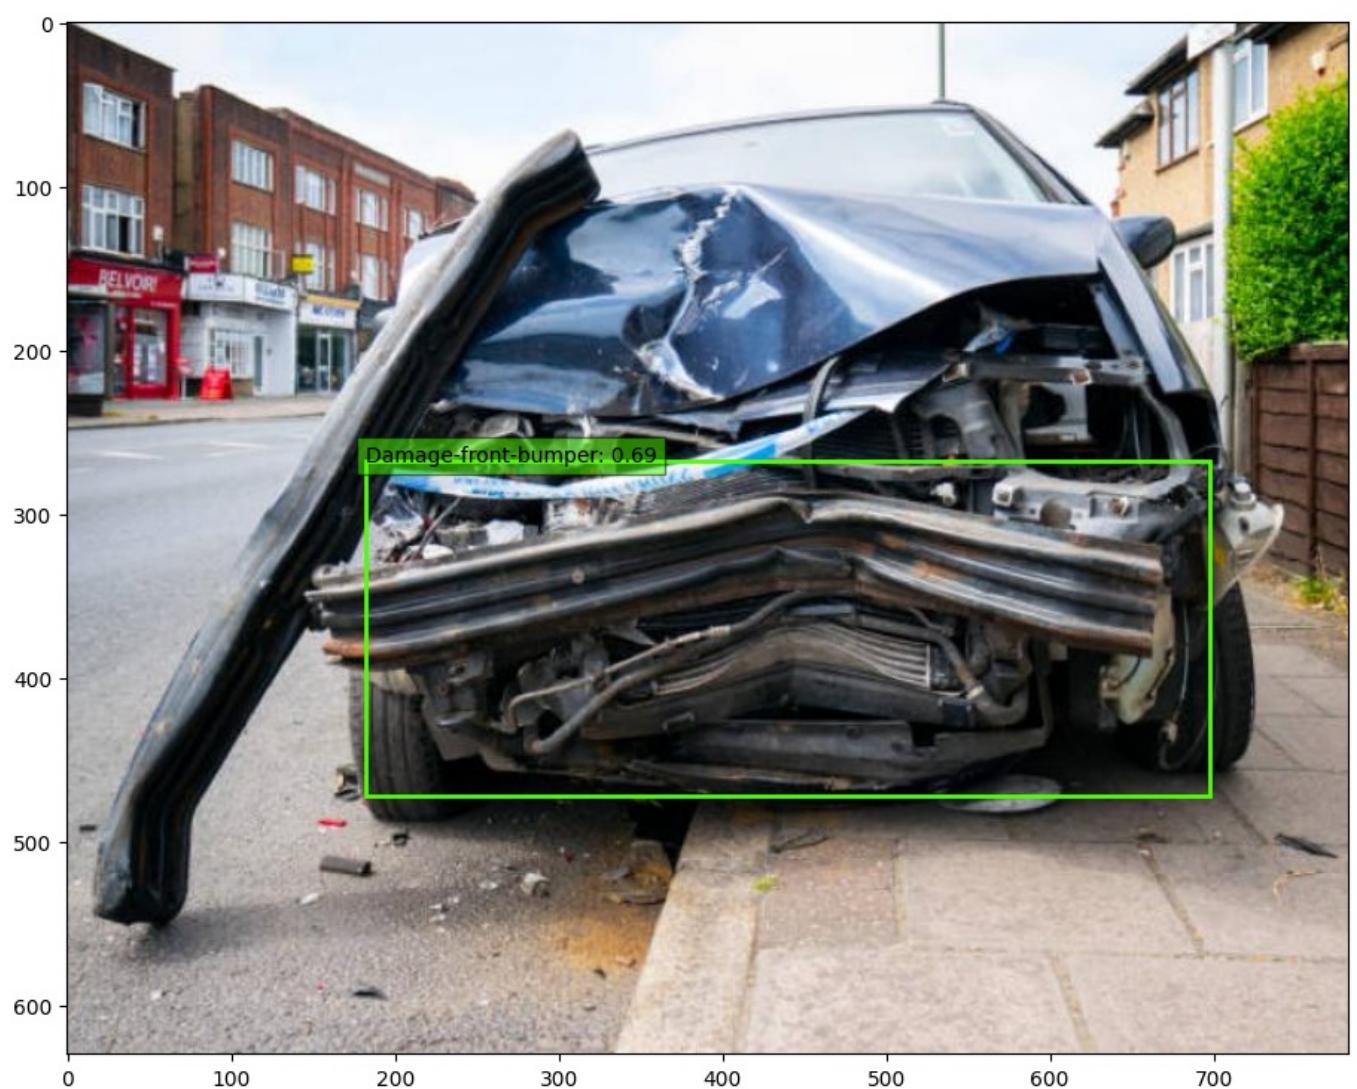
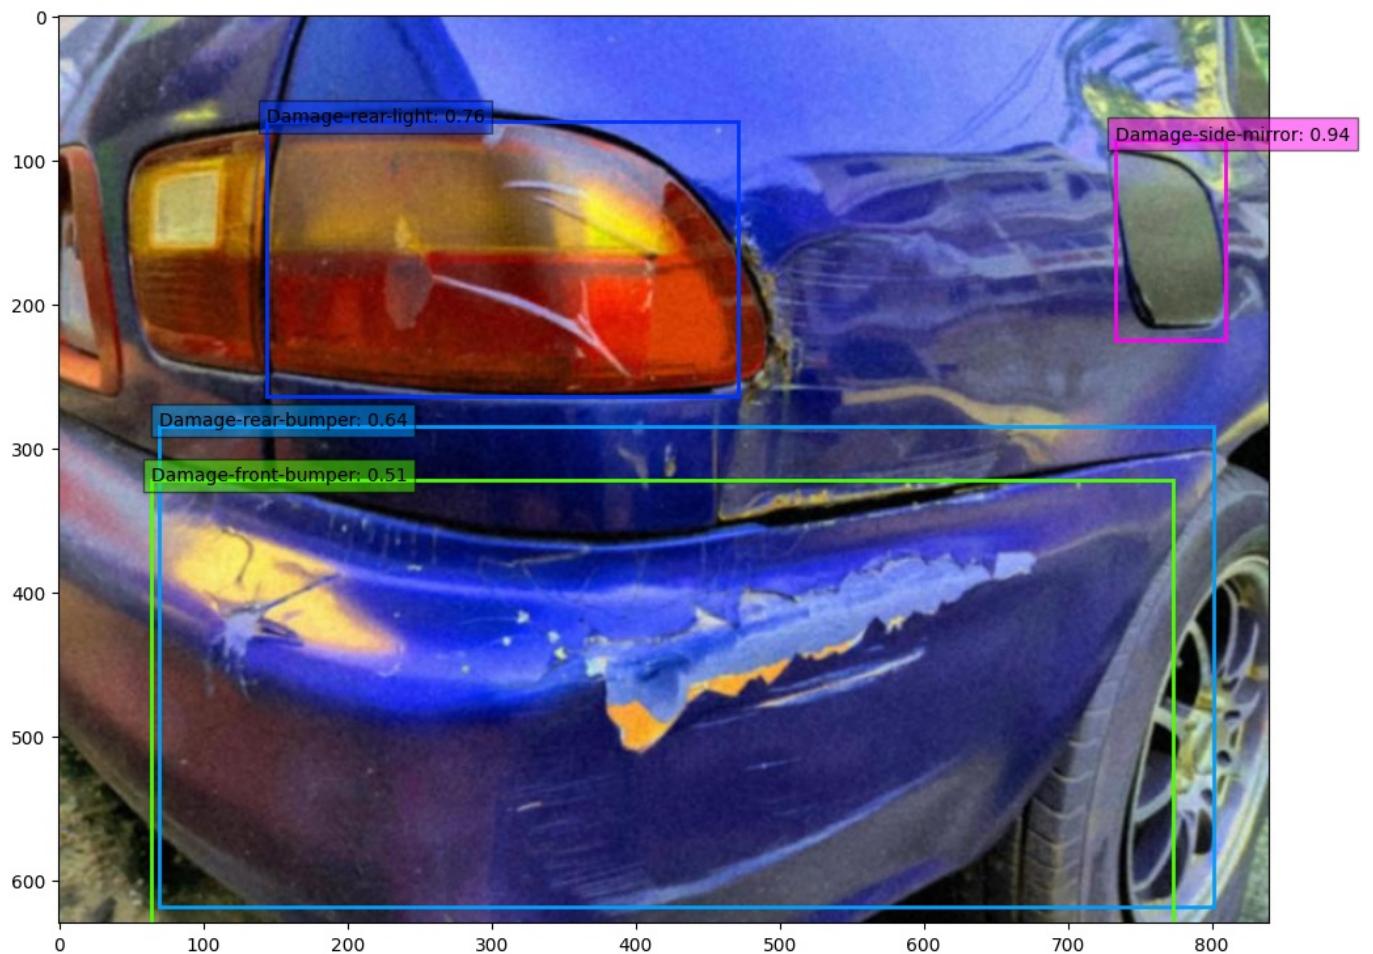


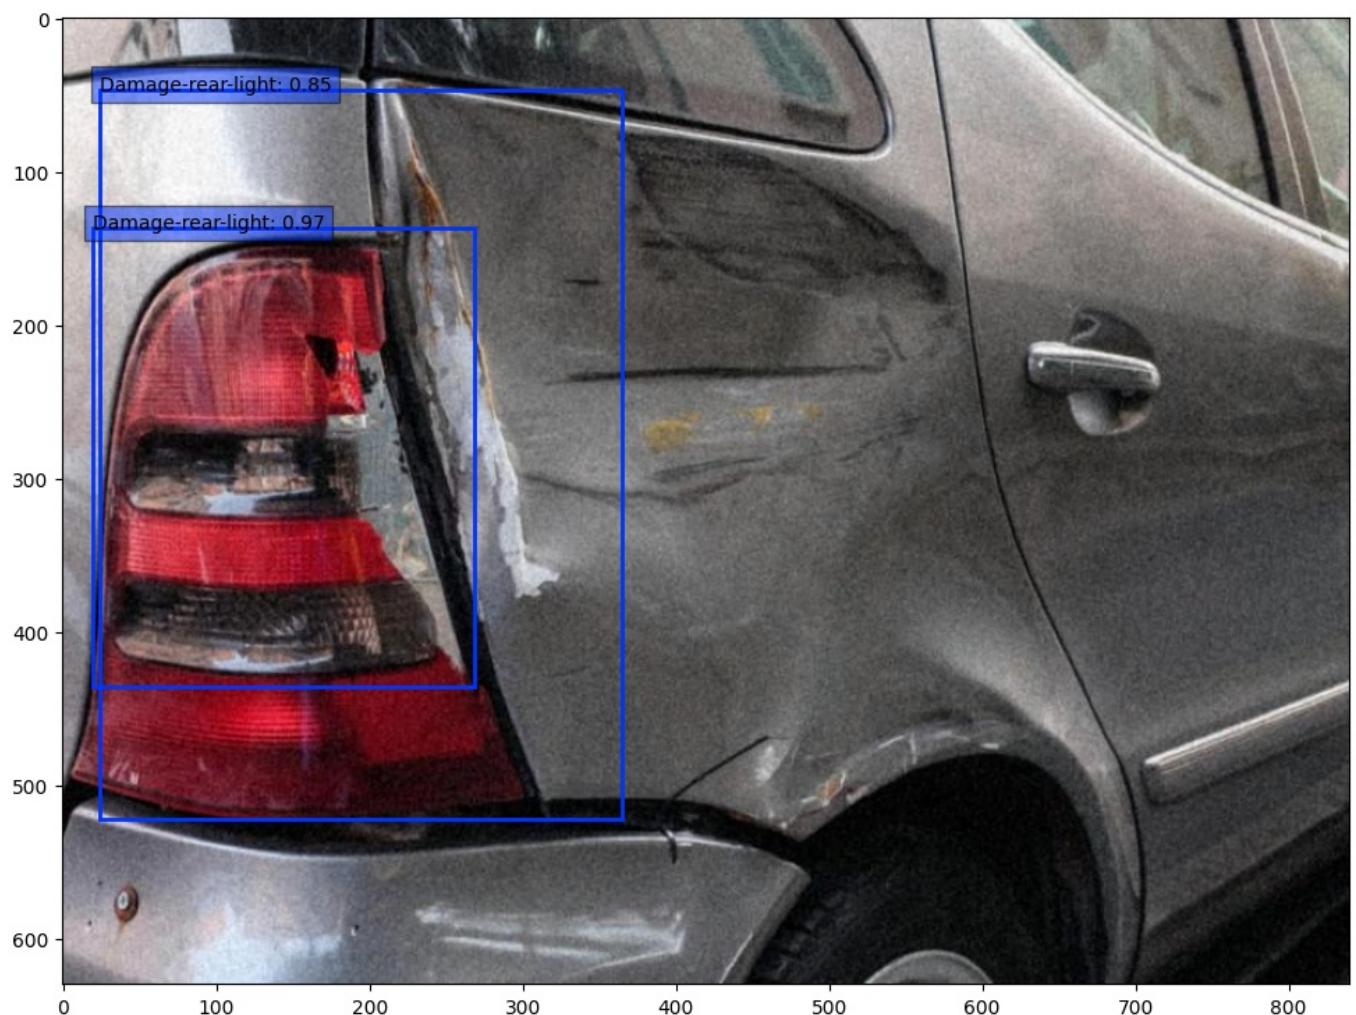
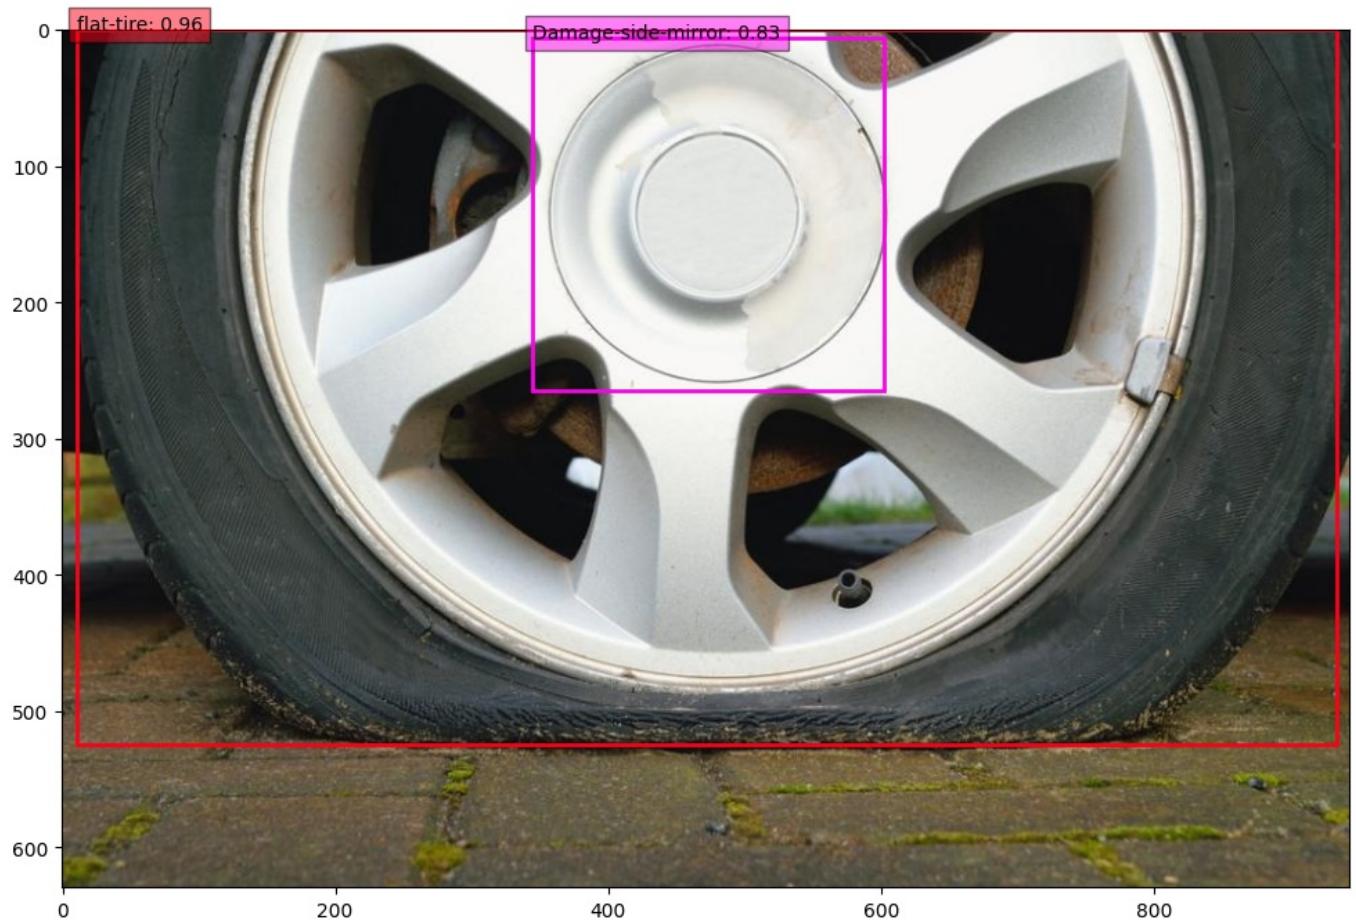


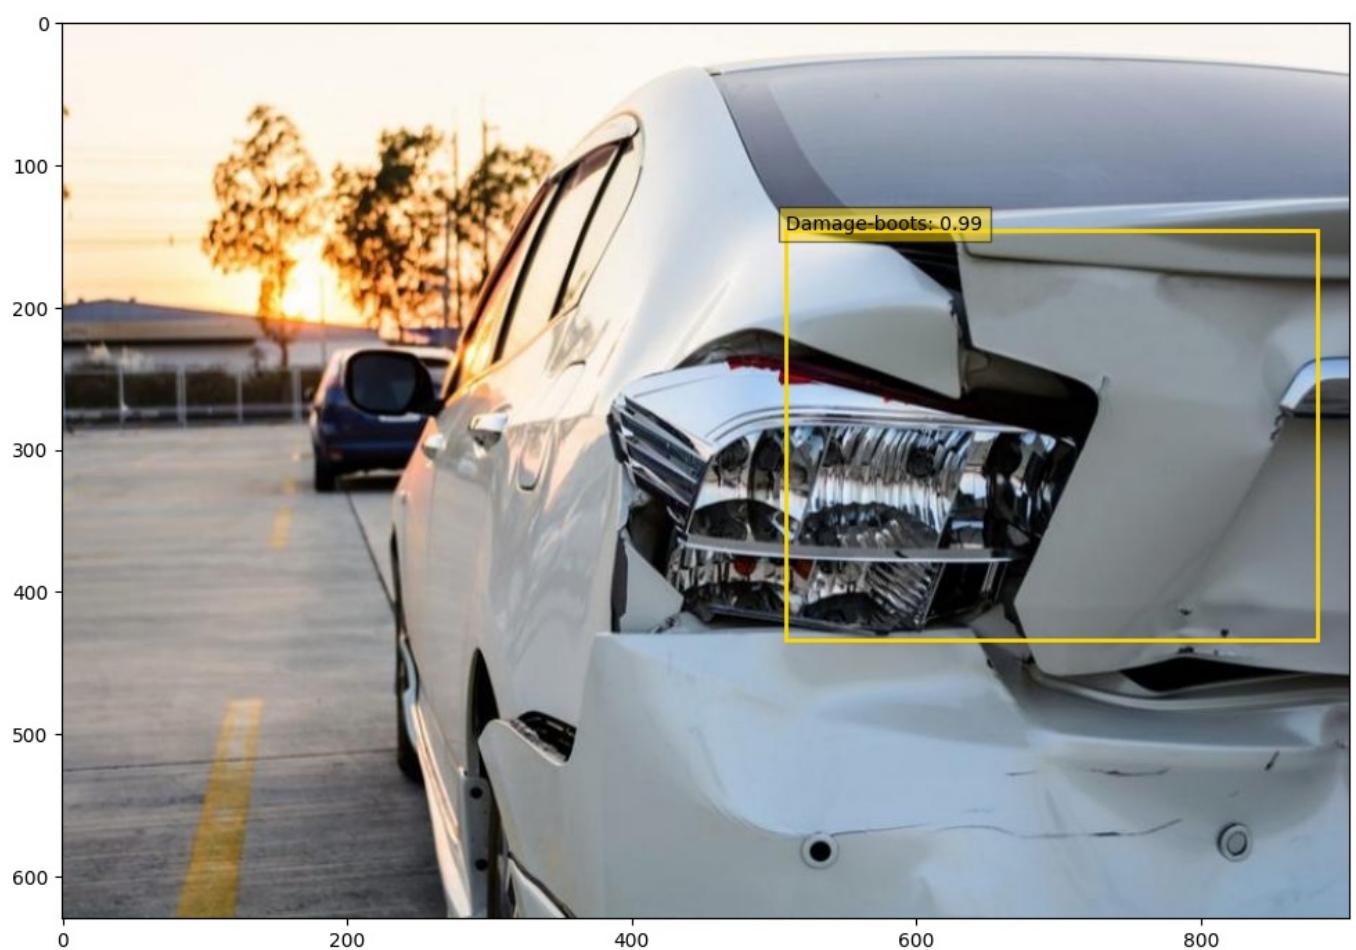






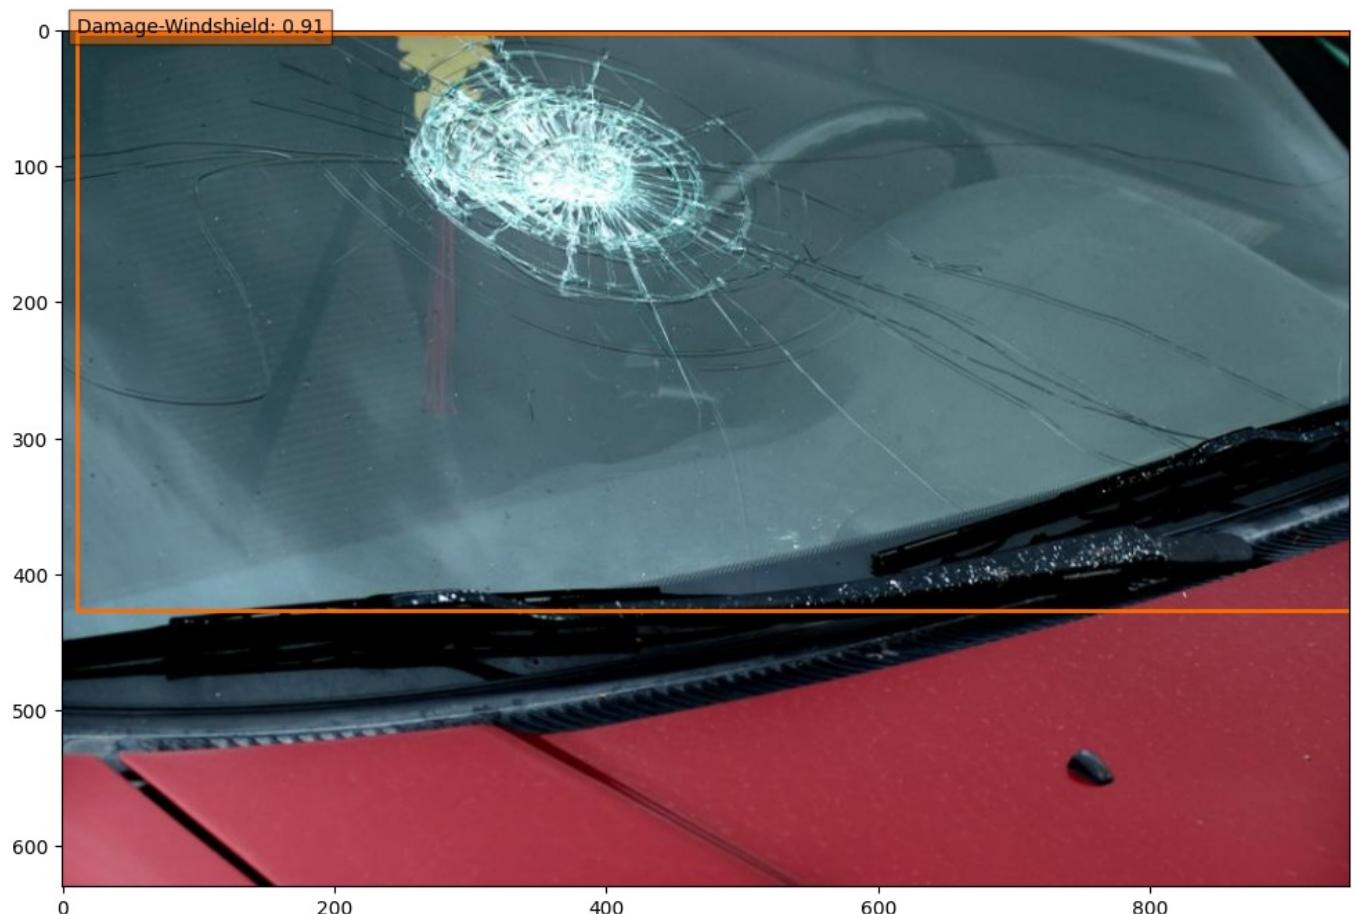
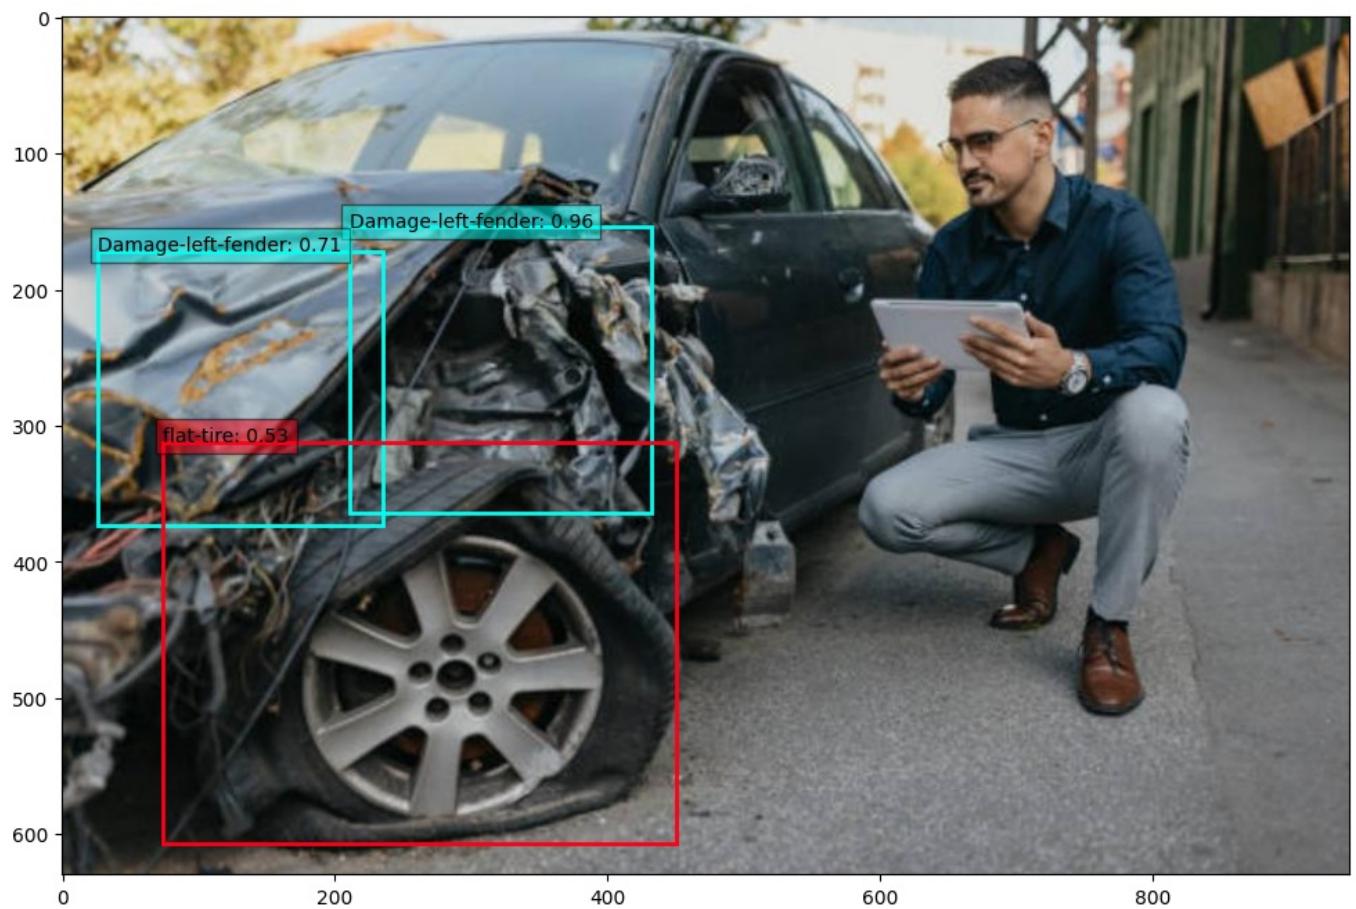




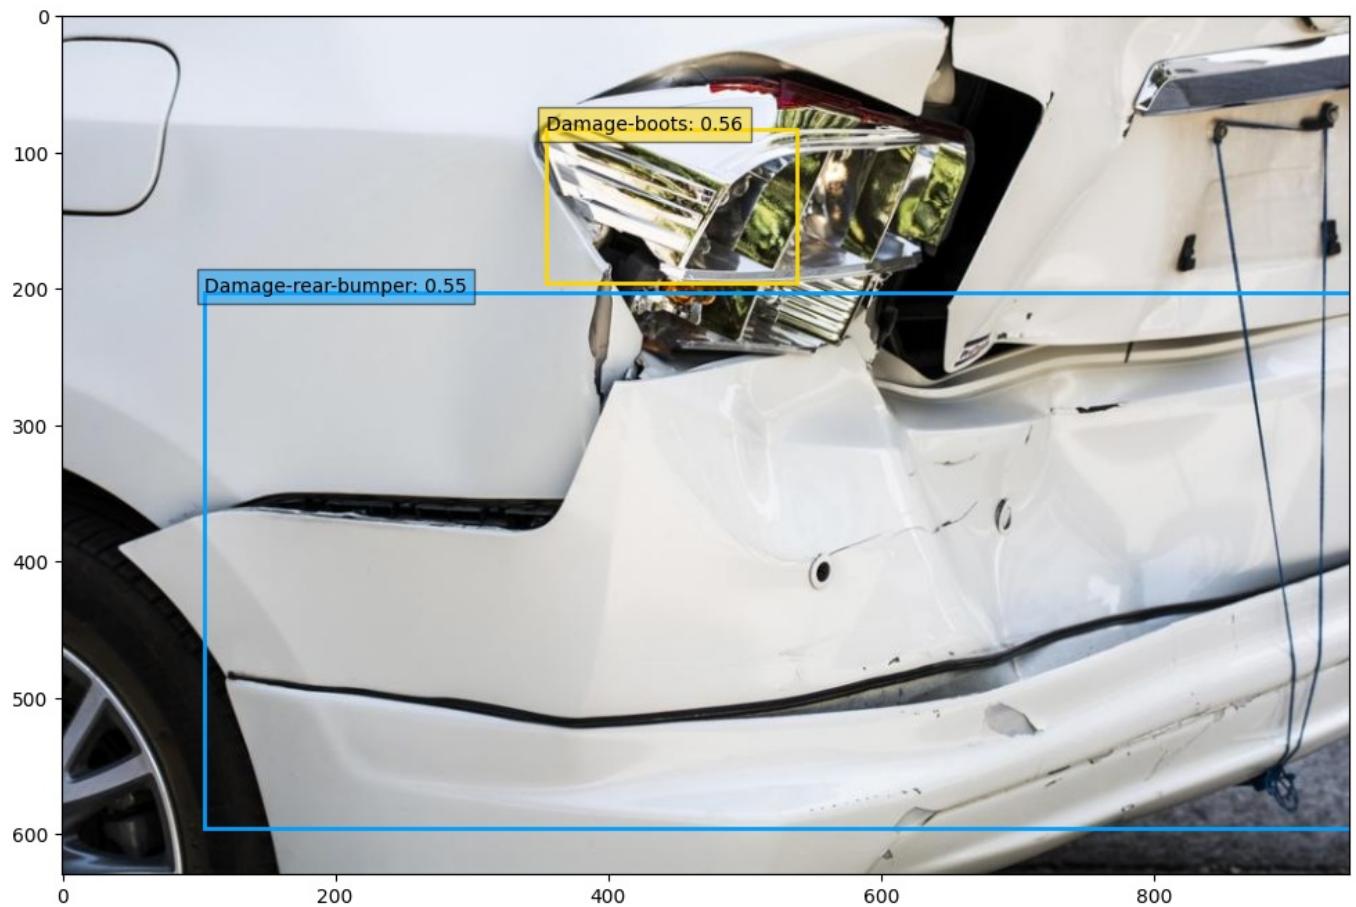


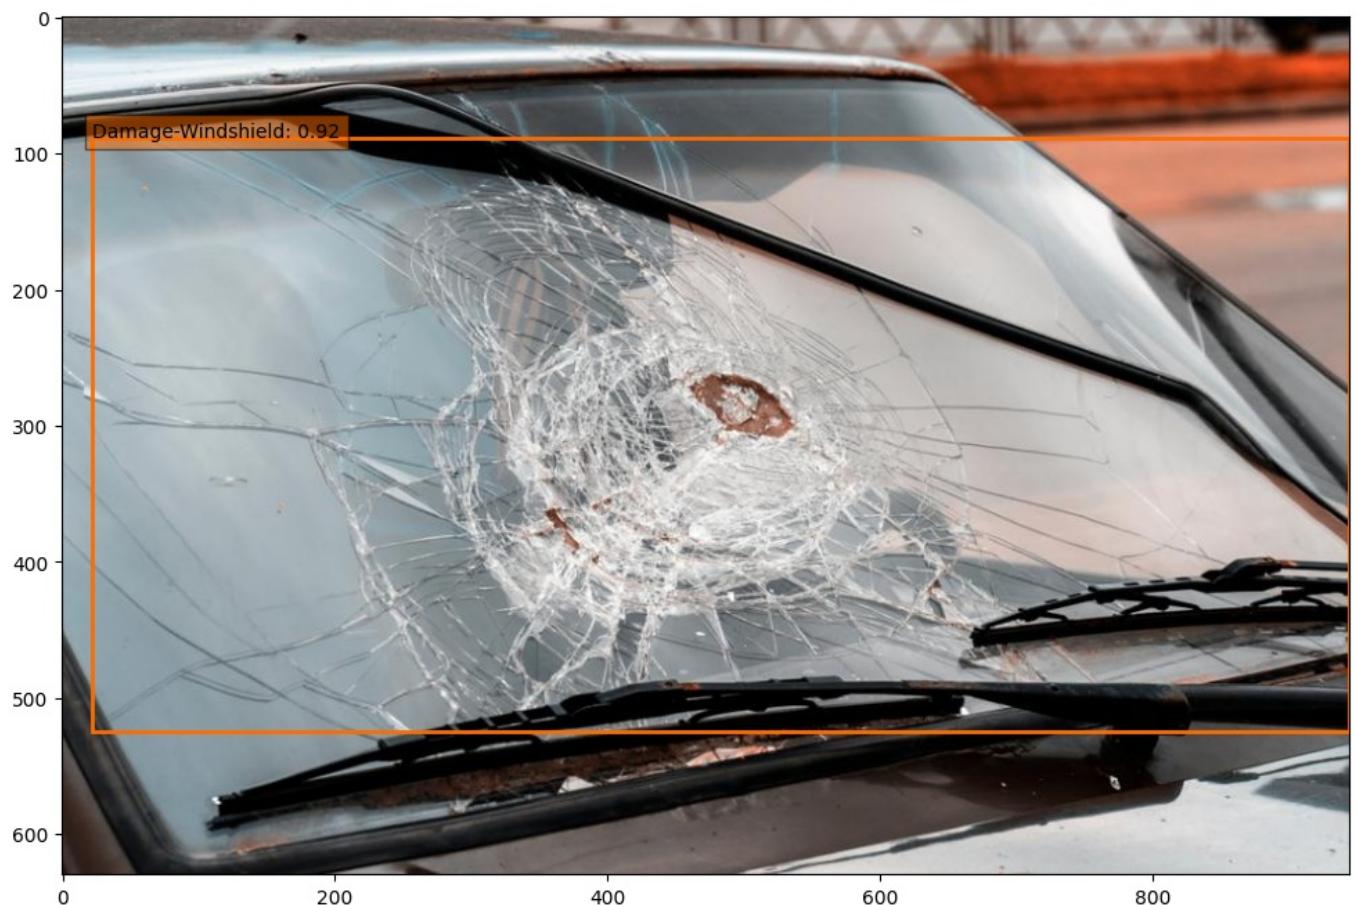


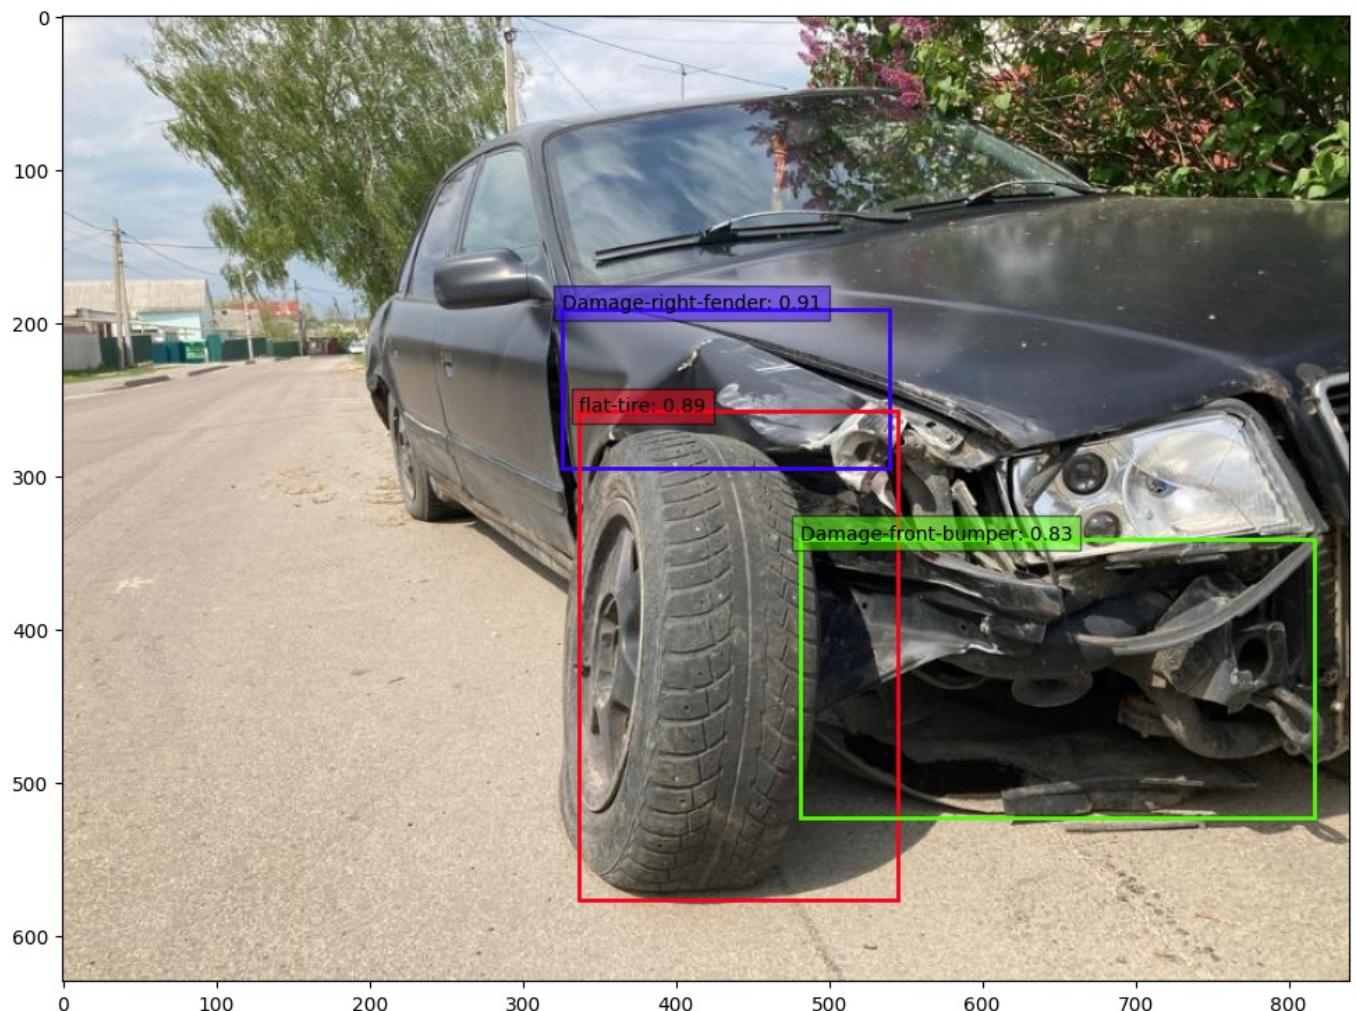
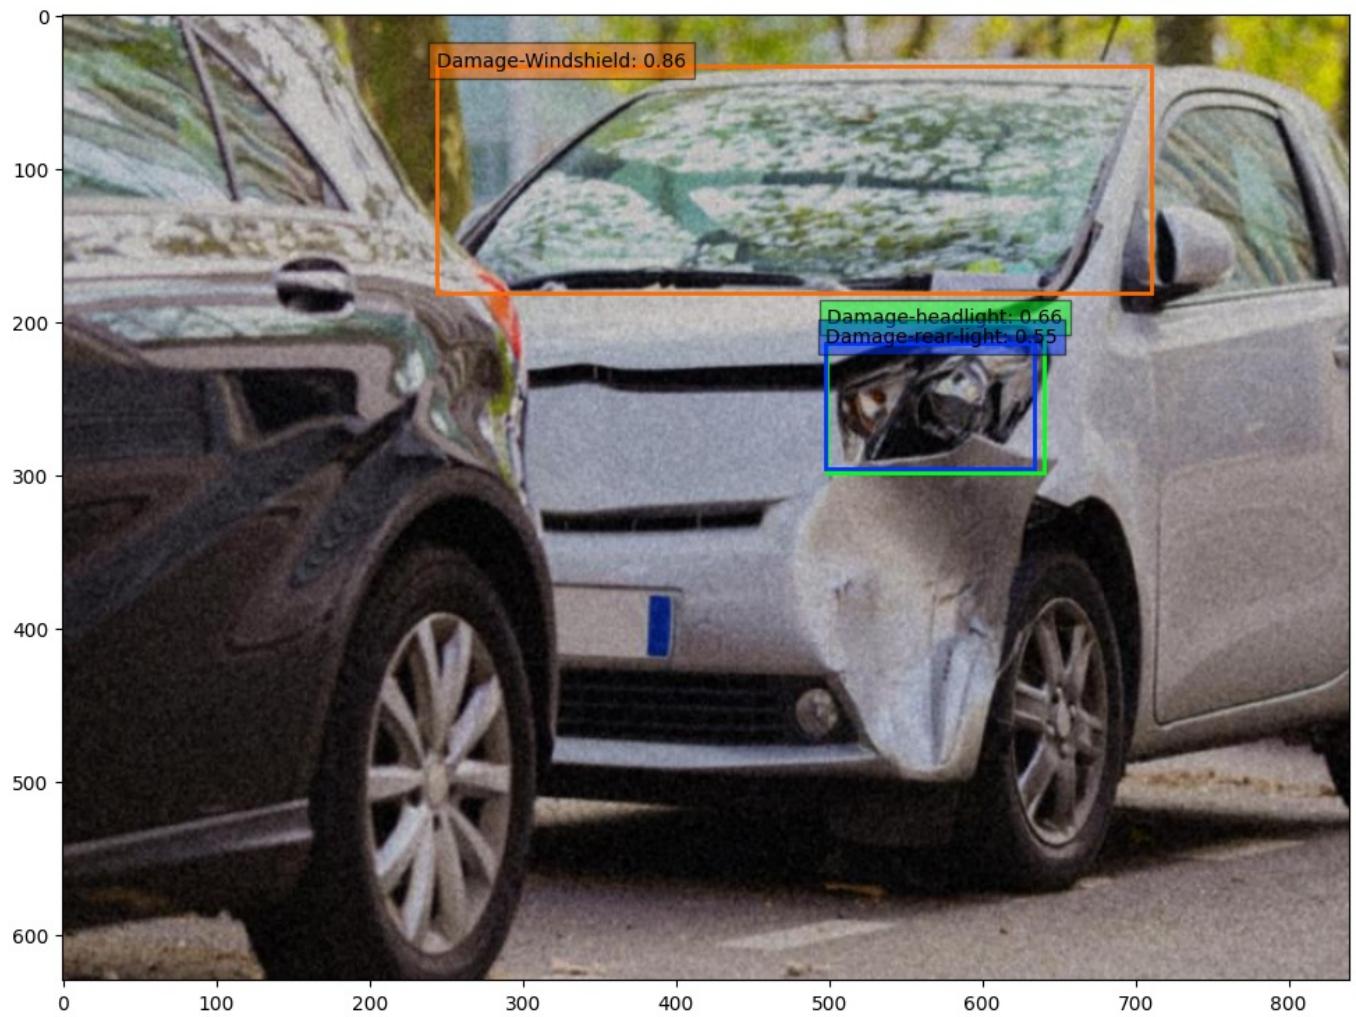


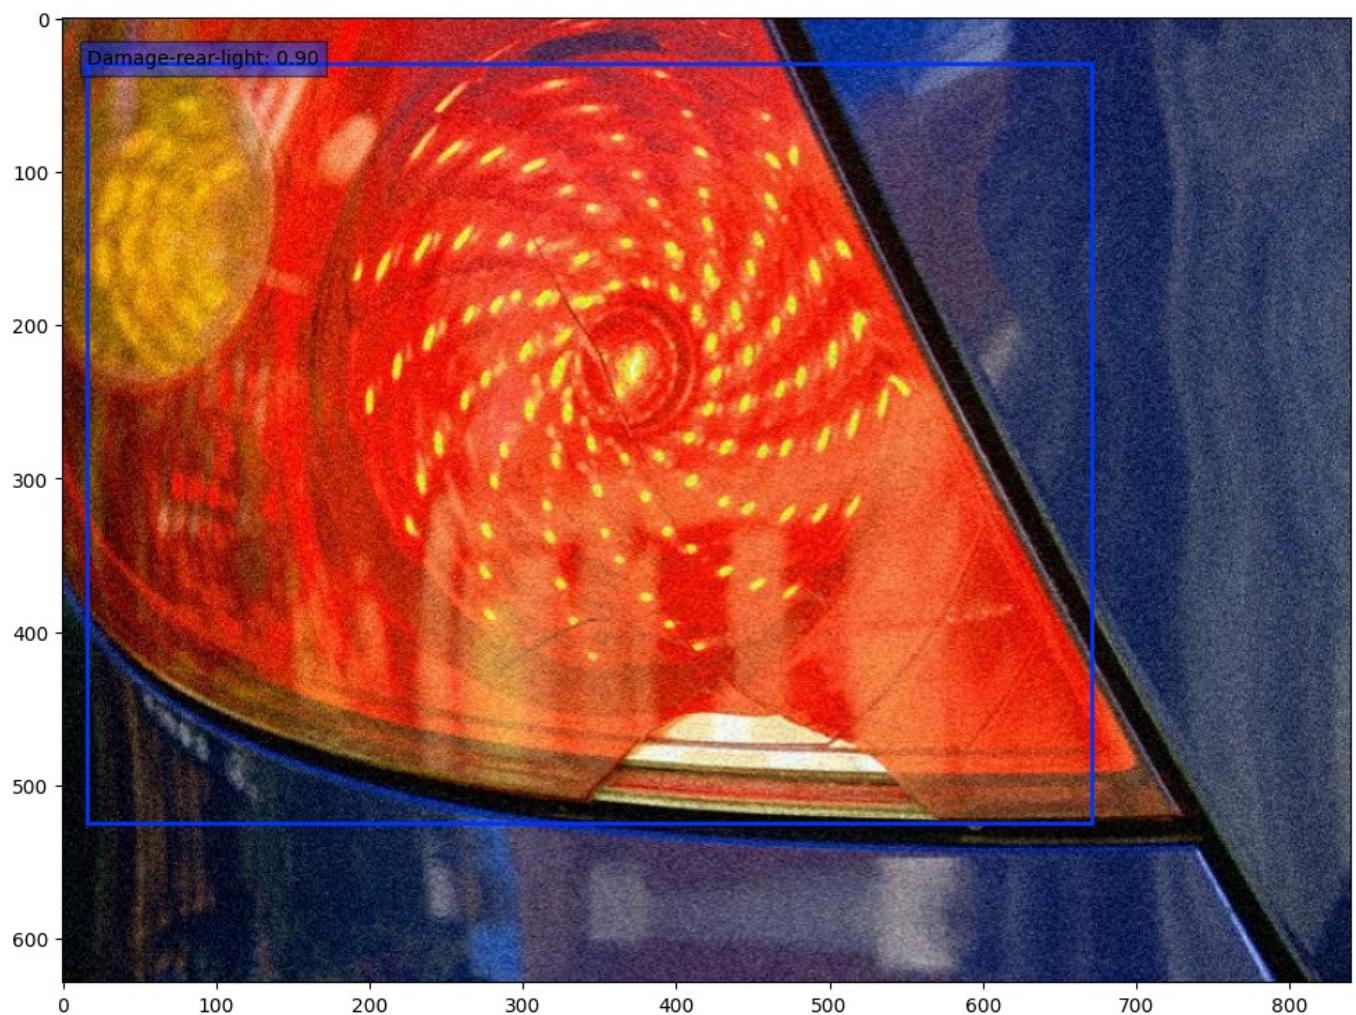


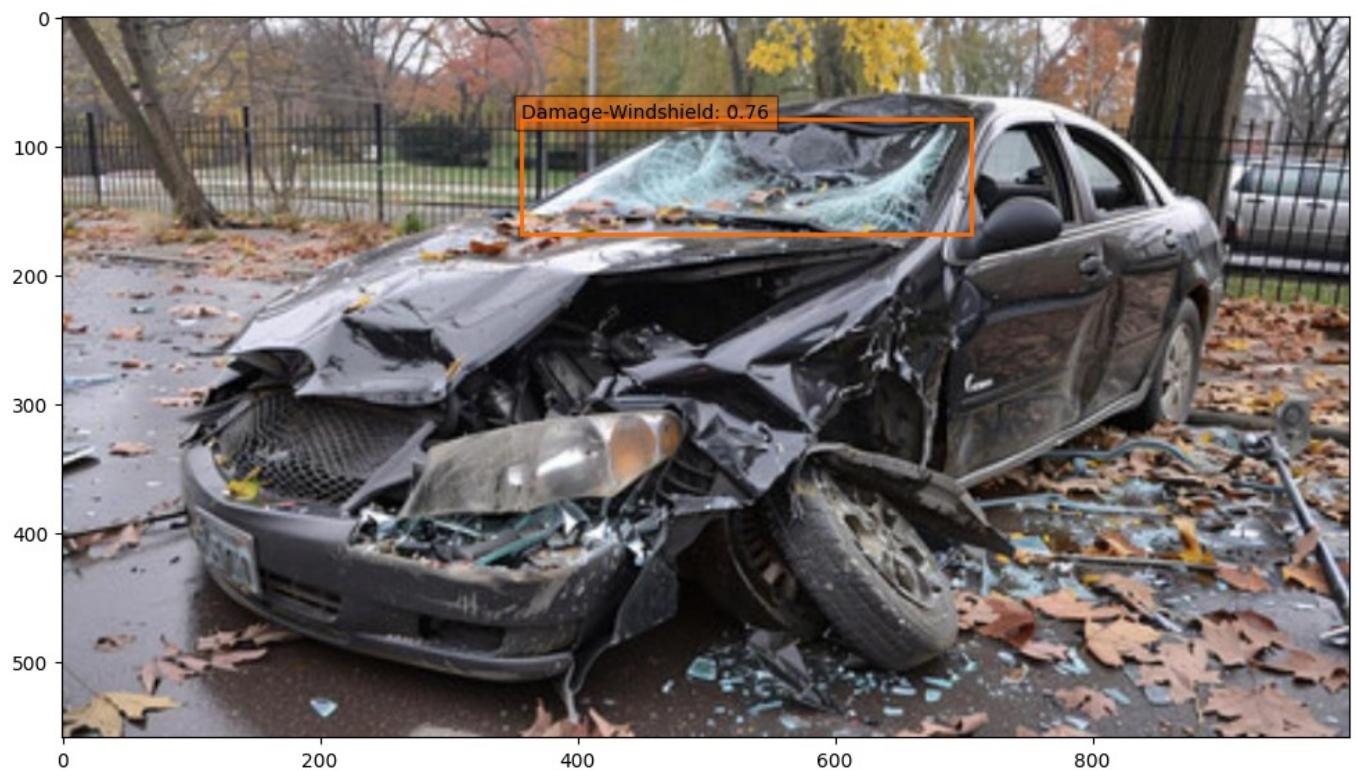
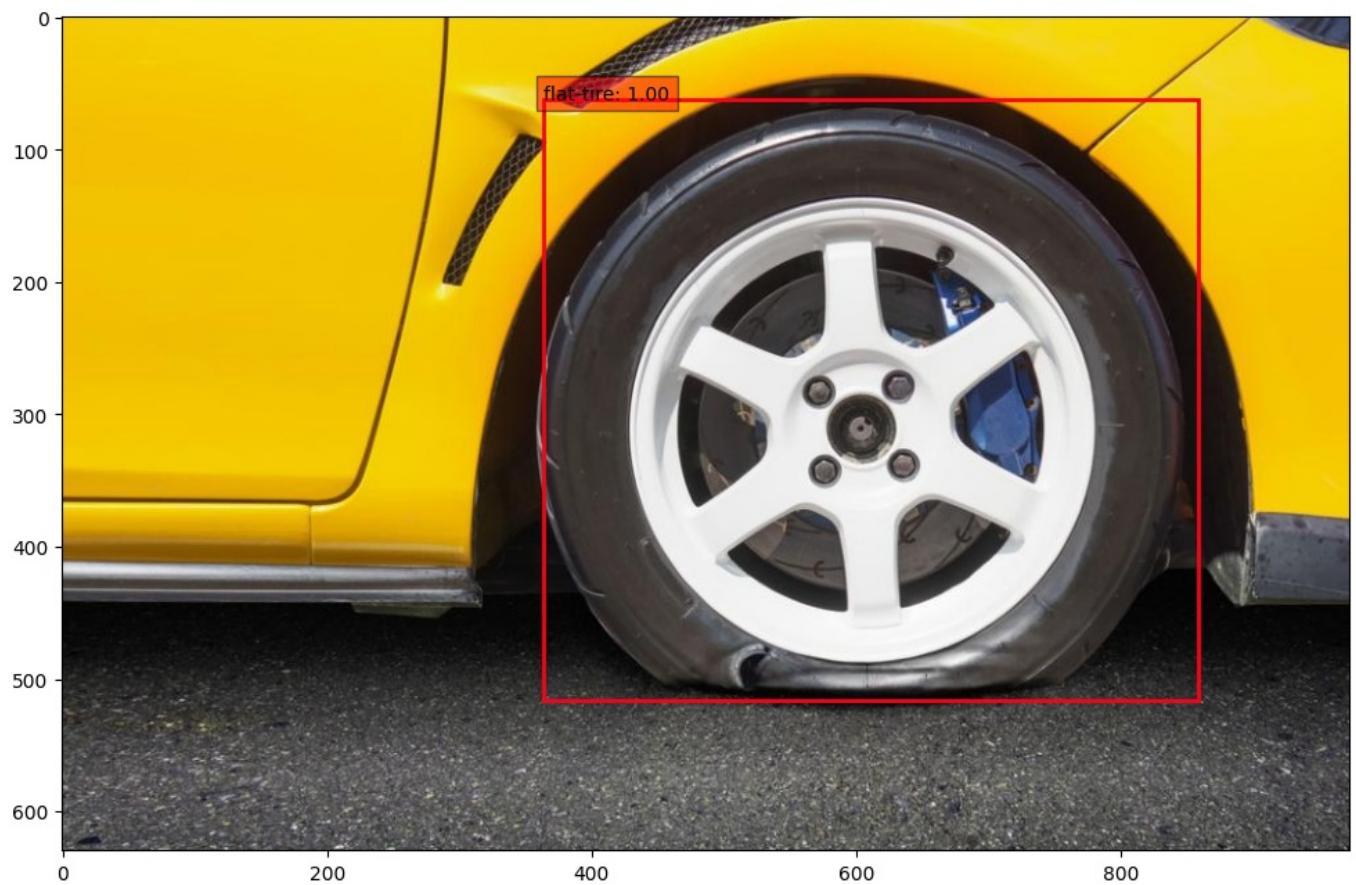




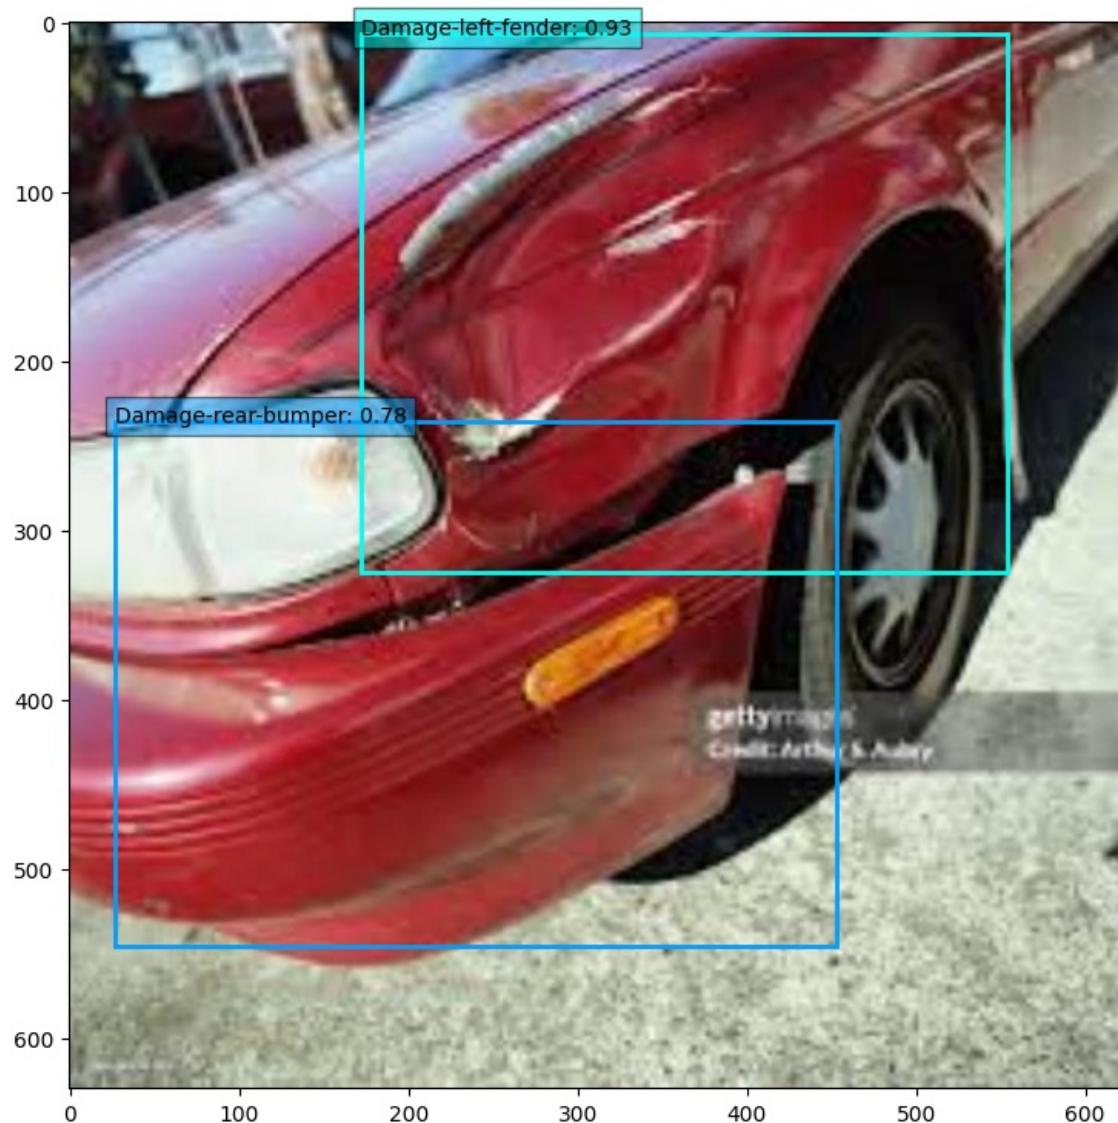


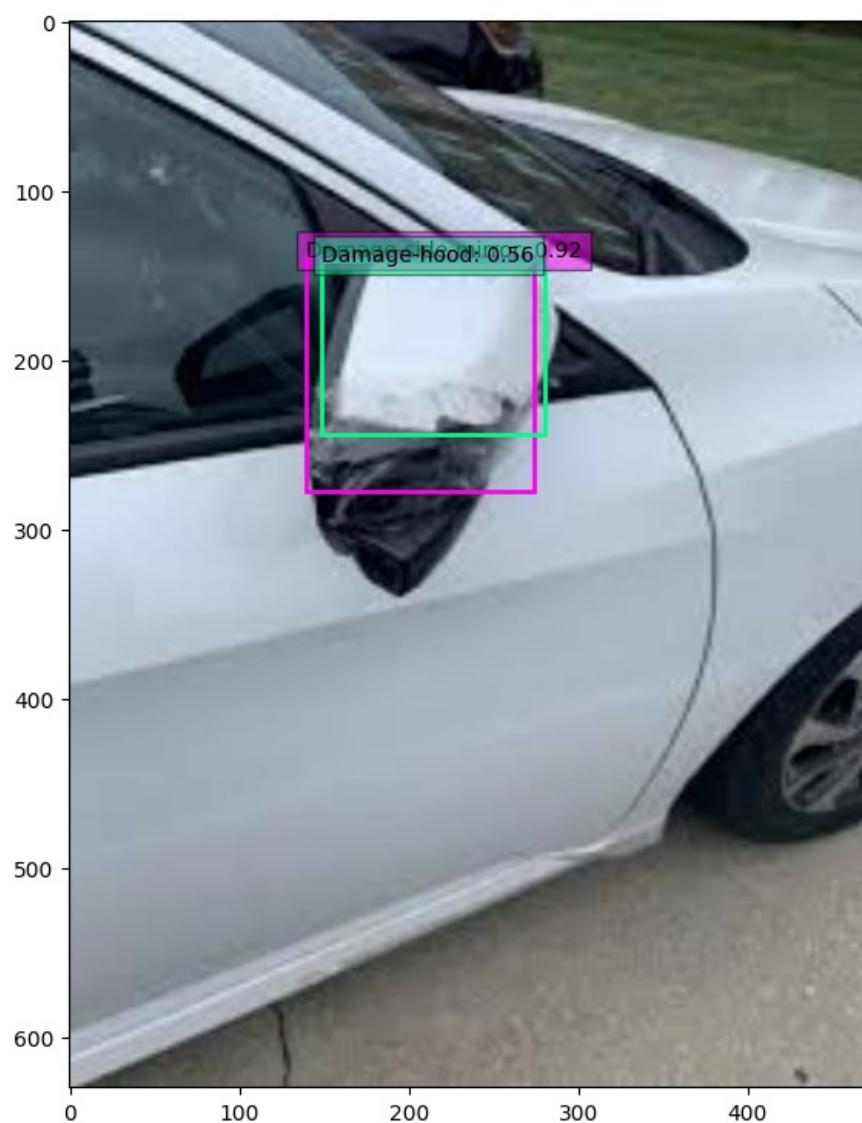


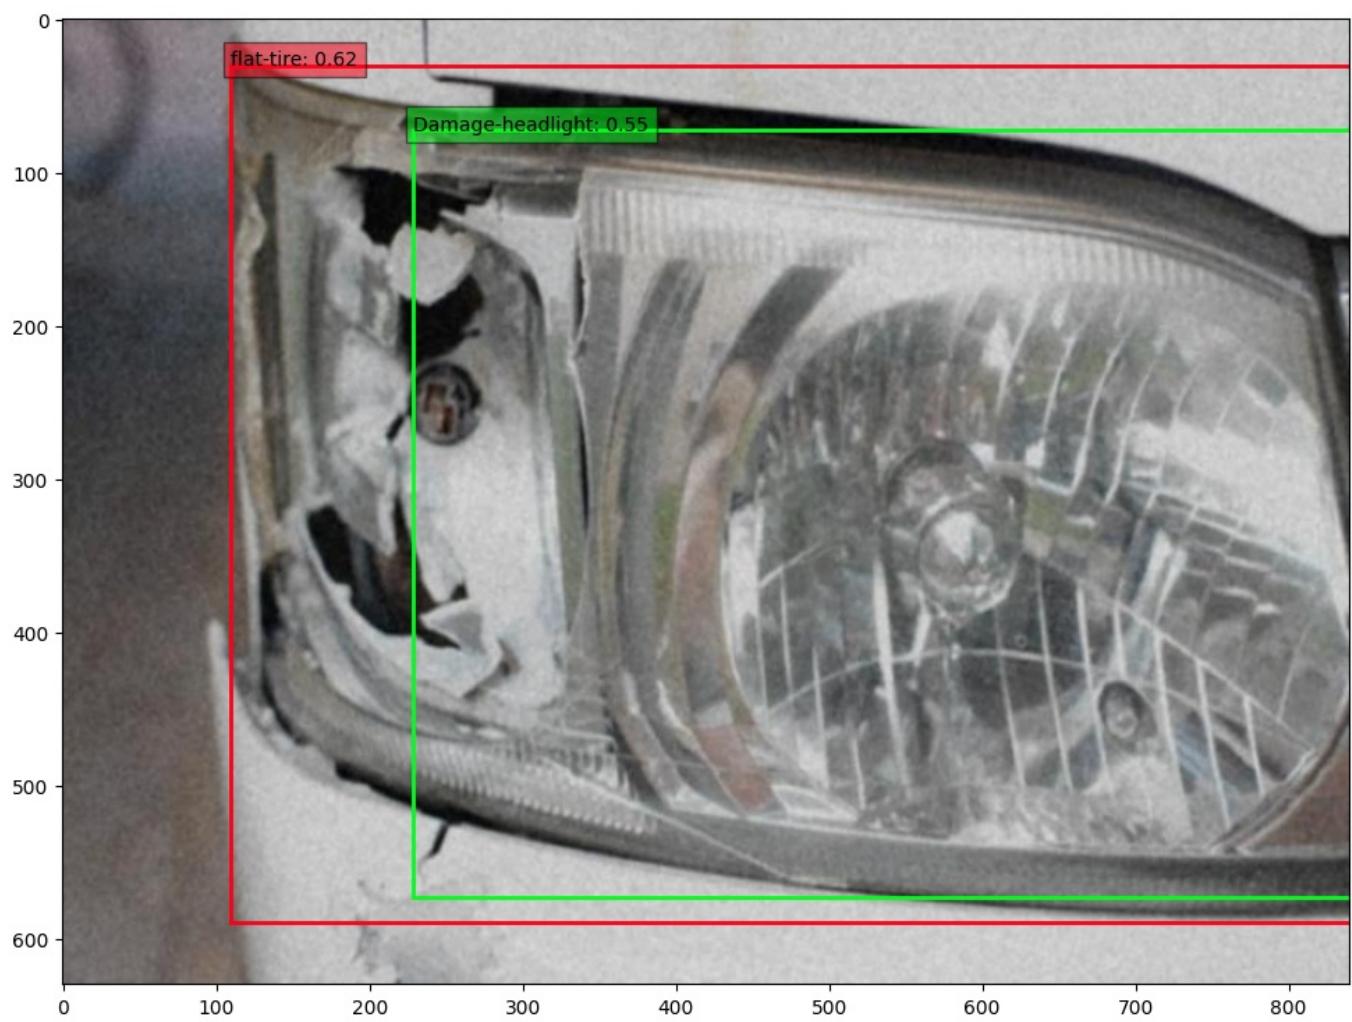
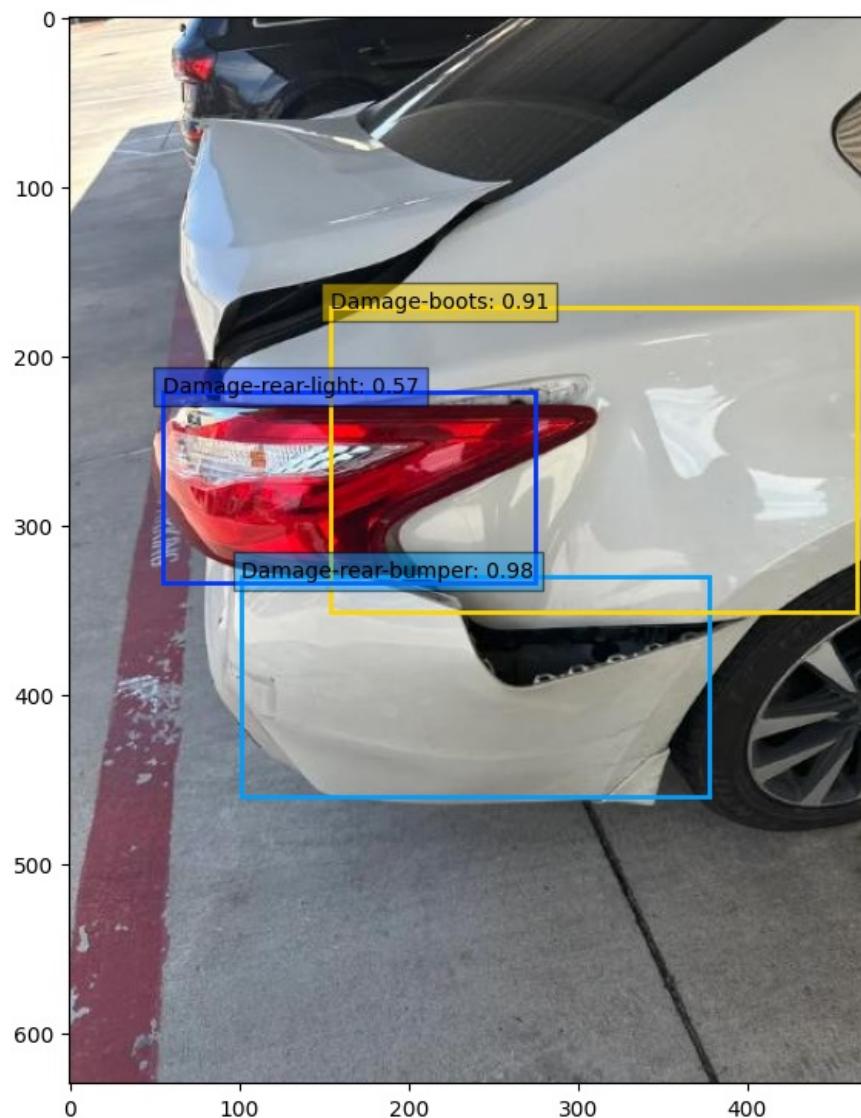


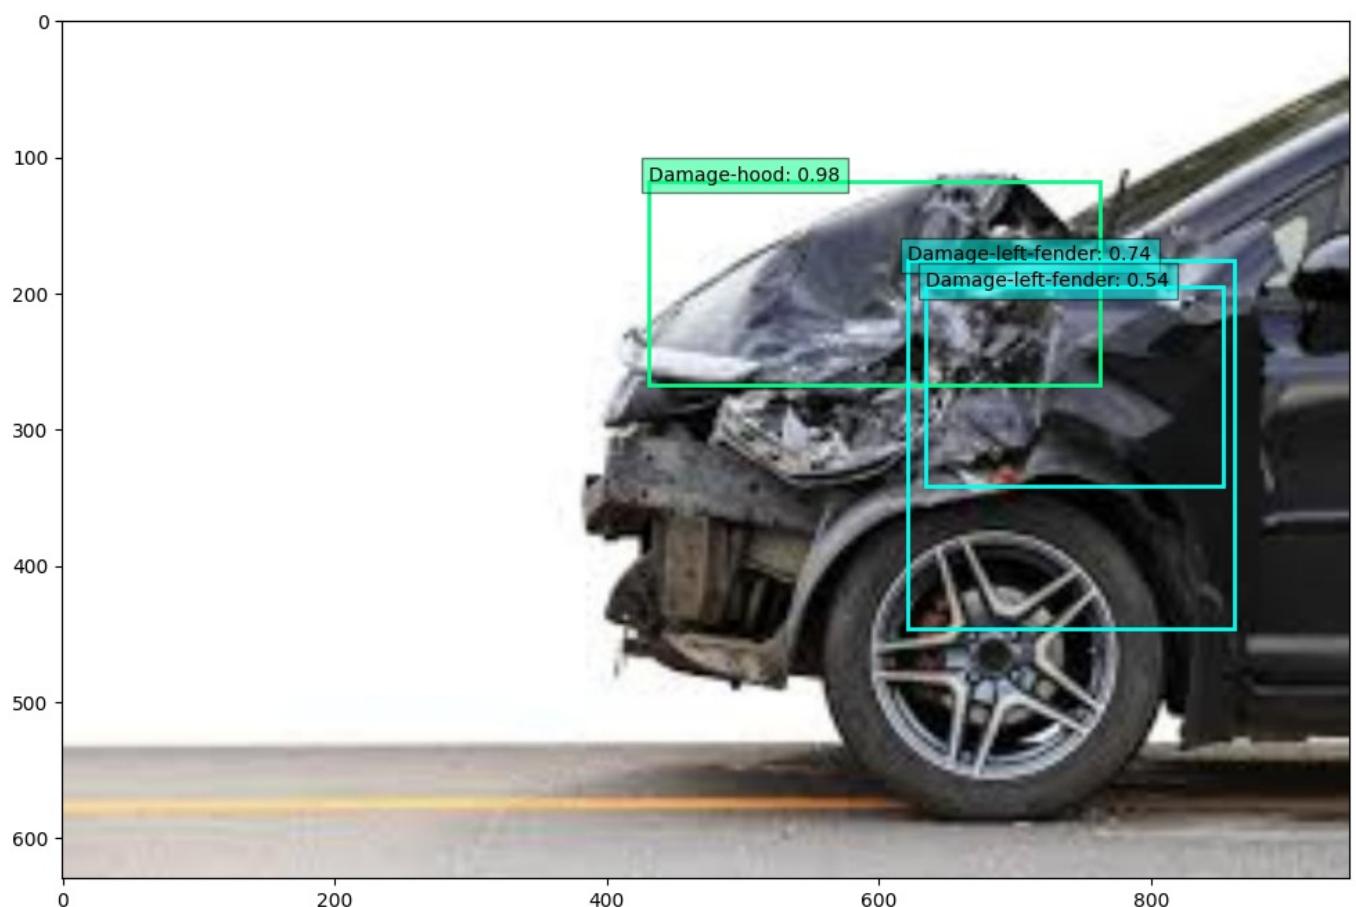
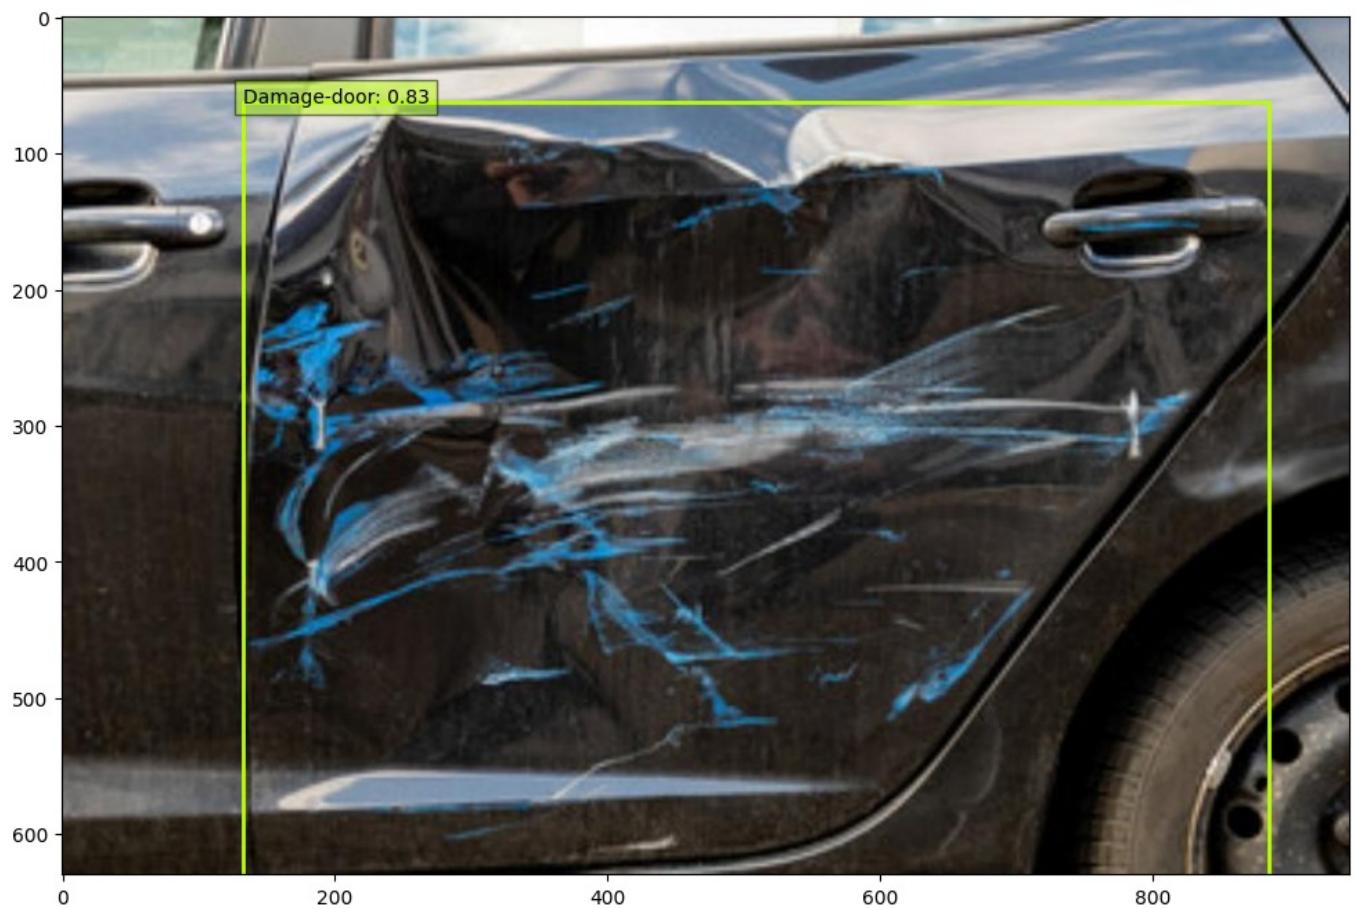


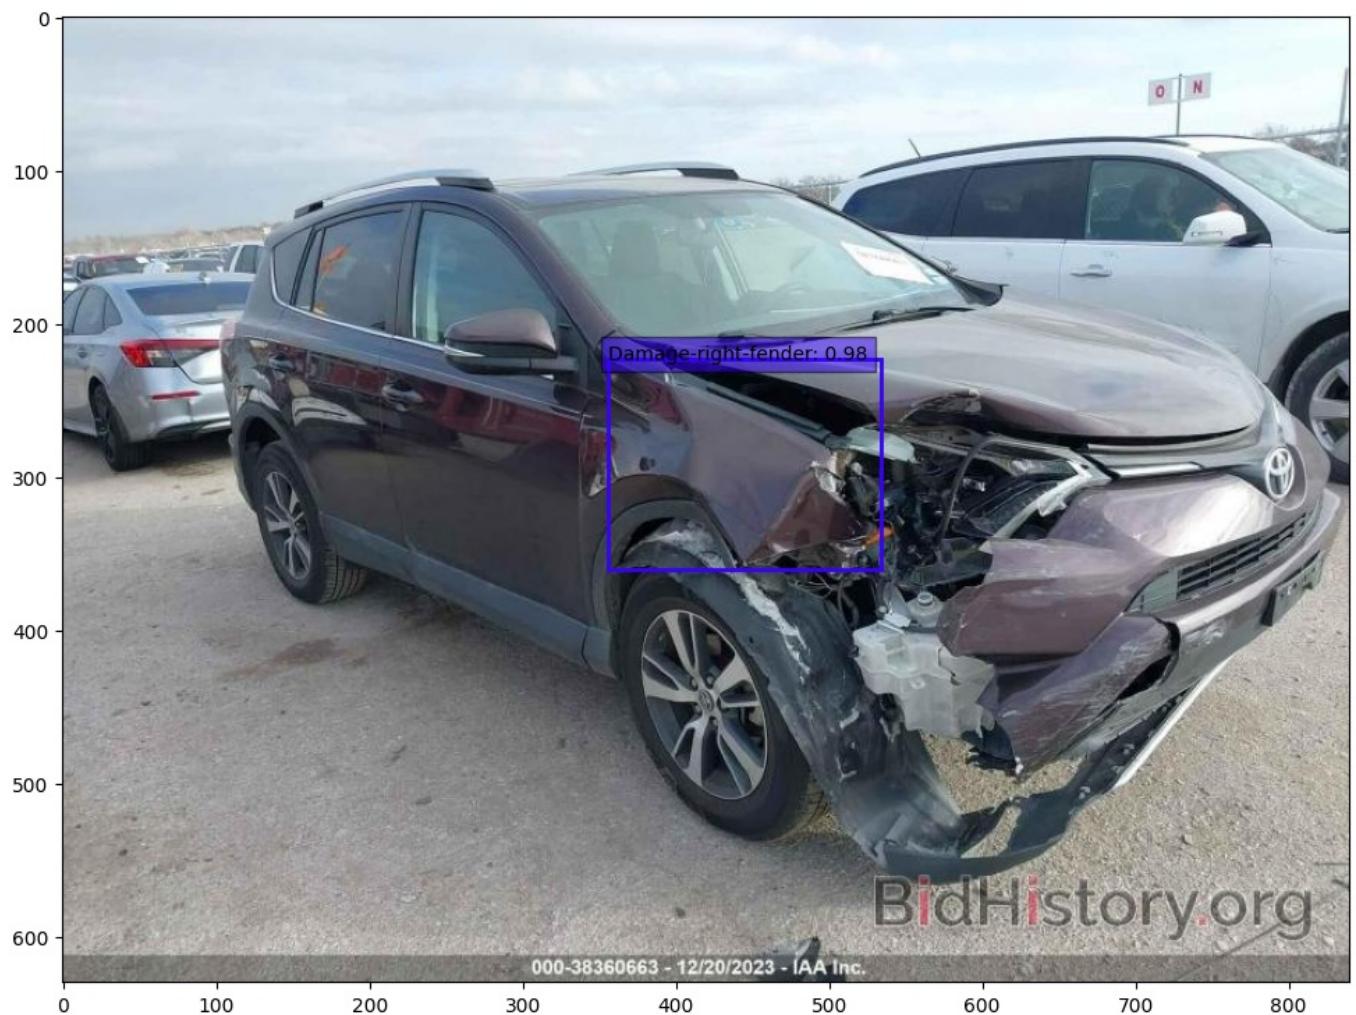
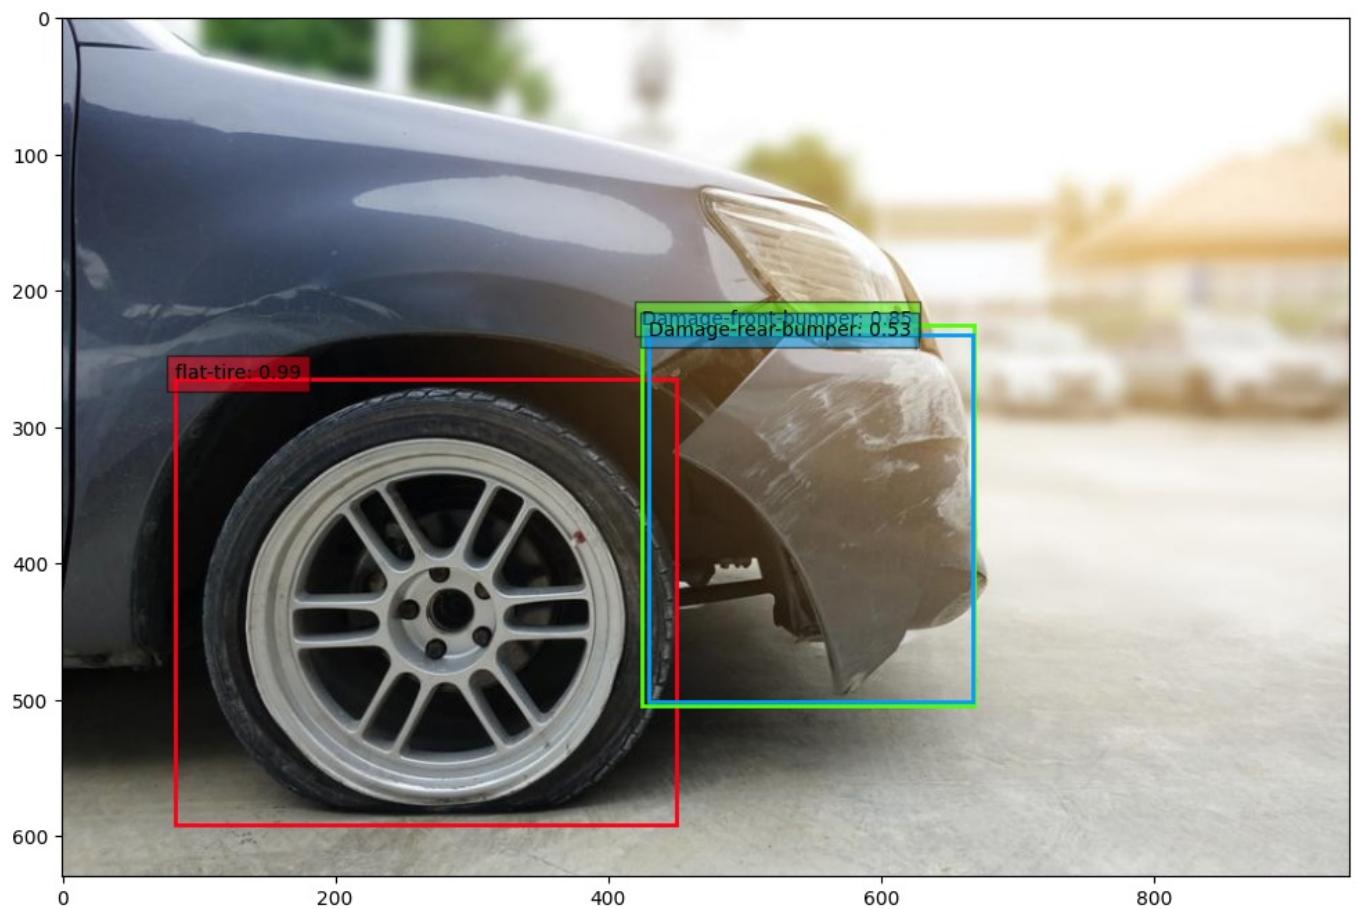


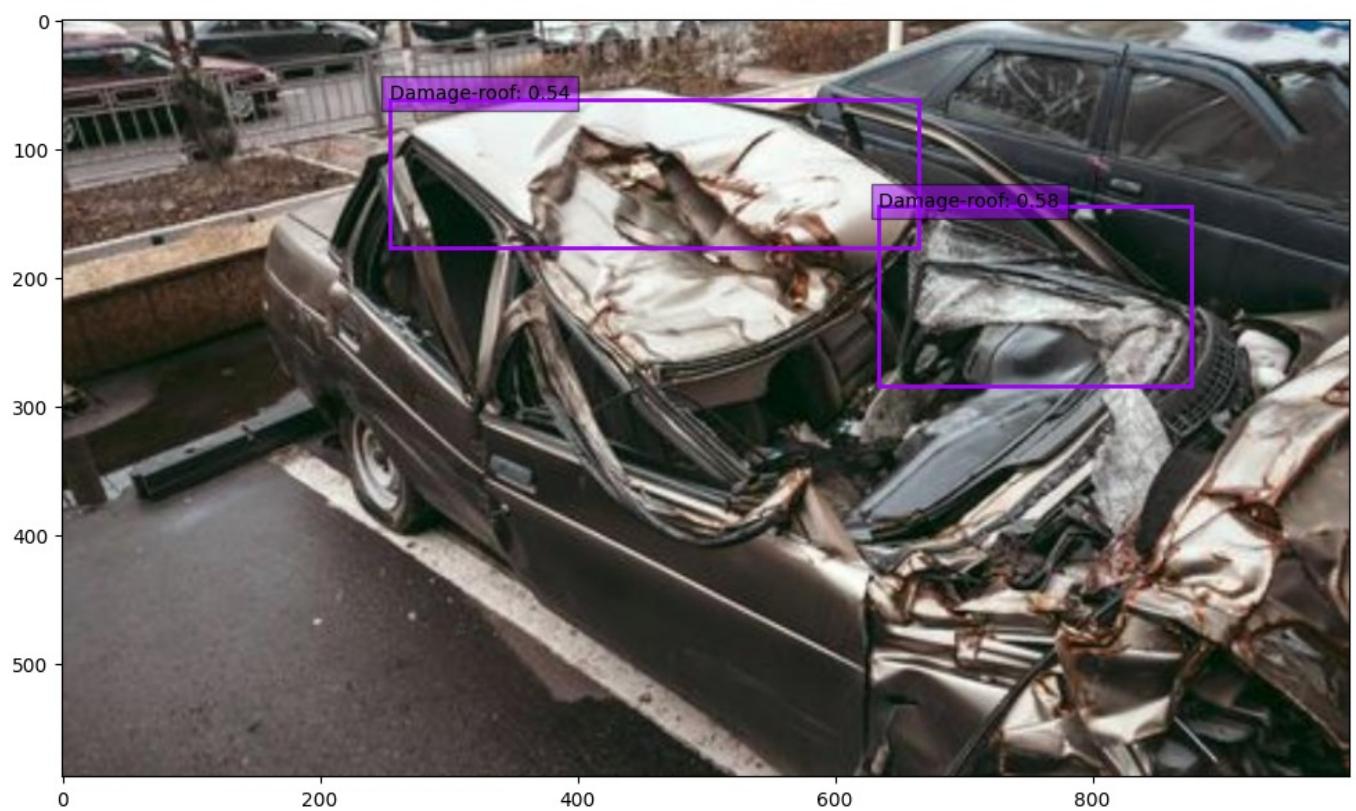


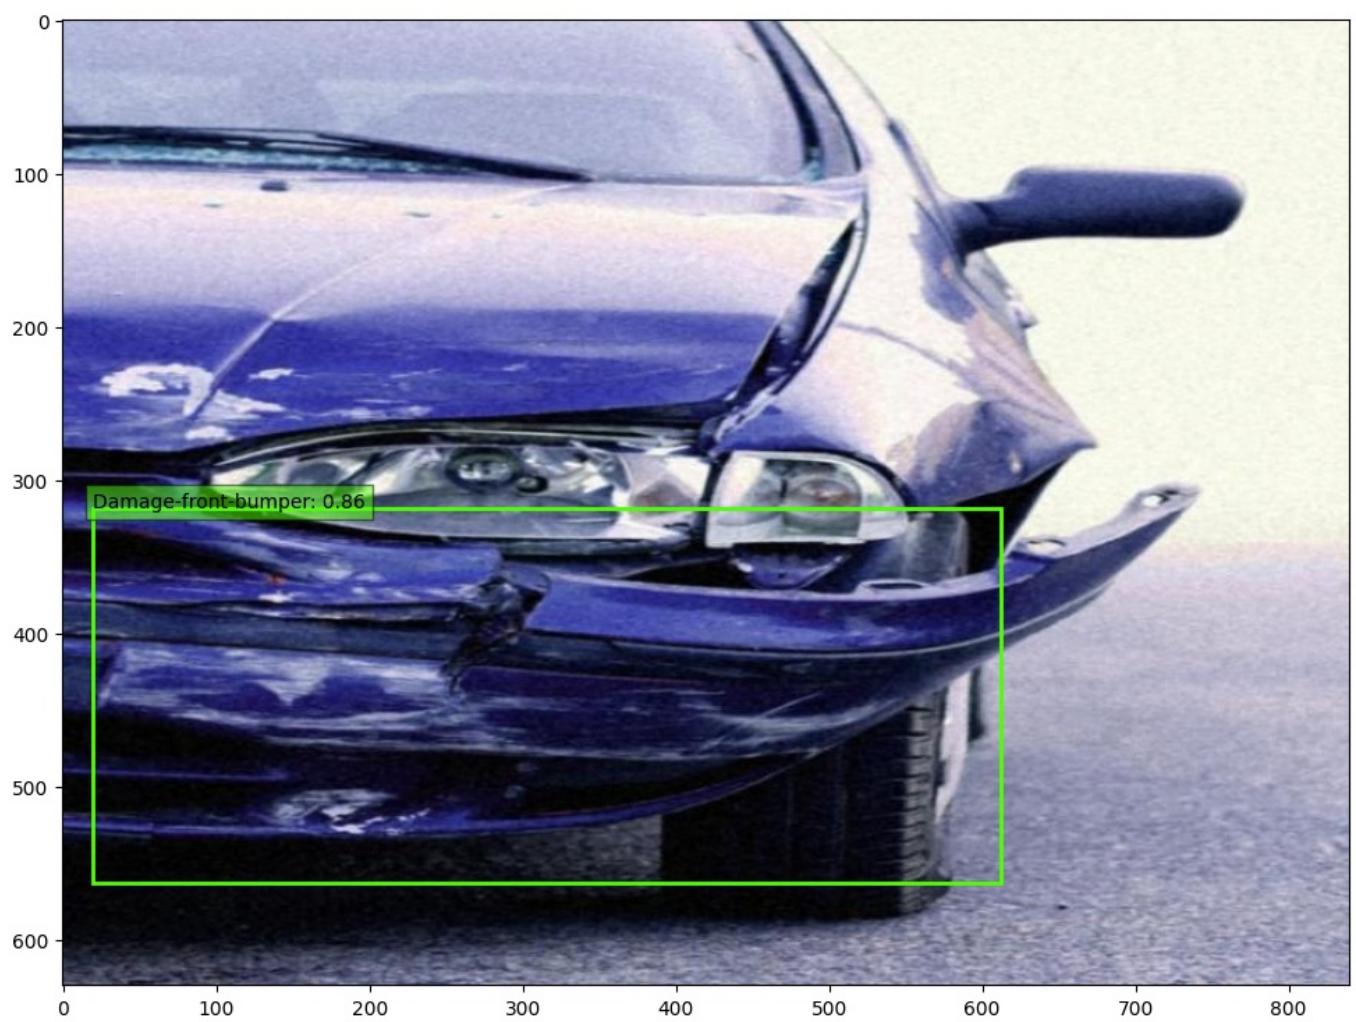


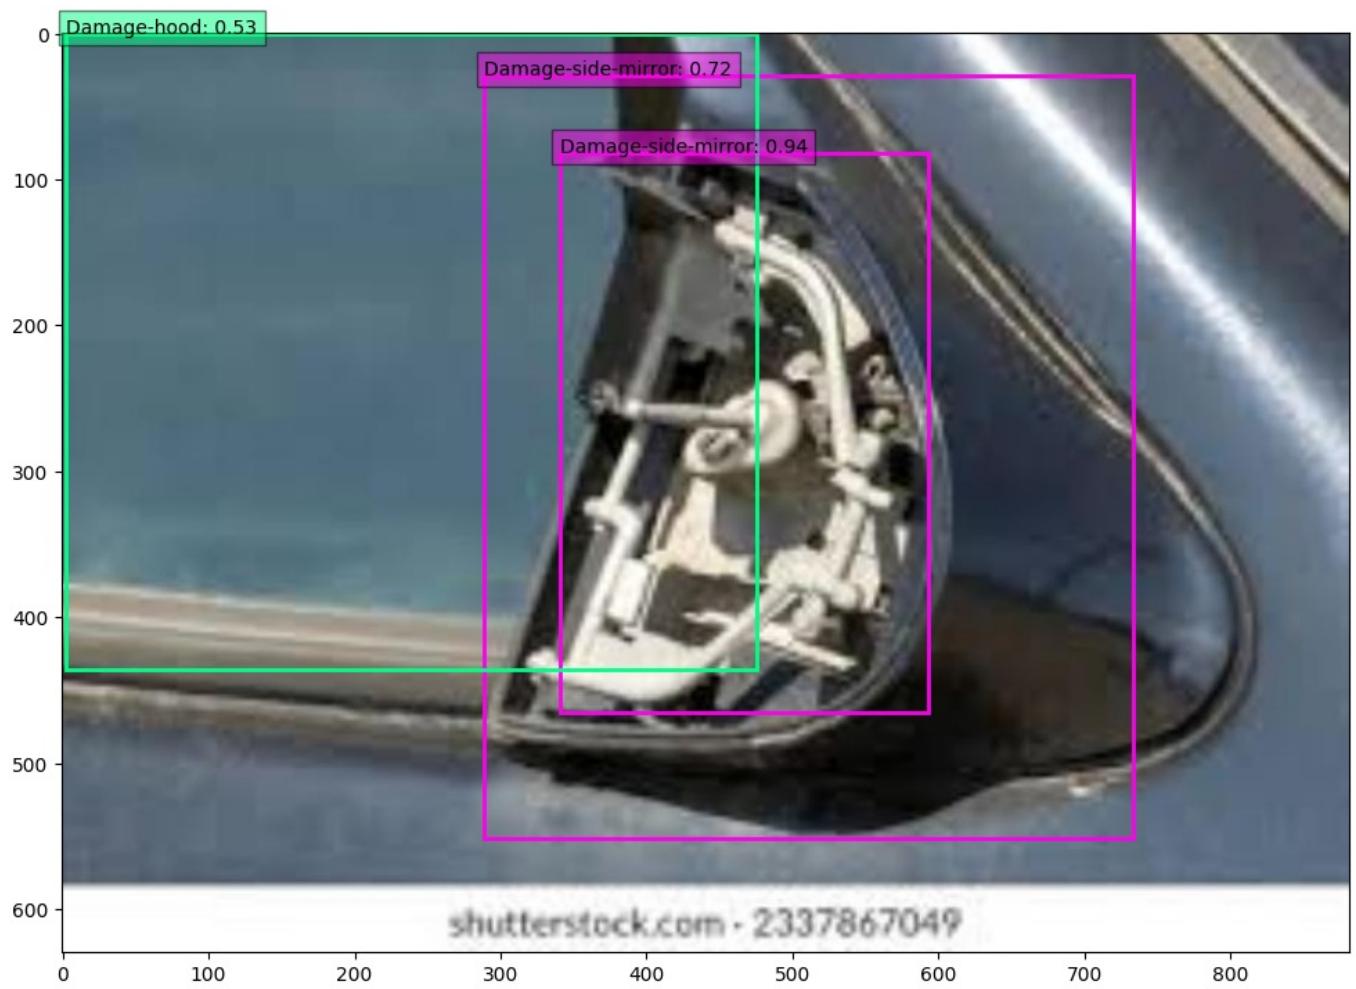
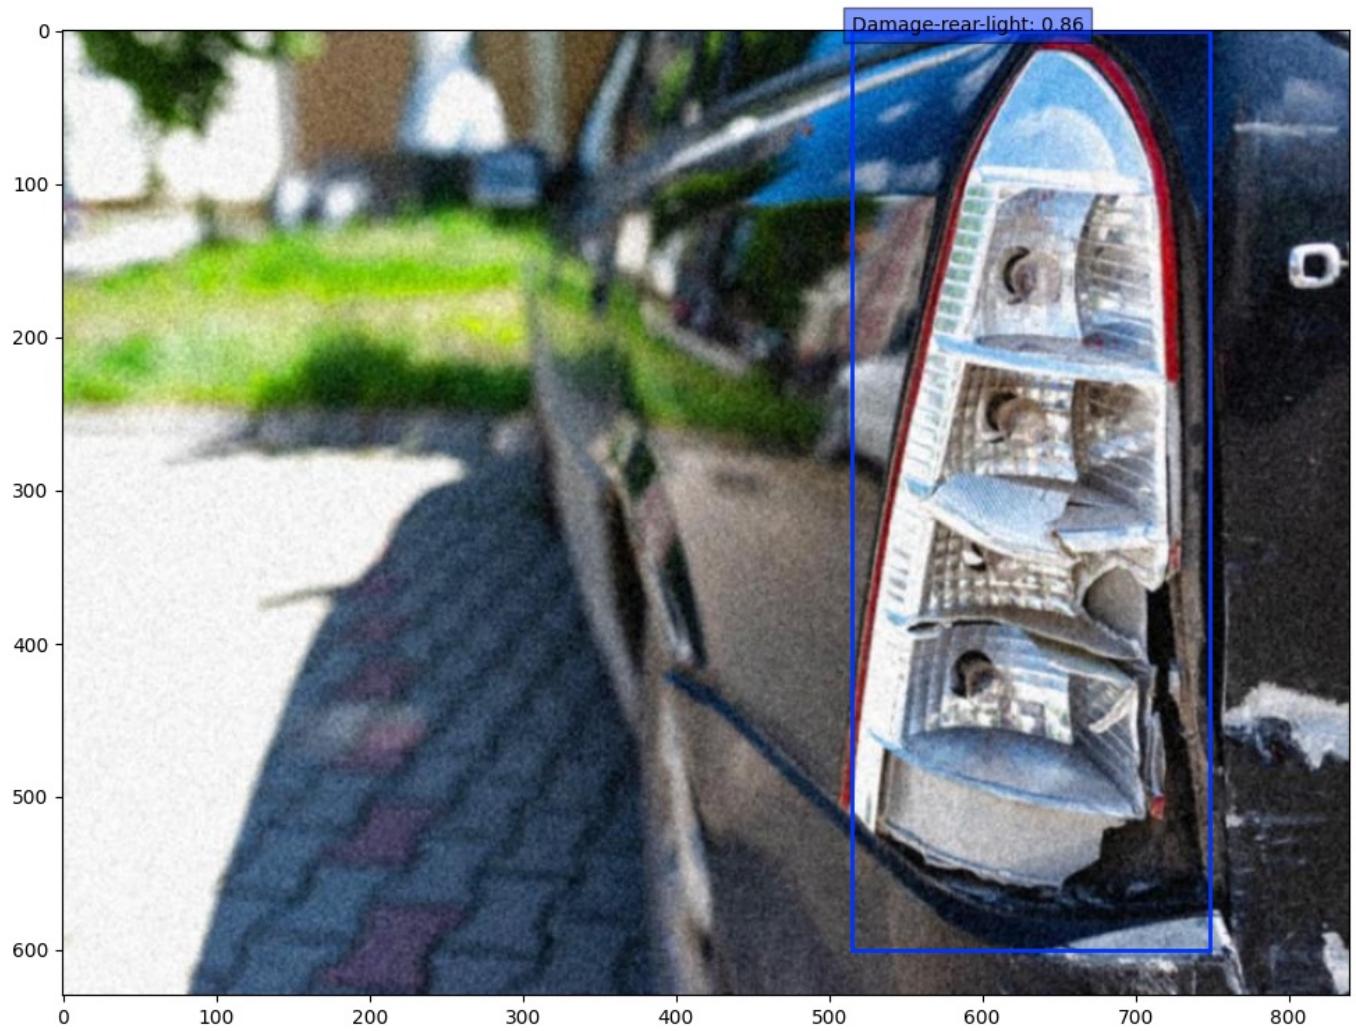


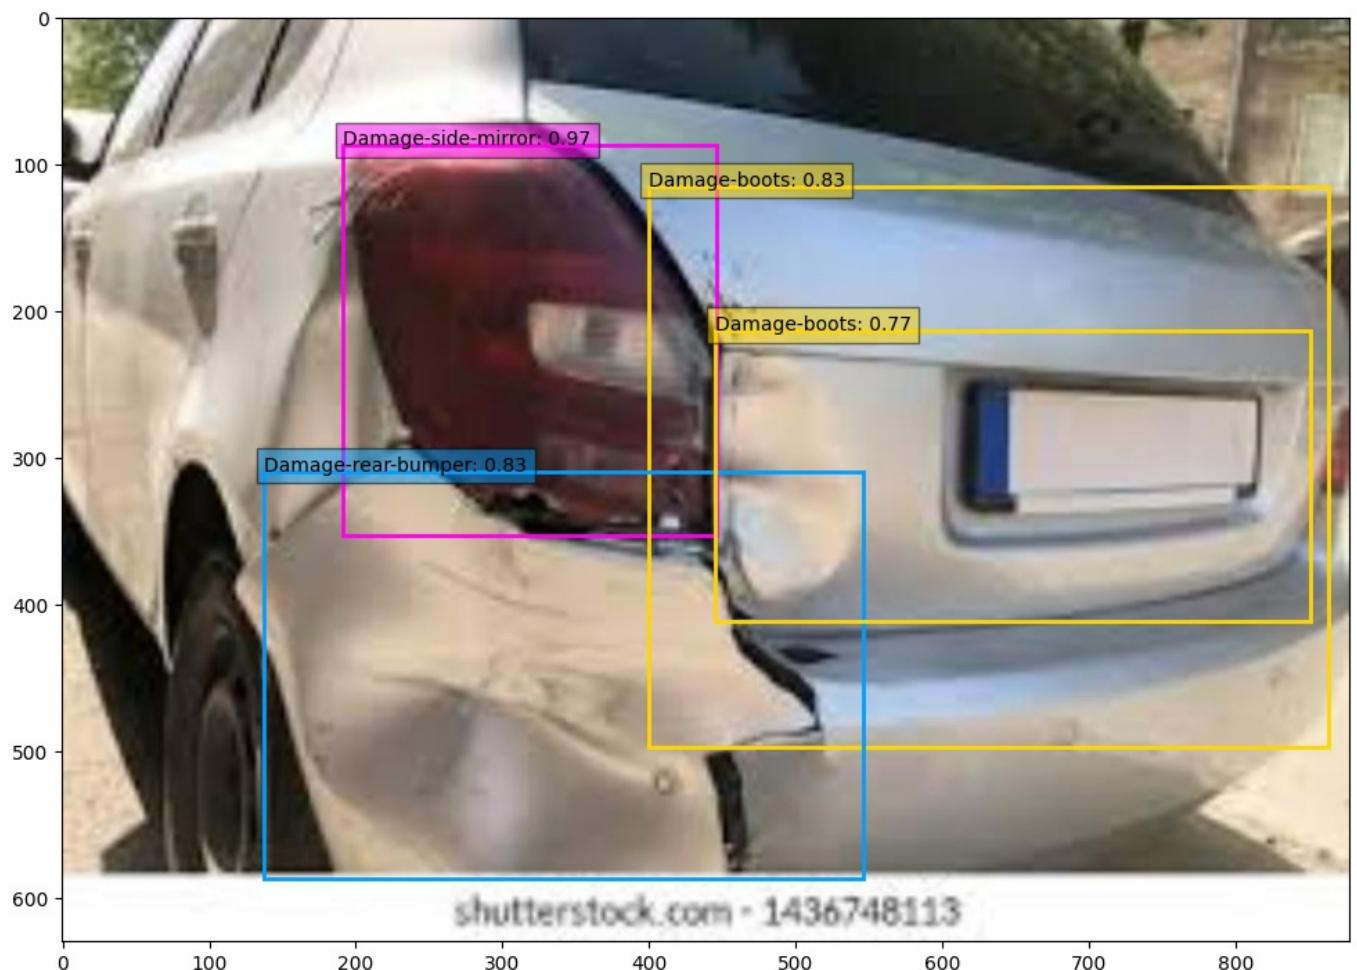


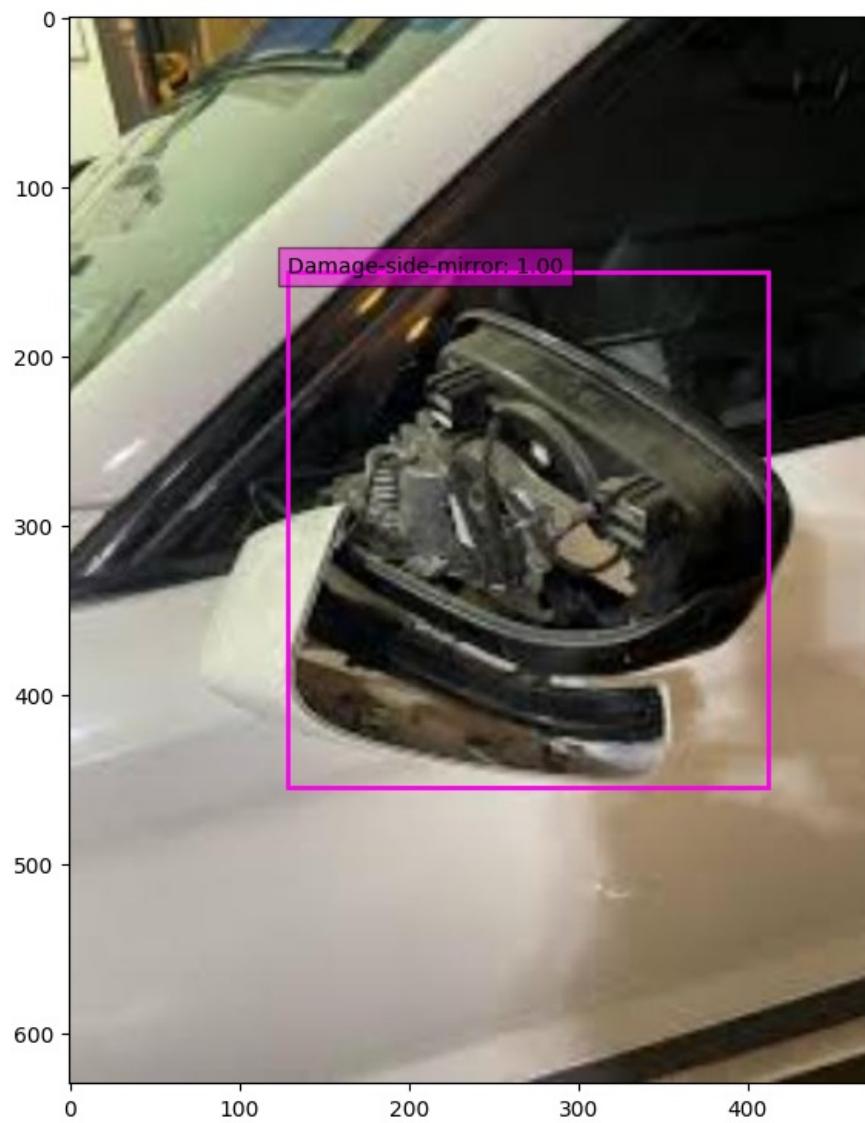


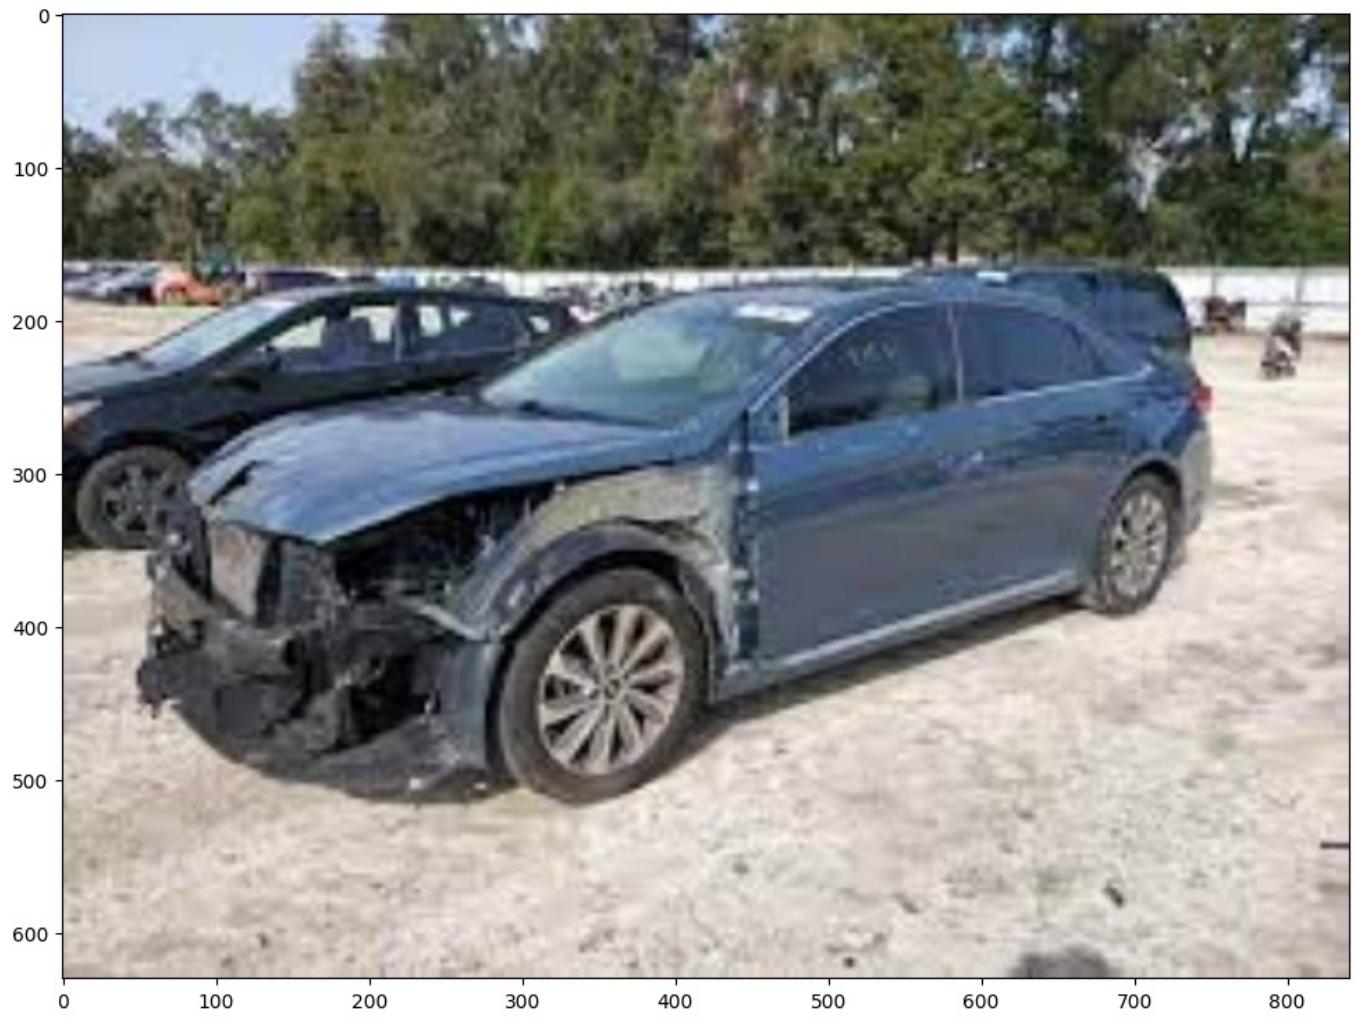
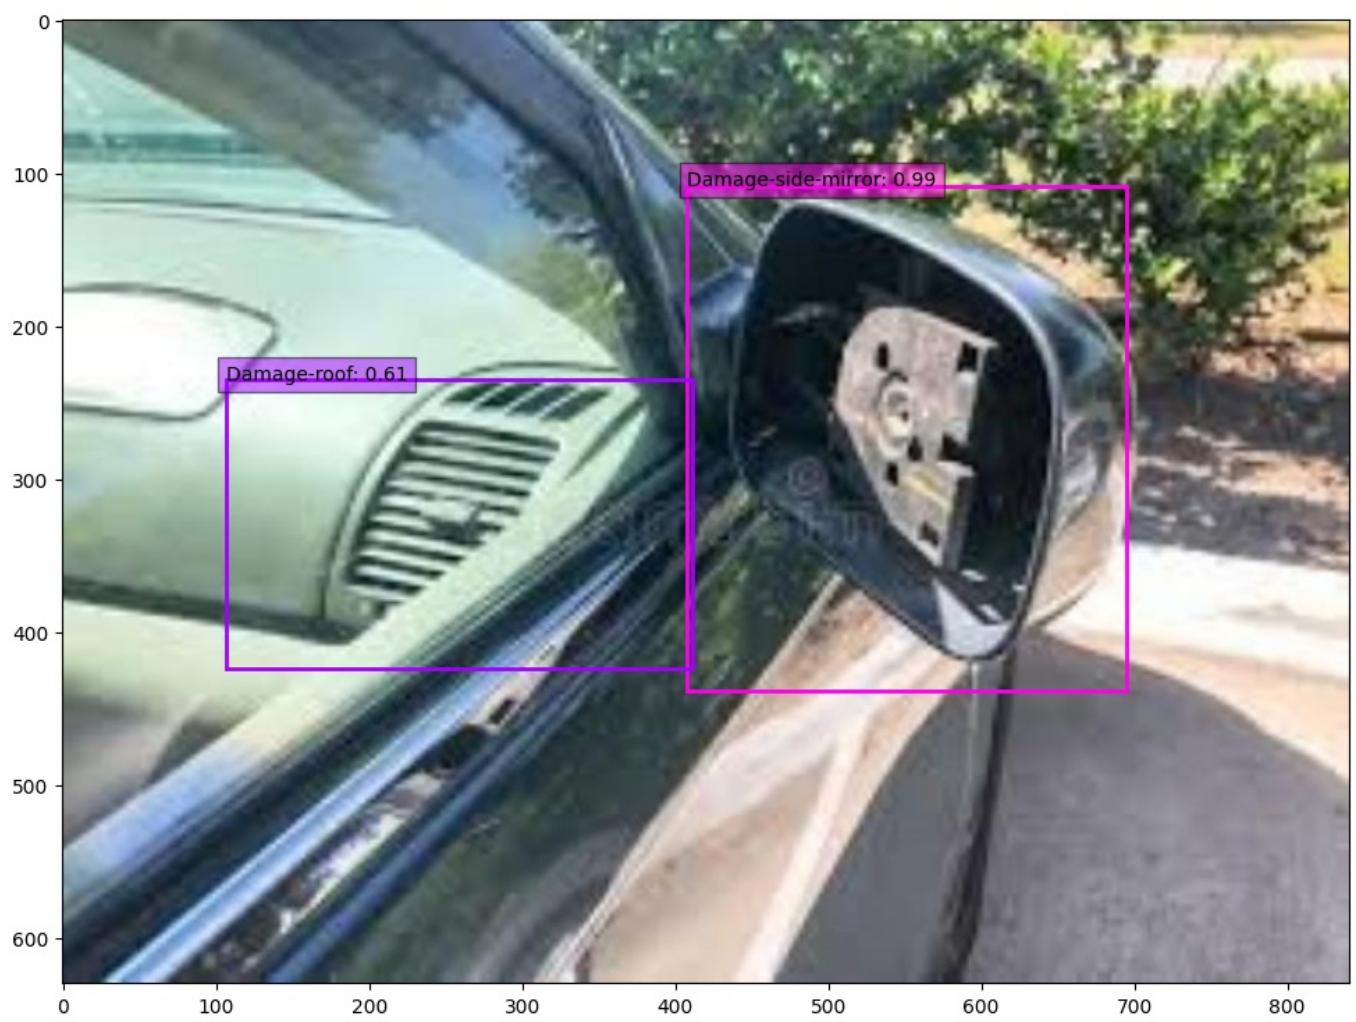




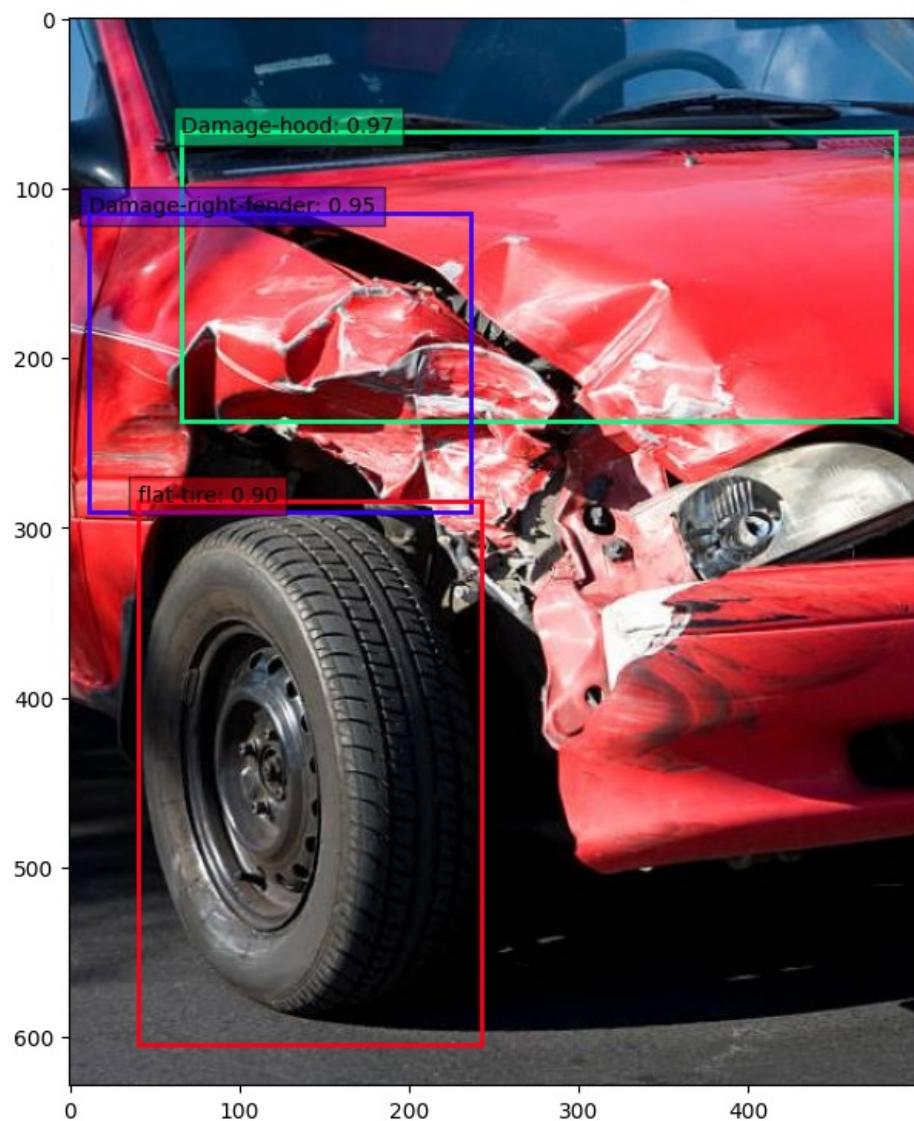


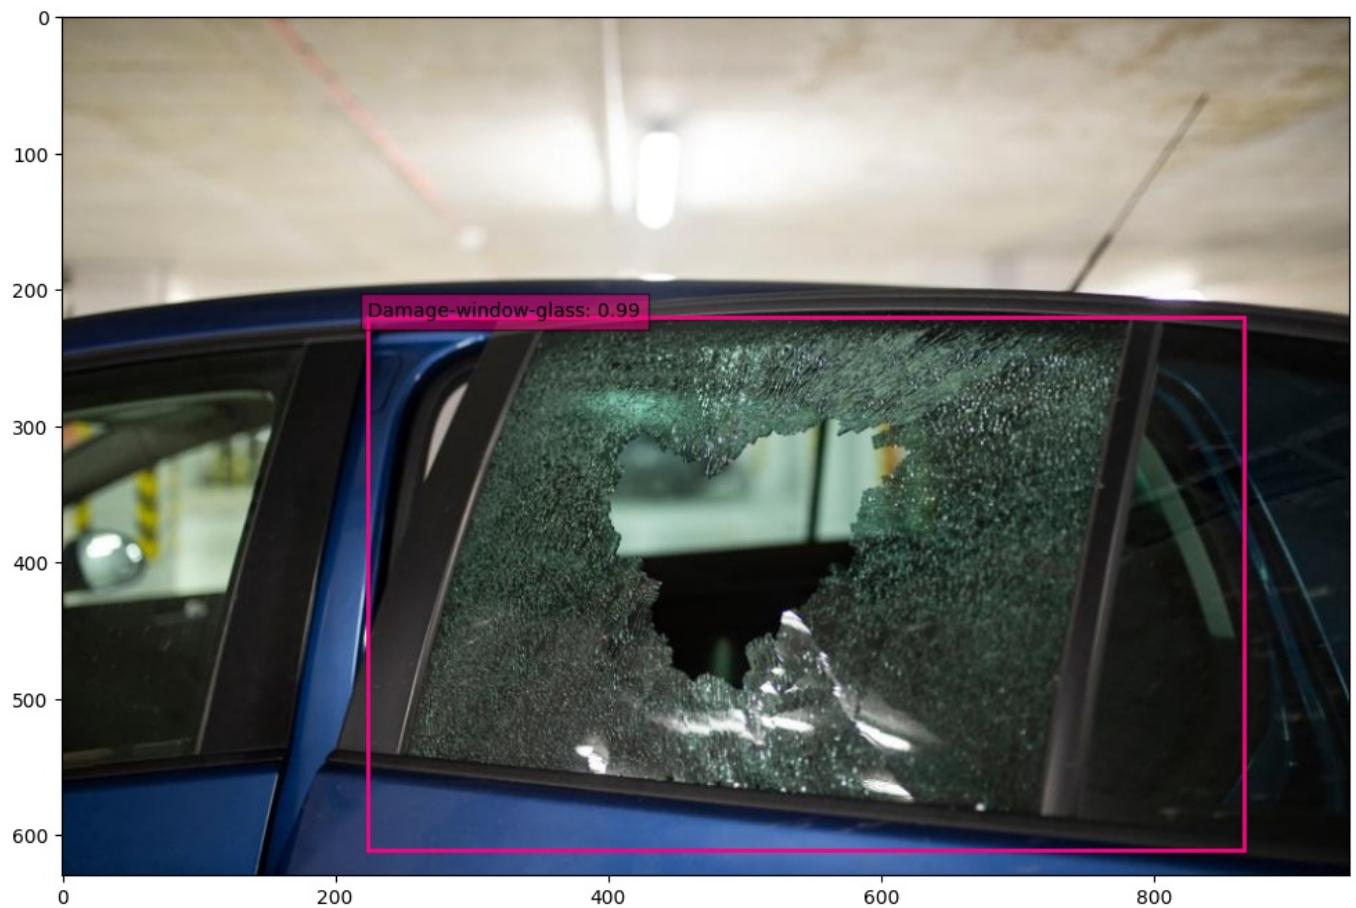
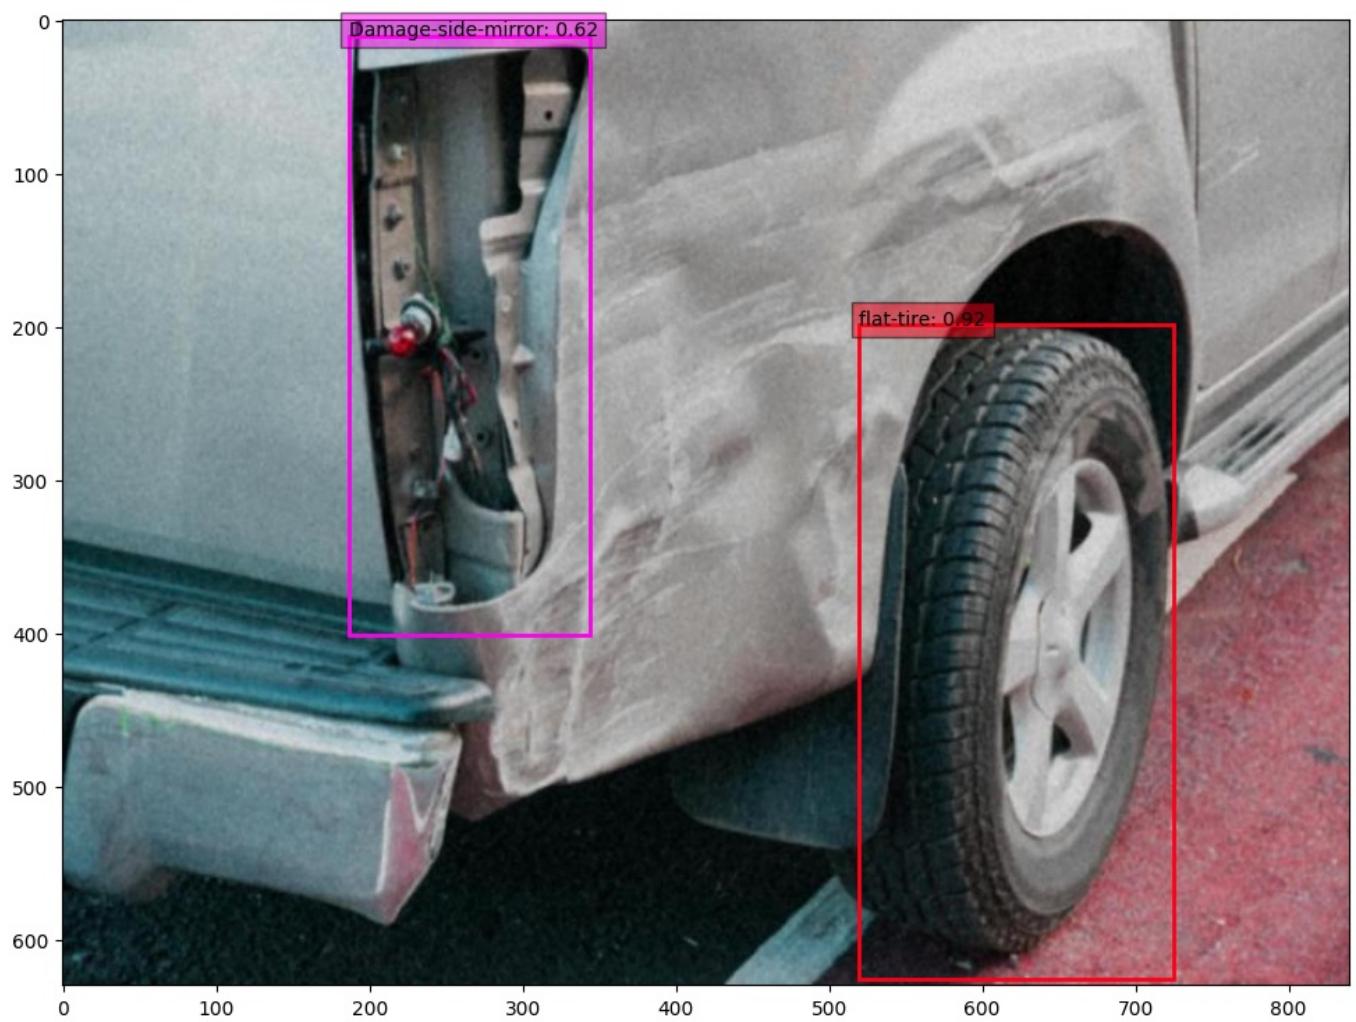




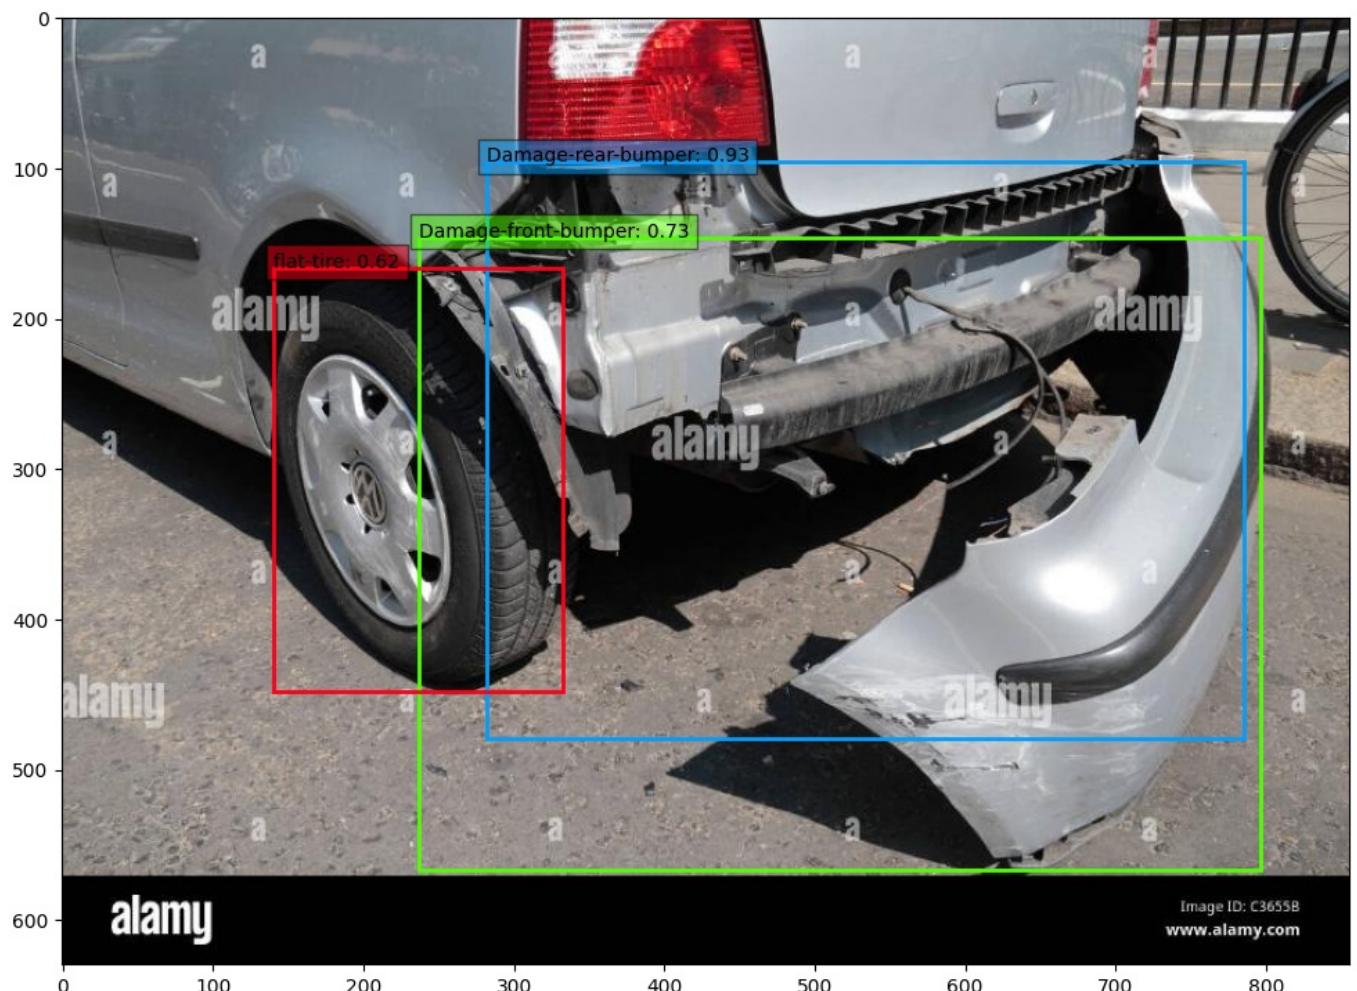
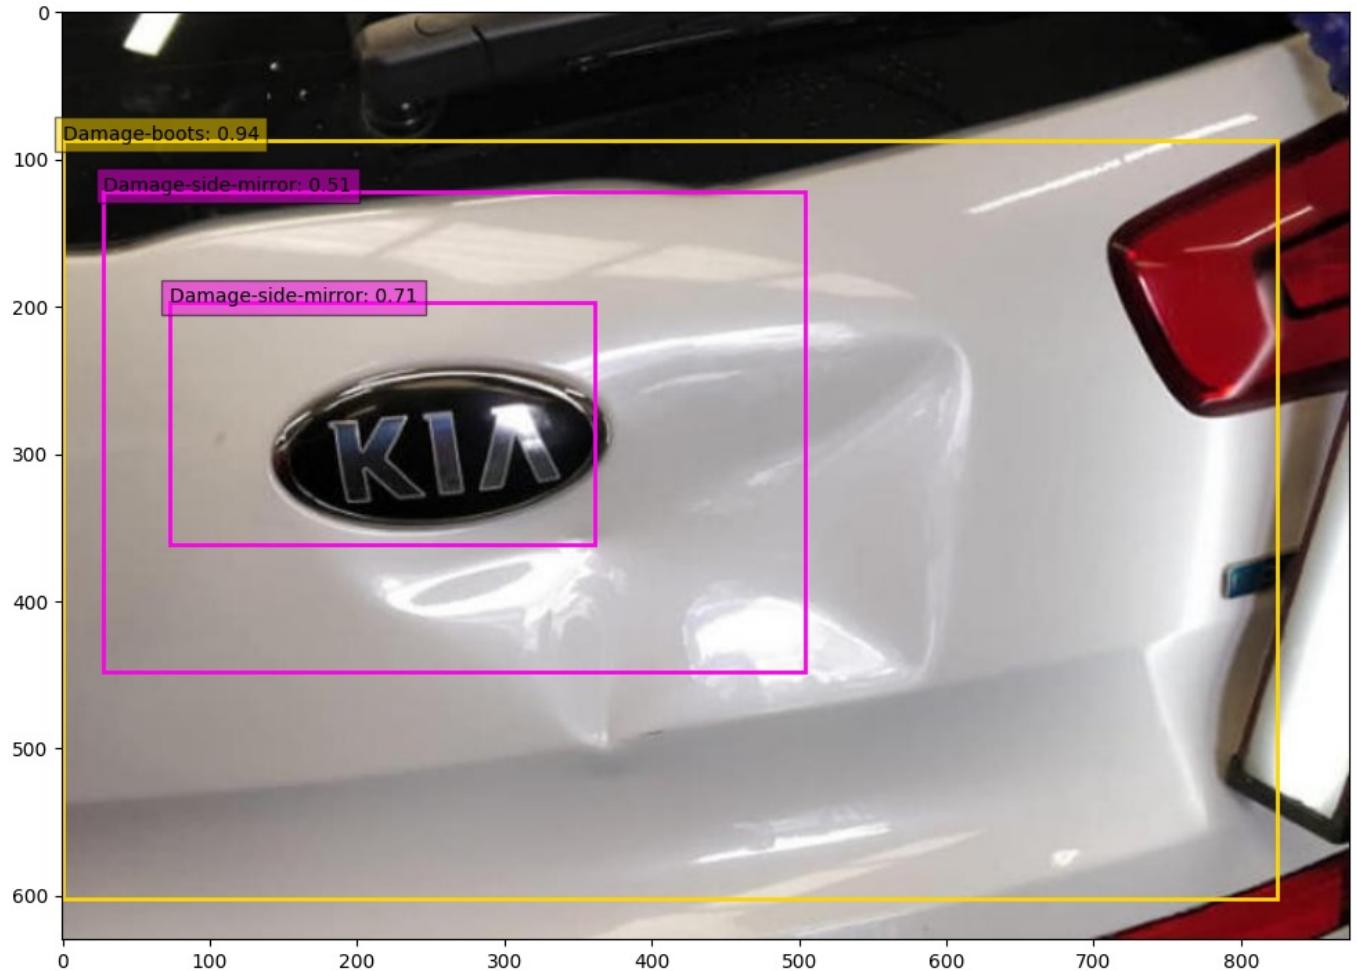










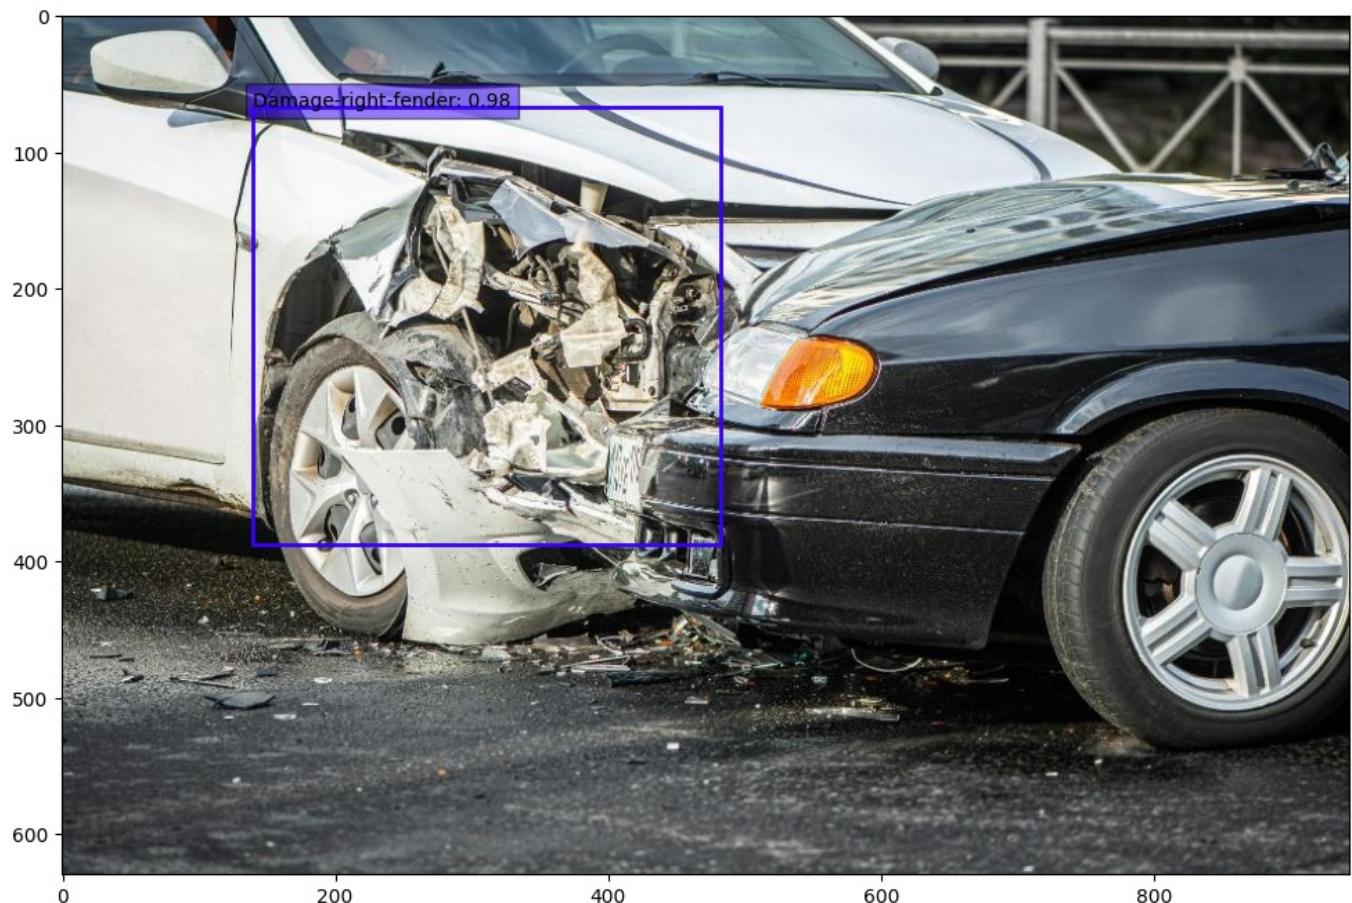
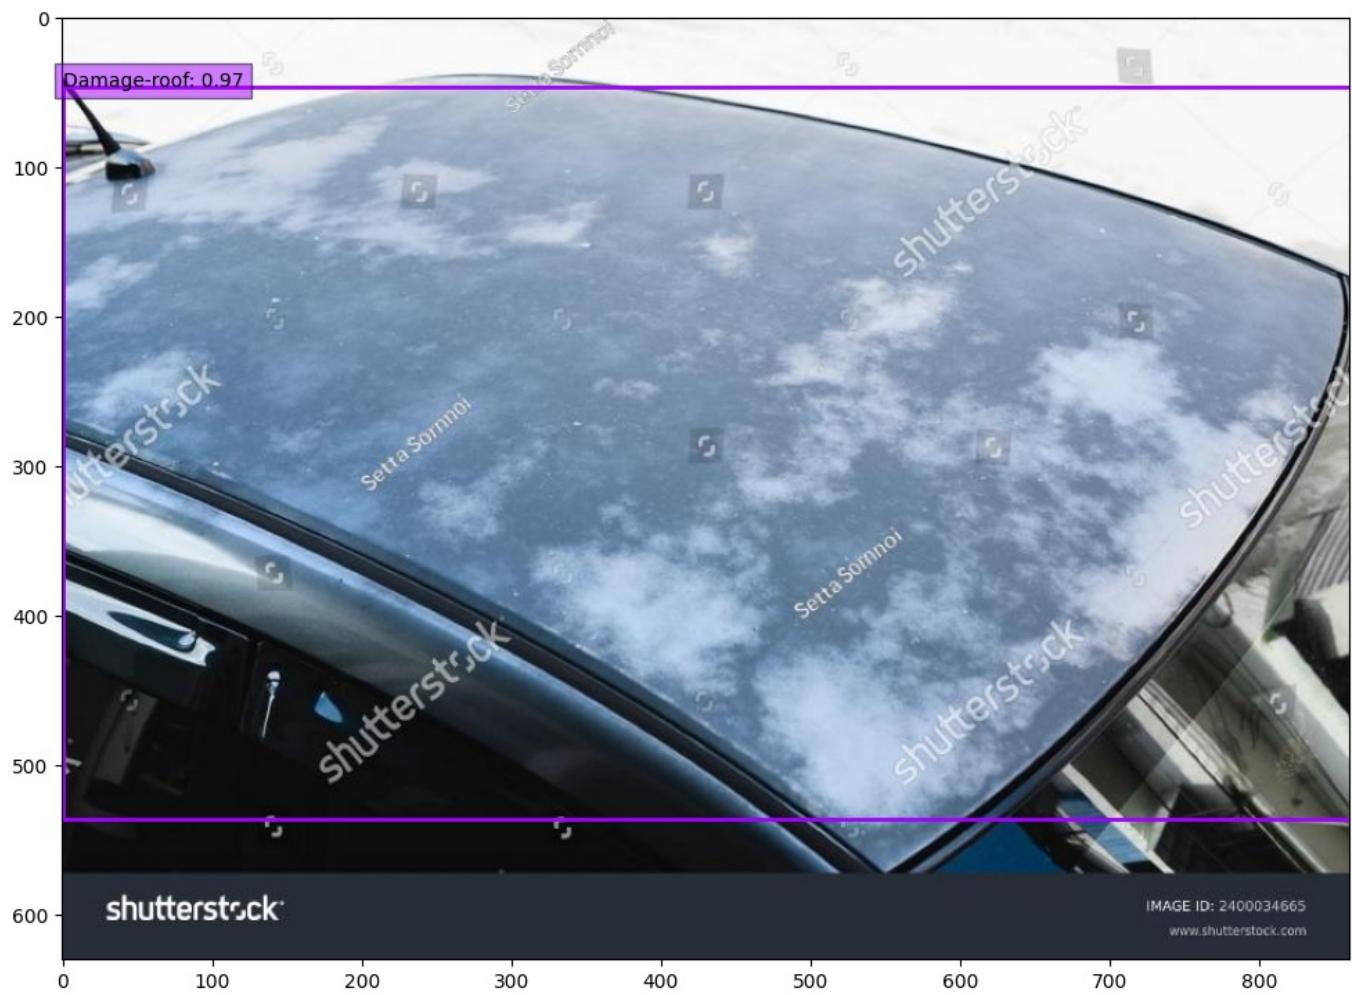


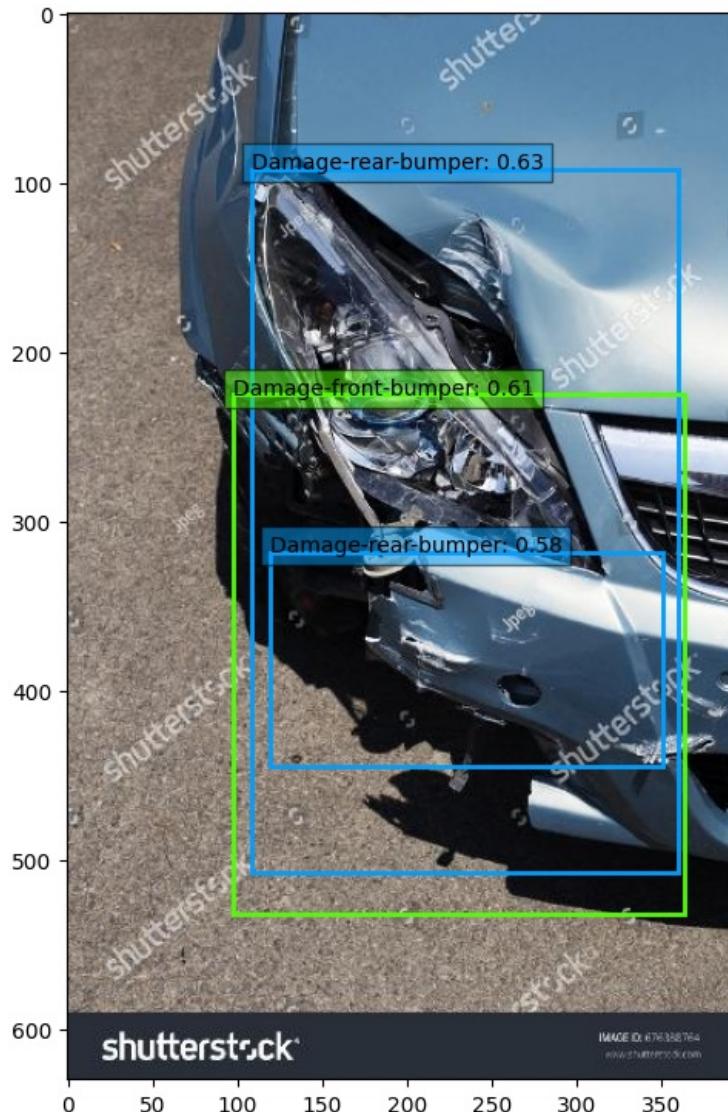


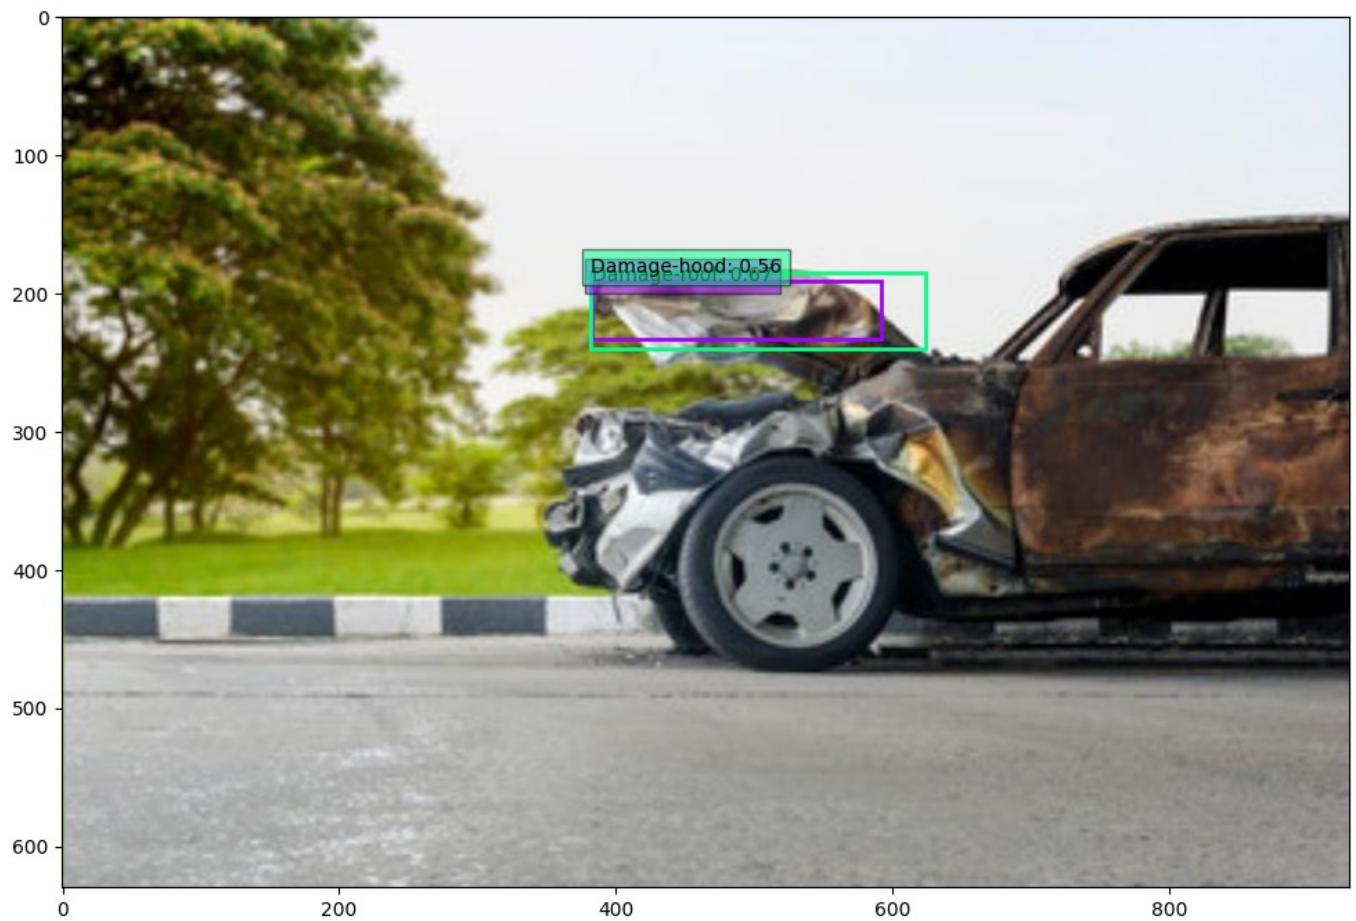




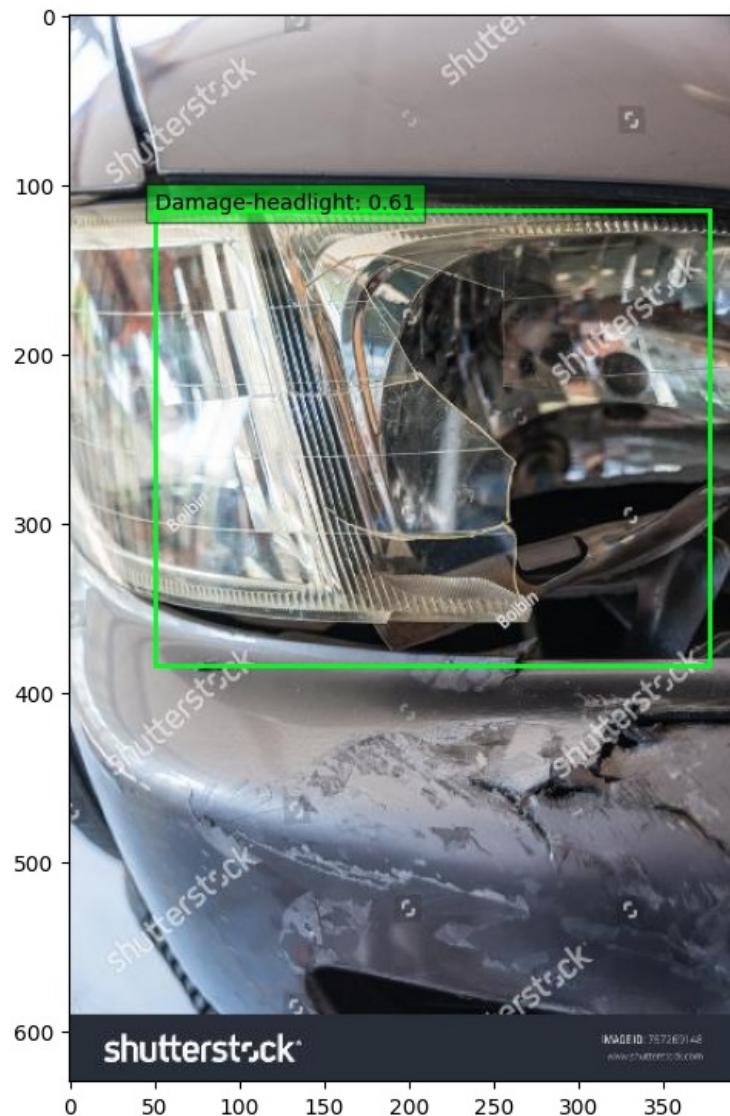


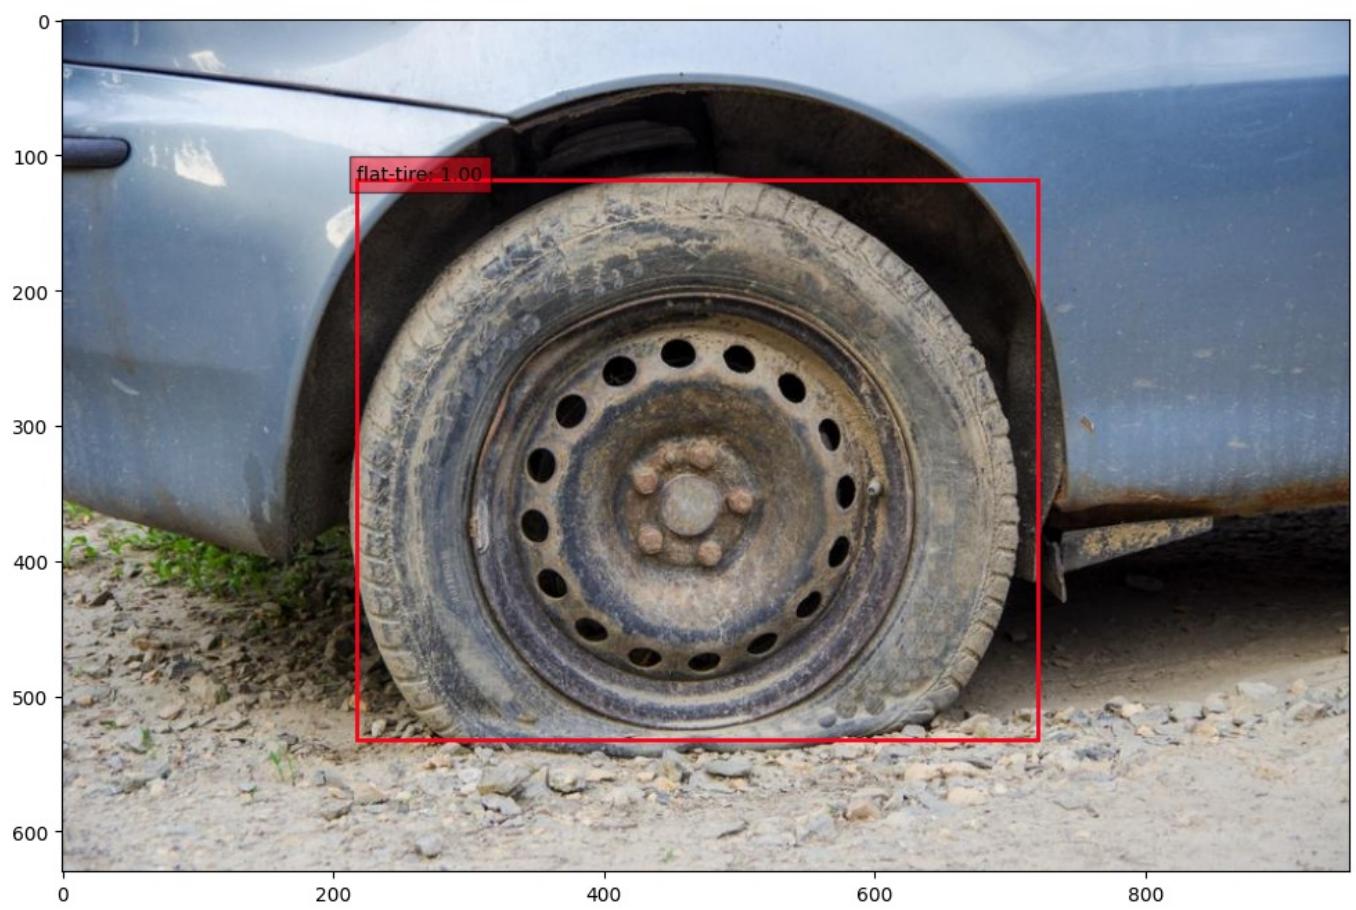






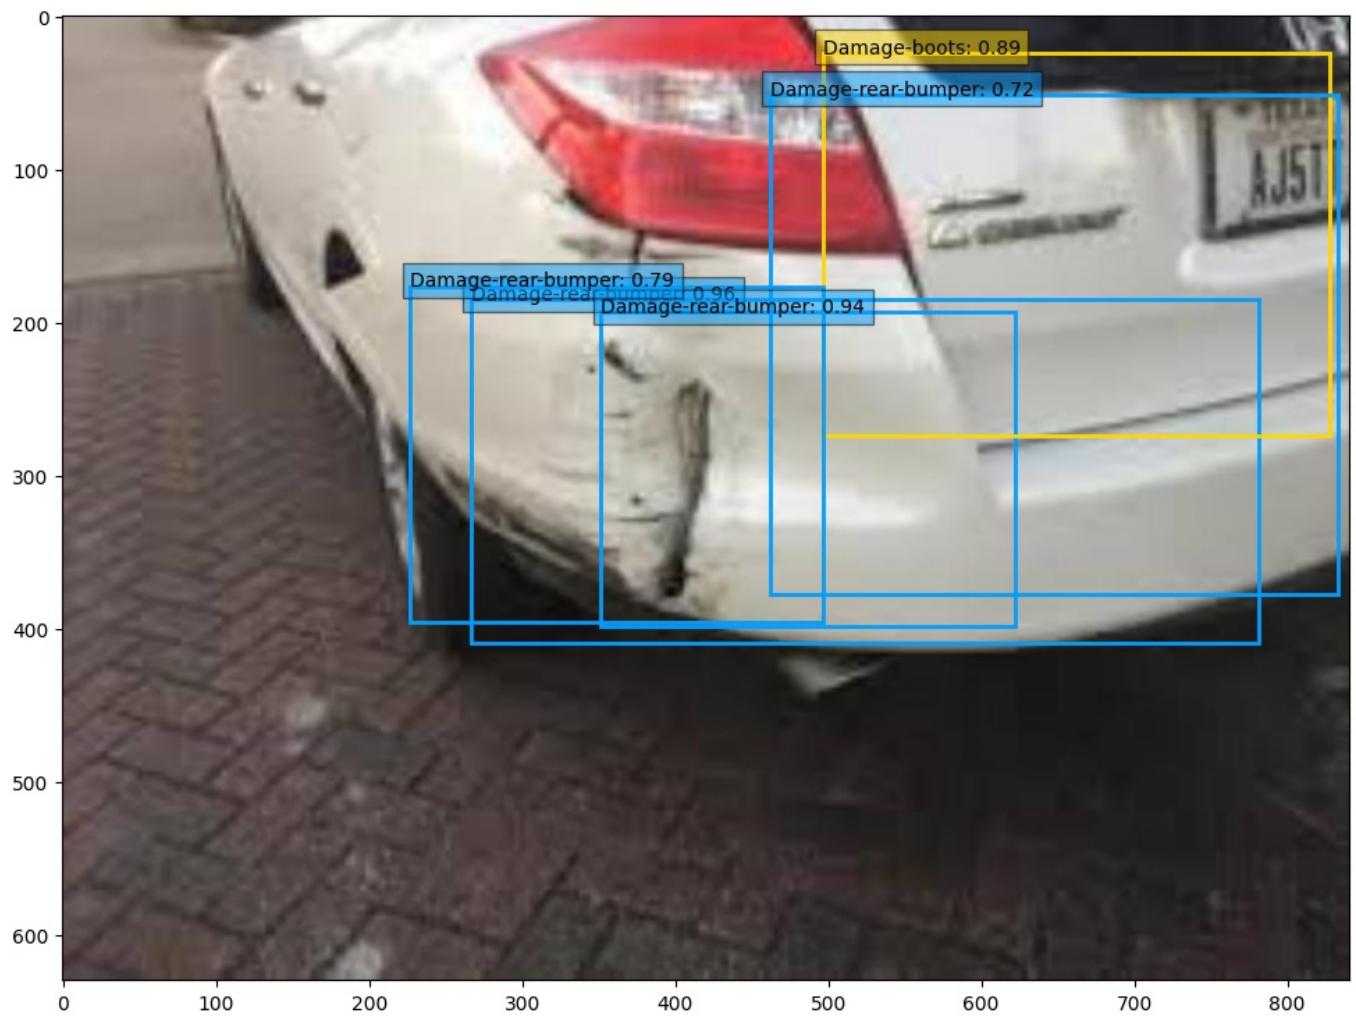


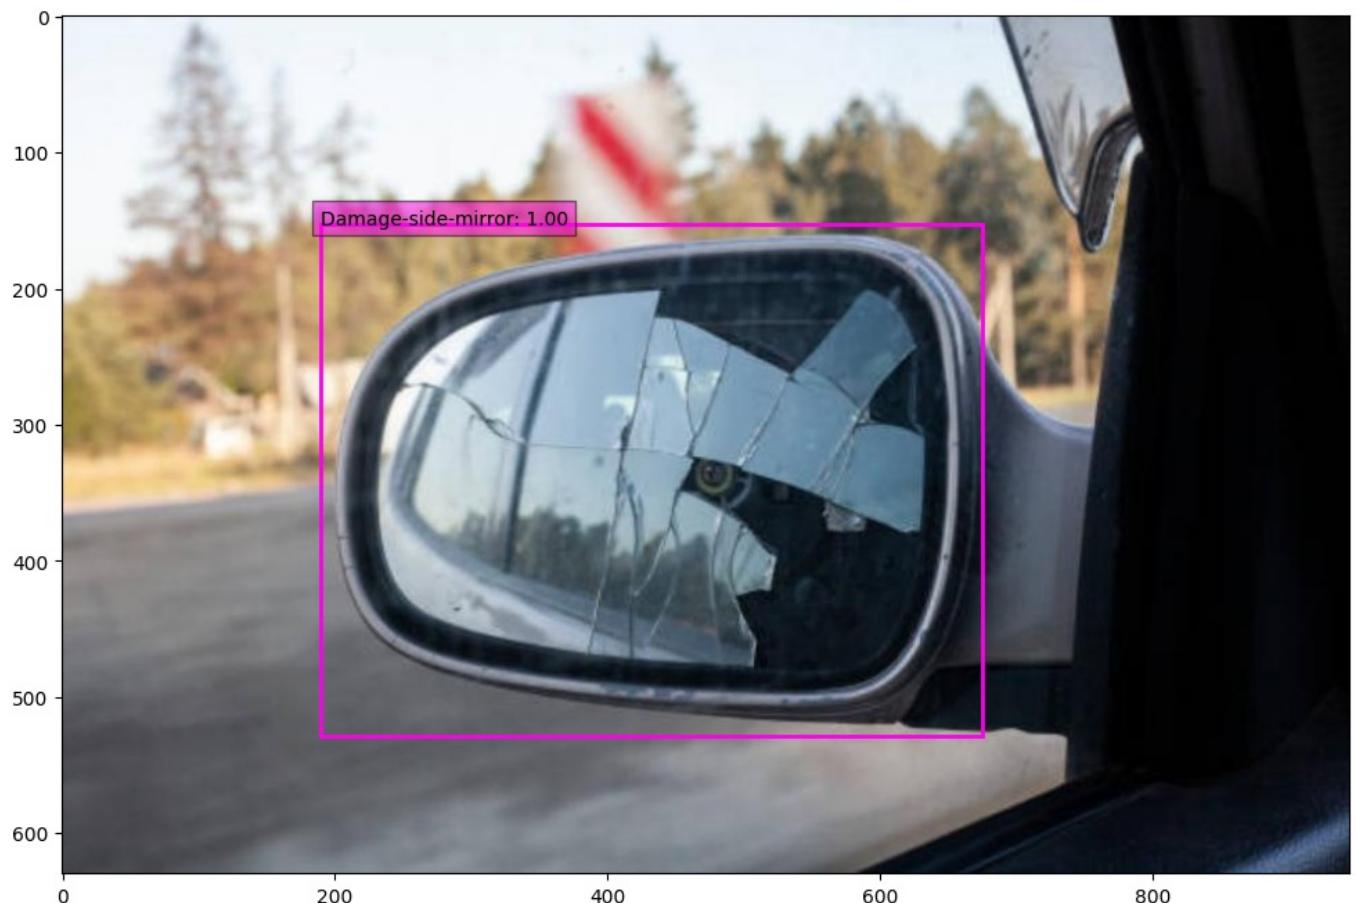














```
KeyboardInterrupt                                     Traceback (most recent call last)
Cell In[23], line 16
      13         visualize_predictions(image_pil, output, train_label_map,color_map) # Pass the label ma
p
      15 # Correct function call with train_label_map
--> 16 run_inference(model, test_loader, train_label_map, device)

Cell In[23], line 7, in run_inference(model, data_loader, train_label_map, device)
      5 for images, _ in data_loader:
      6     images = [img.to(device) for img in images]
-->  7     outputs = model(images)
      9     for i, output in enumerate(outputs):
     10         image = images[i].cpu().permute(1, 2, 0).numpy()

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\torch\nn\modules\module.py:1501, in Module._call_impl(self, *args, **kwargs)
```

```
1496 # If we don't have any hooks, we want to skip the rest of the logic in
1497 # this function, and just call forward.
1498 if not (self._backward_hooks or self._backward_pre_hooks or self._forward_hooks or self._forward_pre_hooks):
1499     or _global_backward_hooks or _global_backward_hooks
1500     or global_forward_hooks or global_forward_pre_hooks):
-> 1501     return forward_call(*args, **kwargs)
1502 # Do not call functions when jit is used
1503 full_backward_hooks, non_full_backward_hooks = [], []
```

```
File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python
n311\site-packages\torchvision\models\detection\generalized_rcnn.py:101, in GeneralizedRCNN.forward(self, images
, targets)
 94         degen_bb: List[float] = boxes[bb_idx].tolist()
 95         torch._assert(
 96             False,
 97             "All bounding boxes should have positive height and width."
 98             f" Found invalid box {degen_bb} for target at index {target_idx}.",
 99         )
--> 101 features = self.backbone(images.tensors)
102 if isinstance(features, torch.Tensor):
103     features = OrderedDict([("0", features)])
```

```
File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python
n311\site-packages\torch\nn\modules\module.py:1501, in Module._call_impl(self, *args, **kwargs)
1496 # If we don't have any hooks, we want to skip the rest of the logic in
1497 # this function, and just call forward.
1498 if not (self._backward_hooks or self._backward_pre_hooks or self._forward_hooks or self._forward_pre_hooks):
ks
1499     or _global_backward_hooks or _global_backward_hooks
1500     or global_forward_hooks or global_forward_pre_hooks):
-> 1501     return forward_call(*args, **kwargs)
1502 # Do not call functions when jit is used
1503 full_backward_hooks, non_full_backward_hooks = [], []
```

```
File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python
n311\site-packages\torchvision\models\detection\backbone_utils.py:57, in BackboneWithFPN.forward(self, x)
 56 def forward(self, x: Tensor) -> Dict[str, Tensor]:
---> 57     x = self.body(x)
 58     x = self.fpn(x)
 59     return x
```

```
File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python
n311\site-packages\torch\nn\modules\module.py:1501, in Module._call_impl(self, *args, **kwargs)
1496 # If we don't have any hooks, we want to skip the rest of the logic in
1497 # this function, and just call forward.
1498 if not (self._backward_hooks or self._backward_pre_hooks or self._forward_hooks or self._forward_pre_hooks):
ks
1499     or _global_backward_hooks or _global_backward_hooks
1500     or global_forward_hooks or global_forward_pre_hooks):
-> 1501     return forward_call(*args, **kwargs)
1502 # Do not call functions when jit is used
1503 full_backward_hooks, non_full_backward_hooks = [], []
```

```
File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python
n311\site-packages\torchvision\models\_utils.py:69, in IntermediateLayerGetter.forward(self, x)
 67 out = OrderedDict()
 68 for name, module in self.items():
---> 69     x = module(x)
 70     if name in self.return_layers:
 71         out_name = self.return_layers[name]
```

```
File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python
n311\site-packages\torch\nn\modules\module.py:1501, in Module._call_impl(self, *args, **kwargs)
1496 # If we don't have any hooks, we want to skip the rest of the logic in
1497 # this function, and just call forward.
1498 if not (self._backward_hooks or self._backward_pre_hooks or self._forward_hooks or self._forward_pre_hooks):
ks
1499     or _global_backward_hooks or _global_backward_hooks
1500     or global_forward_hooks or global_forward_pre_hooks):
-> 1501     return forward_call(*args, **kwargs)
1502 # Do not call functions when jit is used
1503 full_backward_hooks, non_full_backward_hooks = [], []
```

```
File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python
n311\site-packages\torch\nn\modules\container.py:217, in Sequential.forward(self, input)
 215 def forward(self, input):
 216     for module in self:
---> 217         input = module(input)
 218     return input
```

```
File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python
n311\site-packages\torch\nn\modules\module.py:1501, in Module._call_impl(self, *args, **kwargs)
```

```

1496 # If we don't have any hooks, we want to skip the rest of the logic in
1497 # this function, and just call forward.
1498 if not (self._backward_hooks or self._backward_pre_hooks or self._forward_hooks or self._forward_pre_hooks):
1499     or _global_backward_hooks or _global_backward_hooks
1500     or global_forward_hooks or global_forward_pre_hooks):
-> 1501     return forward_call(*args, **kwargs)
1502 # Do not call functions when jit is used
1503 full_backward_hooks, non_full_backward_hooks = [], []

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\torchvision\models\resnet.py:150, in Bottleneck.forward(self, x)

```

147 out = self.bn1(out)
148 out = self.relu(out)
--> 150 out = self.conv2(out)
151 out = self.bn2(out)
152 out = self.relu(out)

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\torch\nn\modules\module.py:1501, in Module.\_call\_impl(self, \*args, \*\*kwargs)

```

1496 # If we don't have any hooks, we want to skip the rest of the logic in
1497 # this function, and just call forward.
1498 if not (self._backward_hooks or self._backward_pre_hooks or self._forward_hooks or self._forward_pre_hooks):
1499     or _global_backward_hooks or _global_backward_hooks
1500     or global_forward_hooks or global_forward_pre_hooks):
-> 1501     return forward_call(*args, **kwargs)
1502 # Do not call functions when jit is used
1503 full_backward_hooks, non_full_backward_hooks = [], []

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\torch\nn\modules\conv.py:463, in Conv2d.forward(self, input)

```

462 def forward(self, input: Tensor) -> Tensor:
--> 463     return self._conv_forward(input, self.weight, self.bias)

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\torch\nn\modules\conv.py:459, in Conv2d.\_conv\_forward(self, input, weight, bias)

```

455 if self.padding_mode != 'zeros':
456     return F.conv2d(F.pad(input, self._reversed_padding_repeated_twice, mode=self.padding_mode),
457                    weight, bias, self.stride,
458                    pair(0), self.dilation, self.groups)
--> 459 return F.conv2d(input, weight, bias, self.stride,
460                     self.padding, self.dilation, self.groups)

```

**KeyboardInterrupt:**

In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js